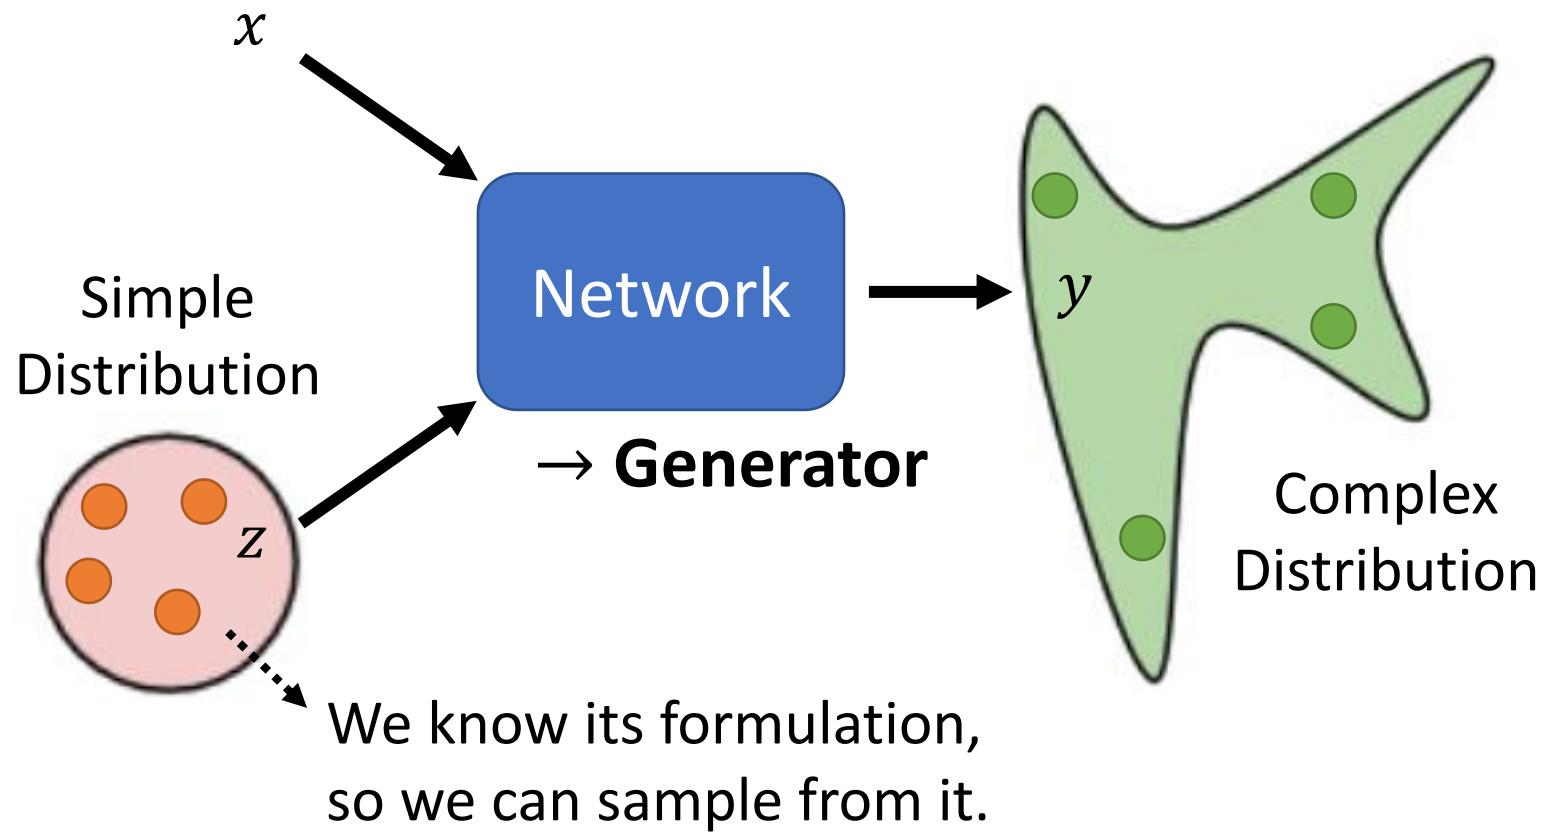


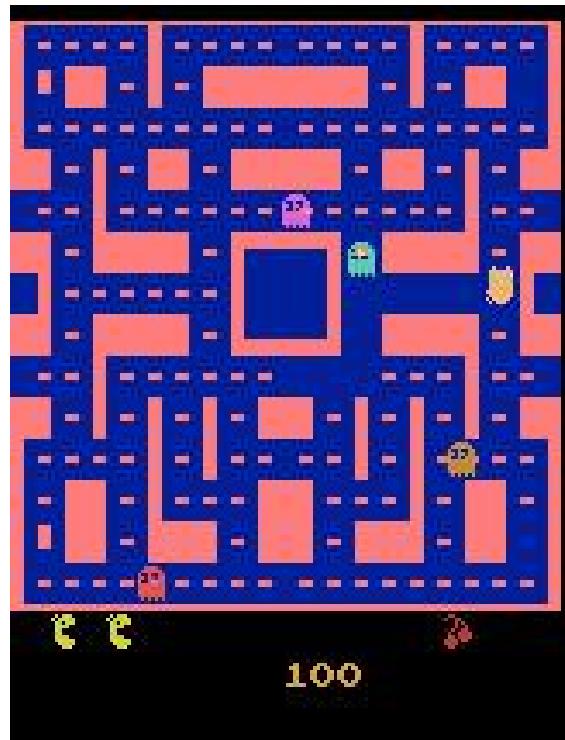
# Generation

Yizhen Lao

# Network as Generator

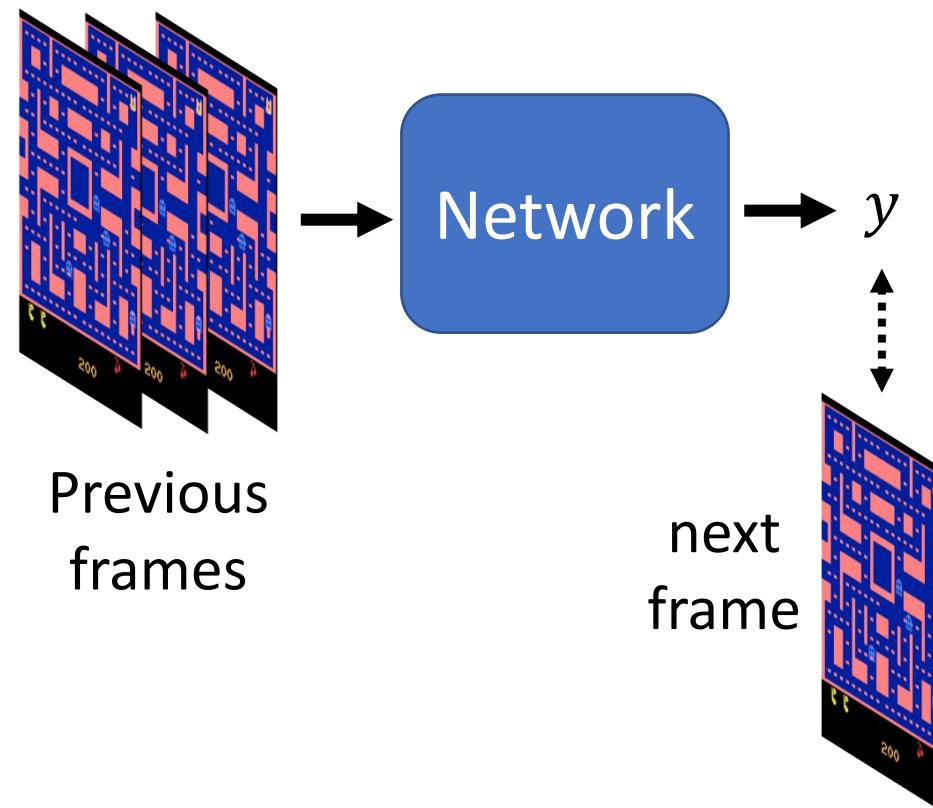


# Why distribution?



Real Video

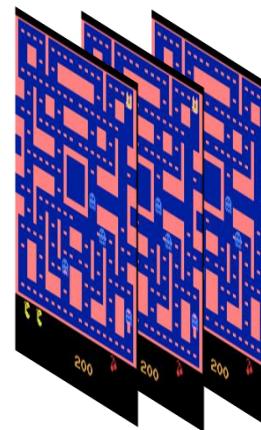
## Video Prediction



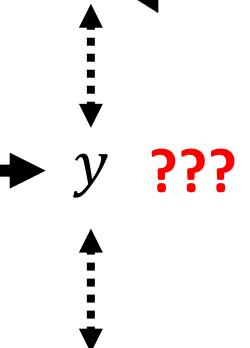
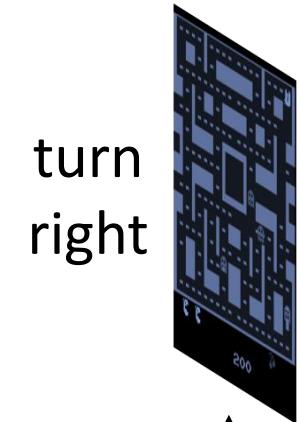
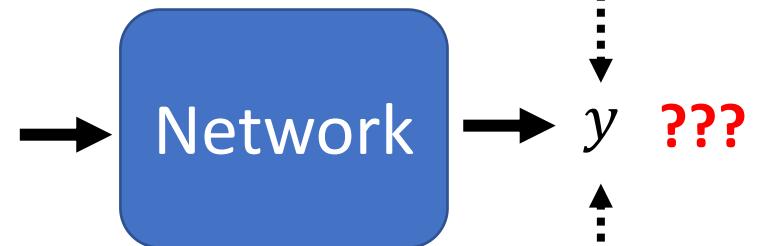
# Why distribution?



## Video Prediction

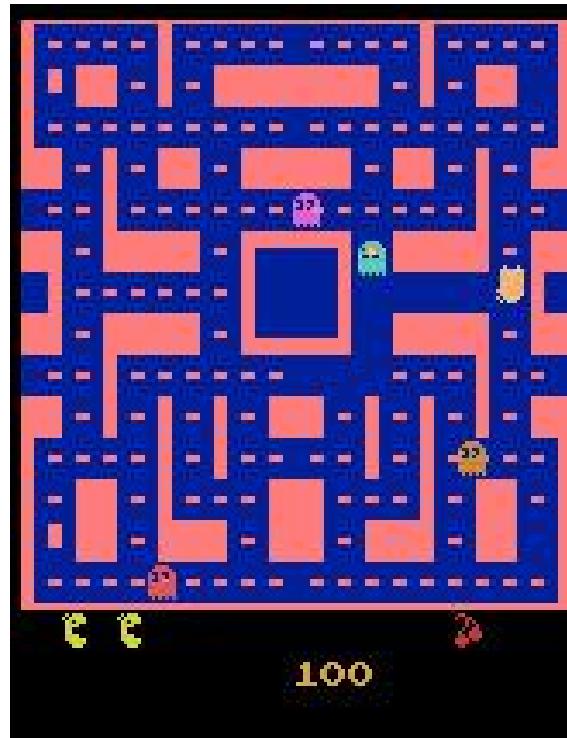


Previous  
frames



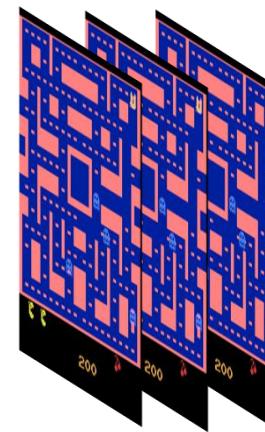
Prediction

# Why distribution?

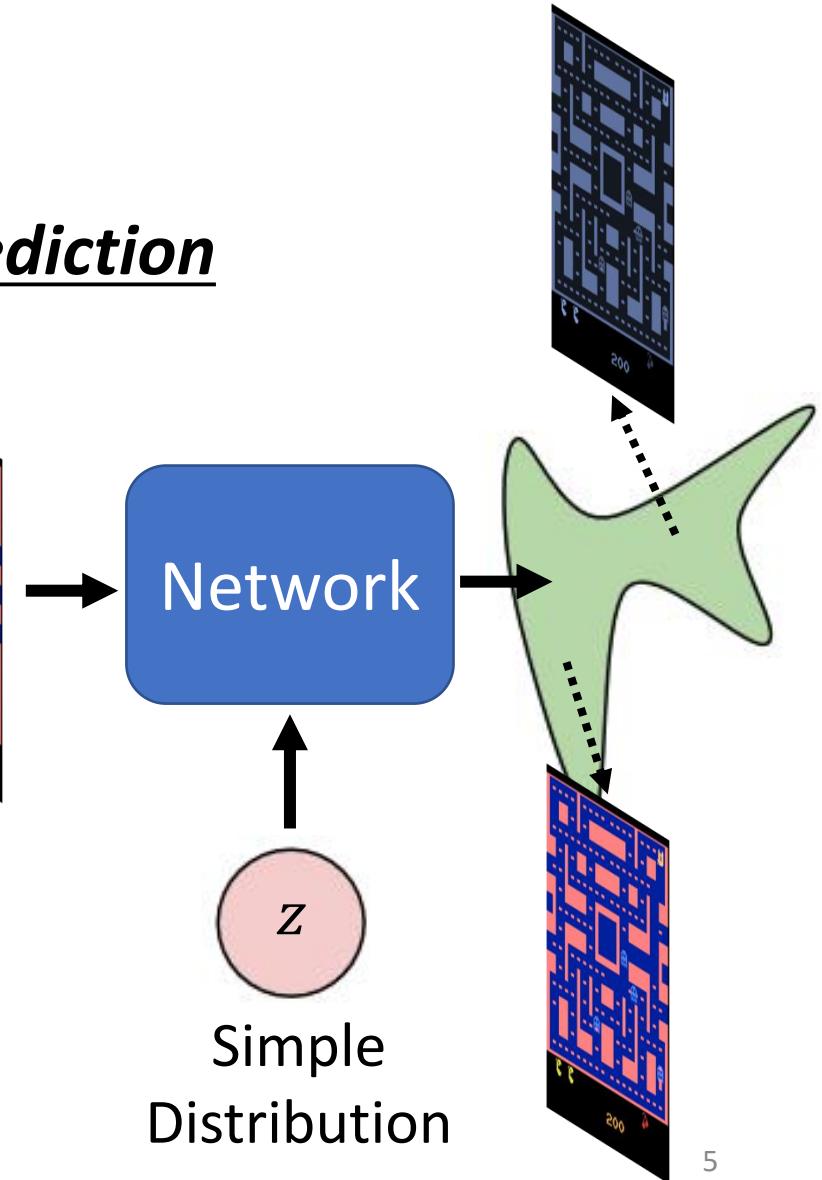


Prediction

## Video Prediction



Previous frames



Simple  
Distribution

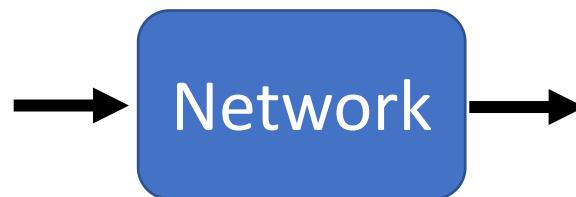
# Why distribution?

(The same input has different outputs.)

- Especially for the tasks needs “*creativity*”

## Drawing

Character  
with red eyes



## Chatbot

你知道辉夜是  
谁吗？



她是秀知院学生会 ...  
她开创了忍者时代 ...

# GAN

# GAN

- How to pronounce “GAN”?



Google 小姐

# All Kinds of GAN ...

<https://github.com/hindupuravinash/the-gan-zoo>

GAN

ACGAN

BGAN

CGAN

DCGAN

EBGAN

fGAN

GoGAN

⋮

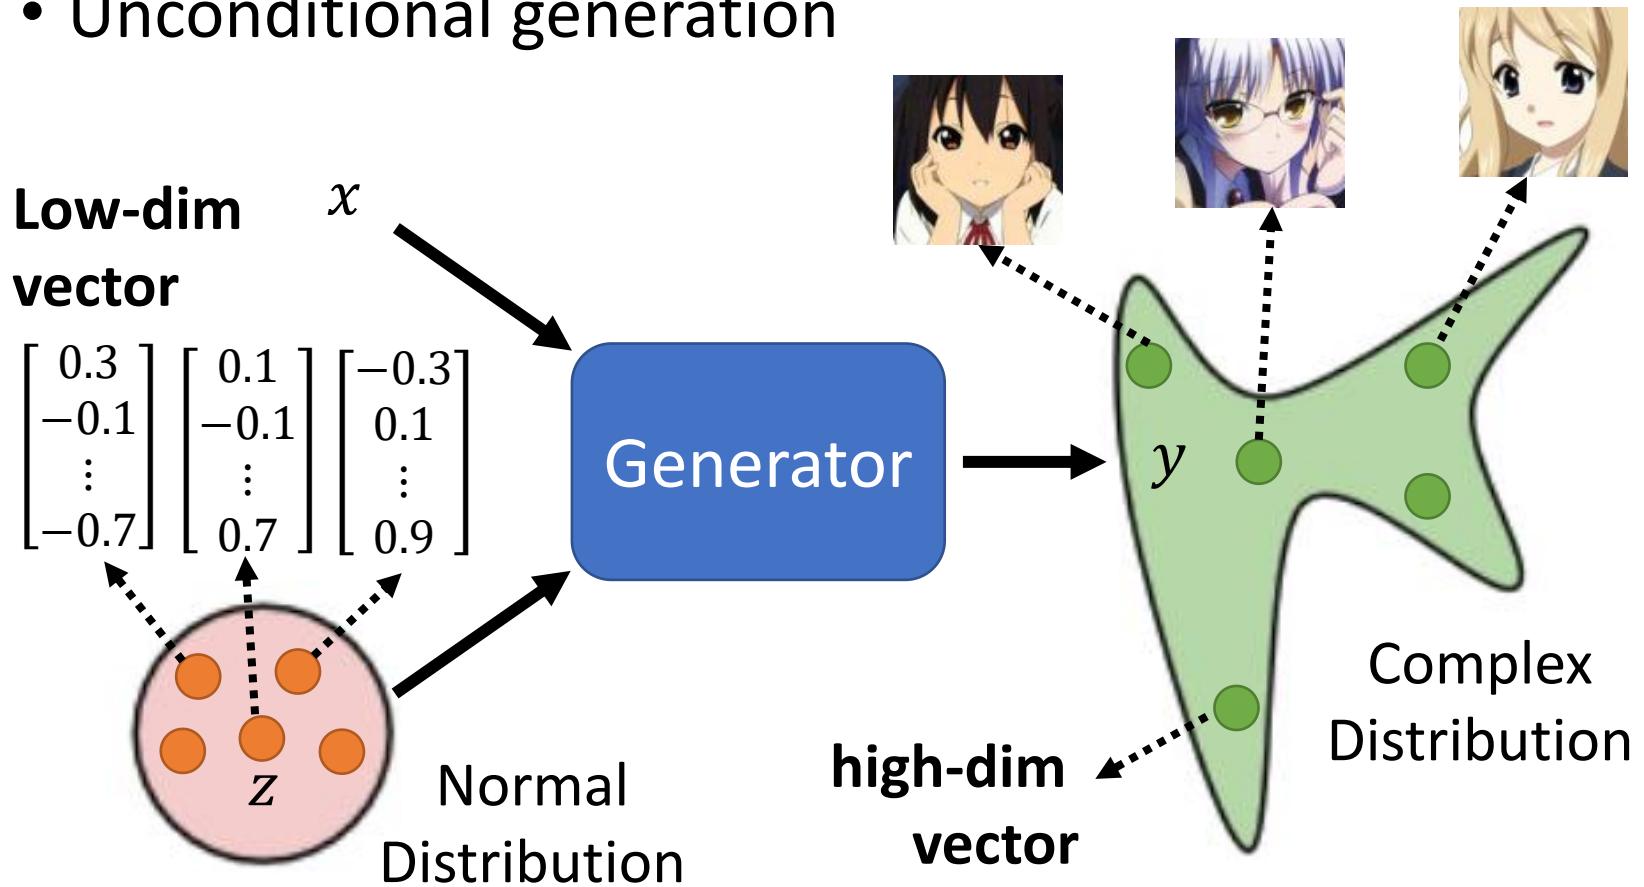
- SeUDA - Semantic-Aware Generative Adversarial Nets for Unsupervised Domain Adaptation Segmentation
- SG-GAN - Semantic-aware Grad-GAN for Virtual-to-Real Urban Scene Adaption ([github](#))
- SG-GAN - Sparsely Grouped Multi-task Generative Adversarial Networks for Facial Attribut
- SGAN - Texture Synthesis with Spatial Generative Adversarial Networks
- SGAN - Stacked Generative Adversarial Networks ([github](#))
- SGAN - Steganographic Generative Adversarial Networks
- SGAN - SGAN: An Alternative Training of Generative Adversarial Networks
- SGAN - CT Image Enhancement Using Stacked Generative Adversarial Networks and T Segmentation Improvement
- sGAN - Generative Adversarial Training for MRA Image Synthesis Using Multi-Contrast
- SiftingGAN - SiftingGAN: Generating and Sifting Labeled Samples to Improve the Rem Classification Baseline in vitro
- SiGAN - SiGAN: Siamese Generative Adversarial Network for Identity-Preserving Face H
- SimGAN - Learning from Simulated and Unsupervised Images through Adversarial Trai
- SisGAN - Semantic Image Synthesis via Adversarial Learning

Mihaela Rosca, Balaji Lakshminarayanan, David Warde-Farley, Shakir Mohamed, “Variational Approaches for Auto-Encoding Generative Adversarial Networks”, arXiv, 2017

<sup>2</sup>We use the Greek  $\alpha$  prefix for  $\alpha$ -GAN, as AEGAN and most other Latin prefixes seem to have been taken  
<https://deephunt.in/the-gan-zoo-79597dc8c347>.

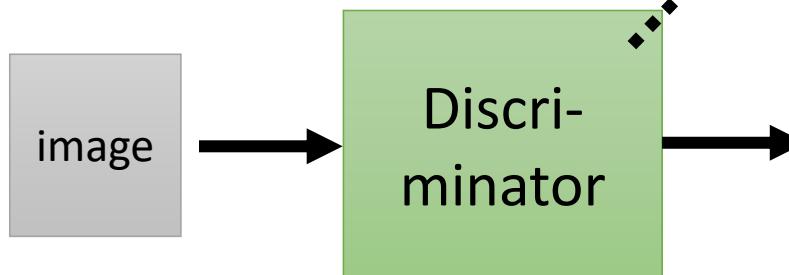
# Anime Face Generation

- Unconditional generation

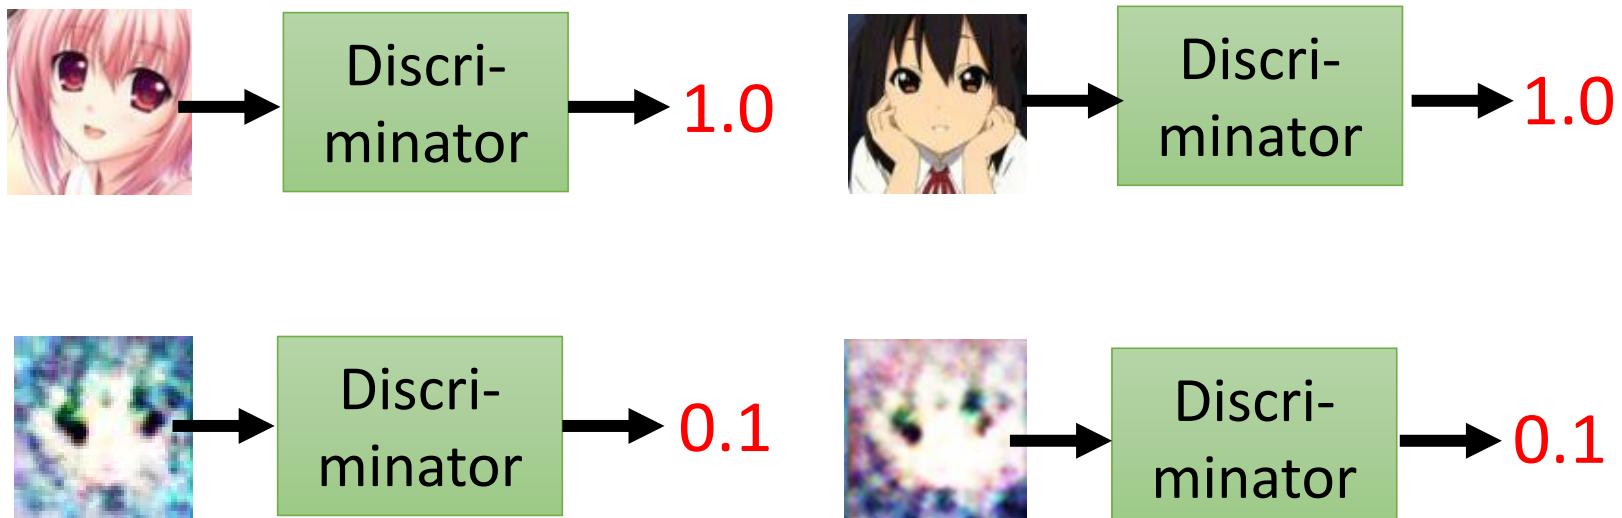


# Discriminator

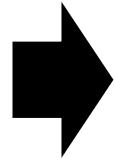
It is a neural network  
(that is, a function).



**Scalar:** Larger means real,  
smaller value fake.



# Basic Idea of GAN



Generator

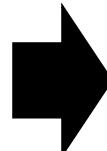
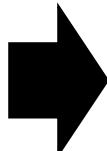
Brown

veins

Butterflies are  
not brown

Butterflies do  
not have veins

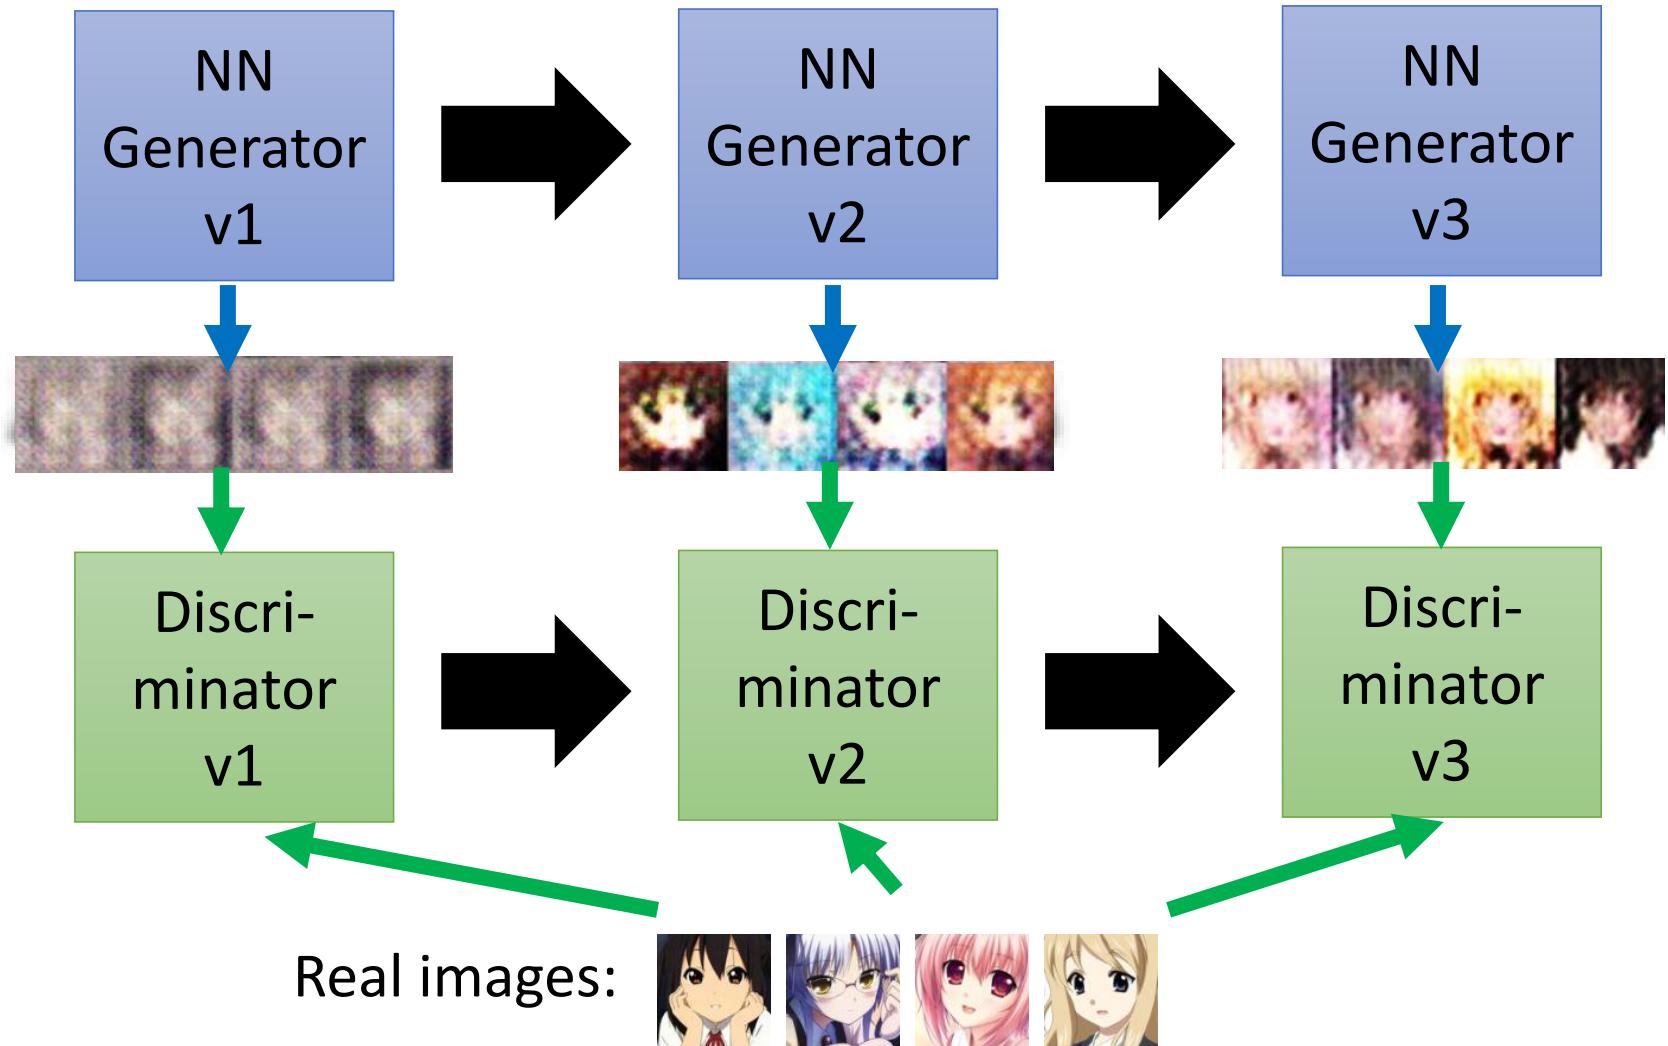
.....



Discriminator

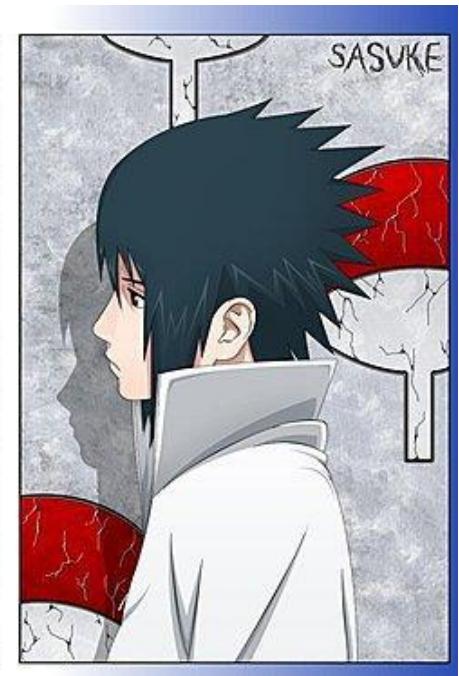
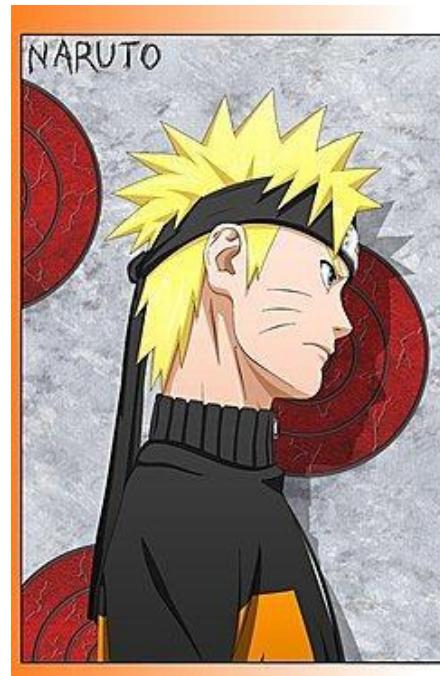
# Basic Idea of GAN

This is where the term  
***“adversarial”*** comes from.



# Basic Idea of GAN

- 写作敌人，念作朋友

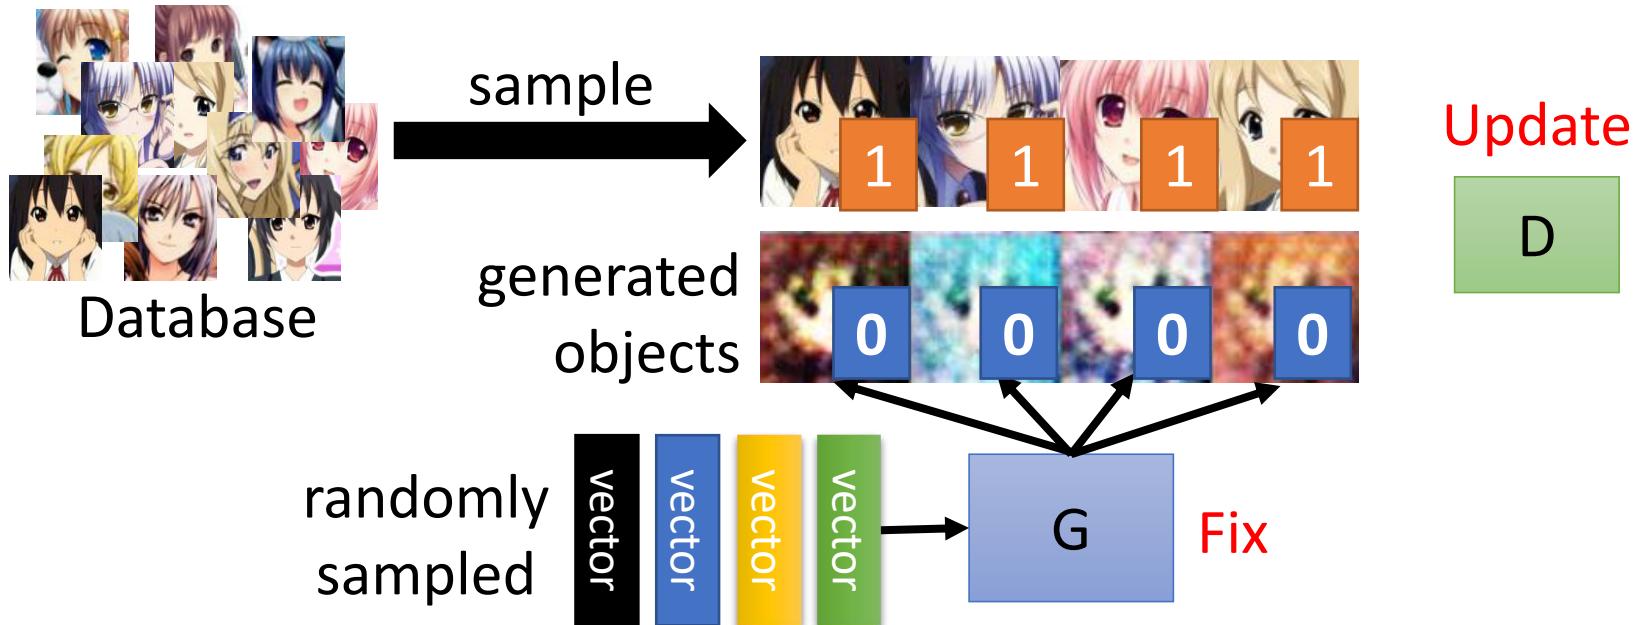


# Algorithm

- Initialize generator and discriminator
- In each training iteration:



**Step 1:** Fix generator G, and update discriminator D



Discriminator learns to assign high scores to real objects and low scores to generated objects.

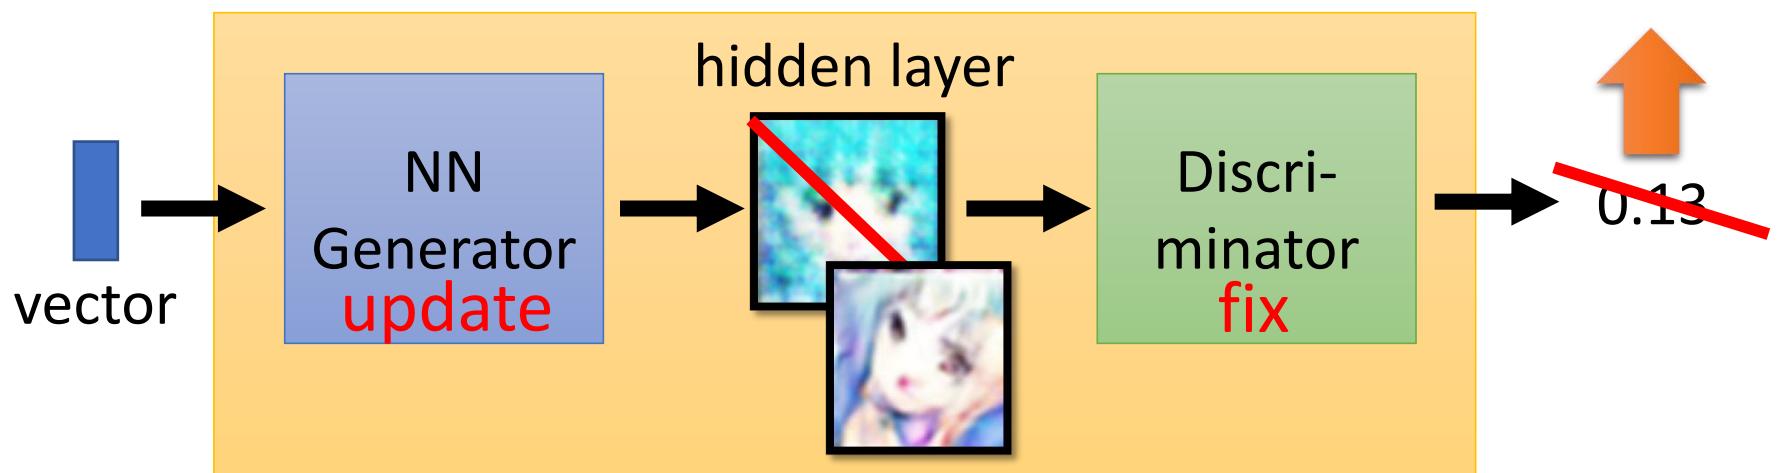
# Algorithm

- Initialize generator and discriminator
- In each training iteration:



**Step 2:** Fix discriminator D, and update generator G

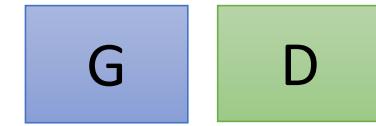
Generator learns to “fool” the discriminator



**large network**

# Algorithm

- Initialize generator and discriminator
- In each training iteration:



Learning  
D

Learning  
G

Sample some  
real objects:

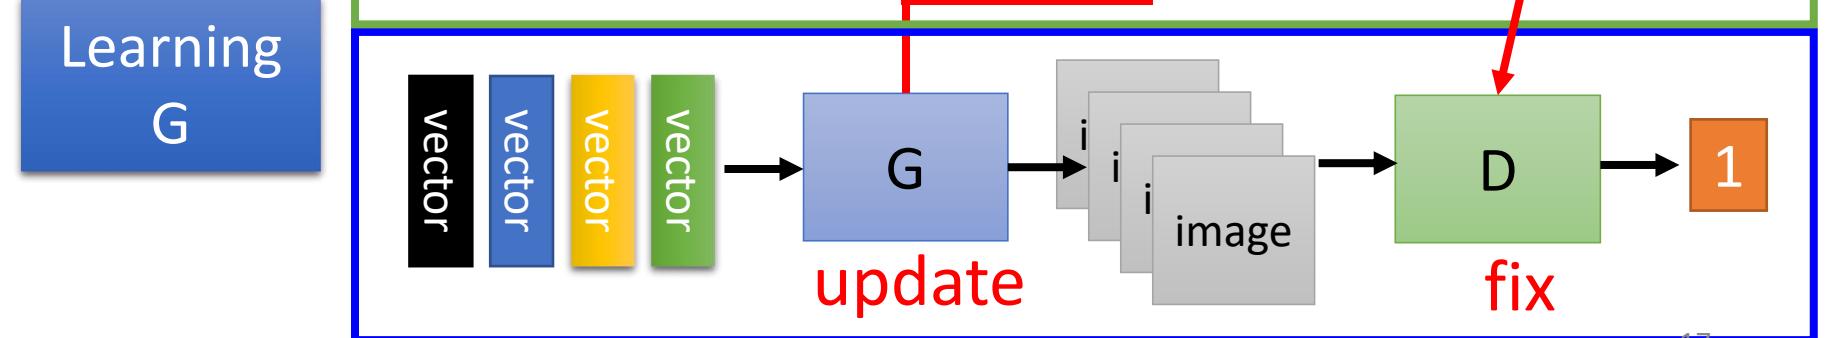
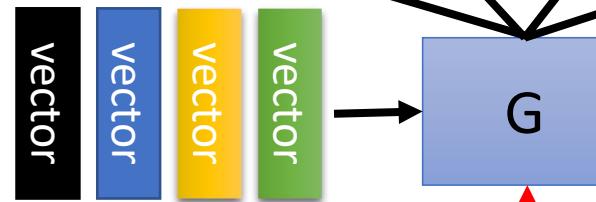


Generate some  
fake objects:



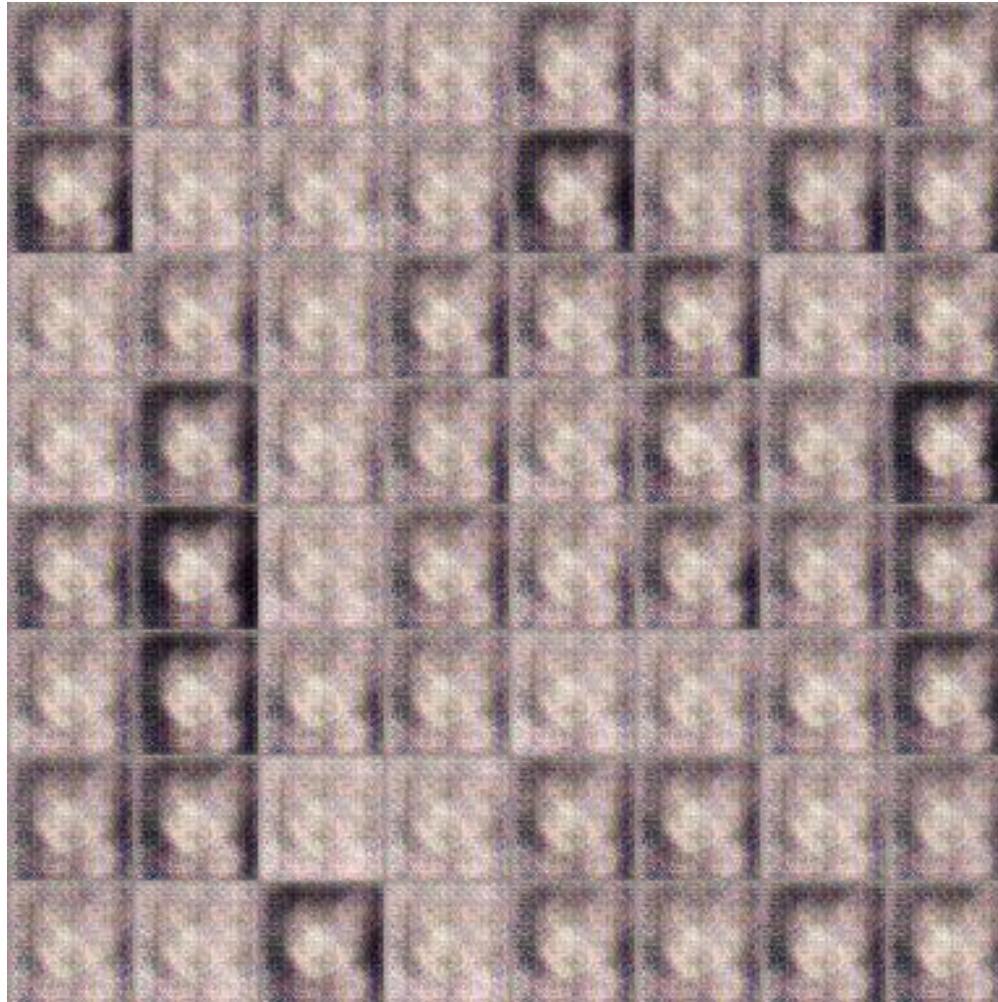
Update  
D

fix



# Anime Face Generation

100 updates



Source of training data: <https://zhuanlan.zhihu.com/p/24767059>

# Anime Face Generation



1000 updates

# Anime Face Generation

2000 updates



# Anime Face Generation

5000 updates



# Anime Face Generation



10,000 updates

# Anime Face Generation



20,000 updates

# Anime Face Generation



50,000 updates



The faces  
generated by  
machine.

# In 2019, with StyleGAN .....



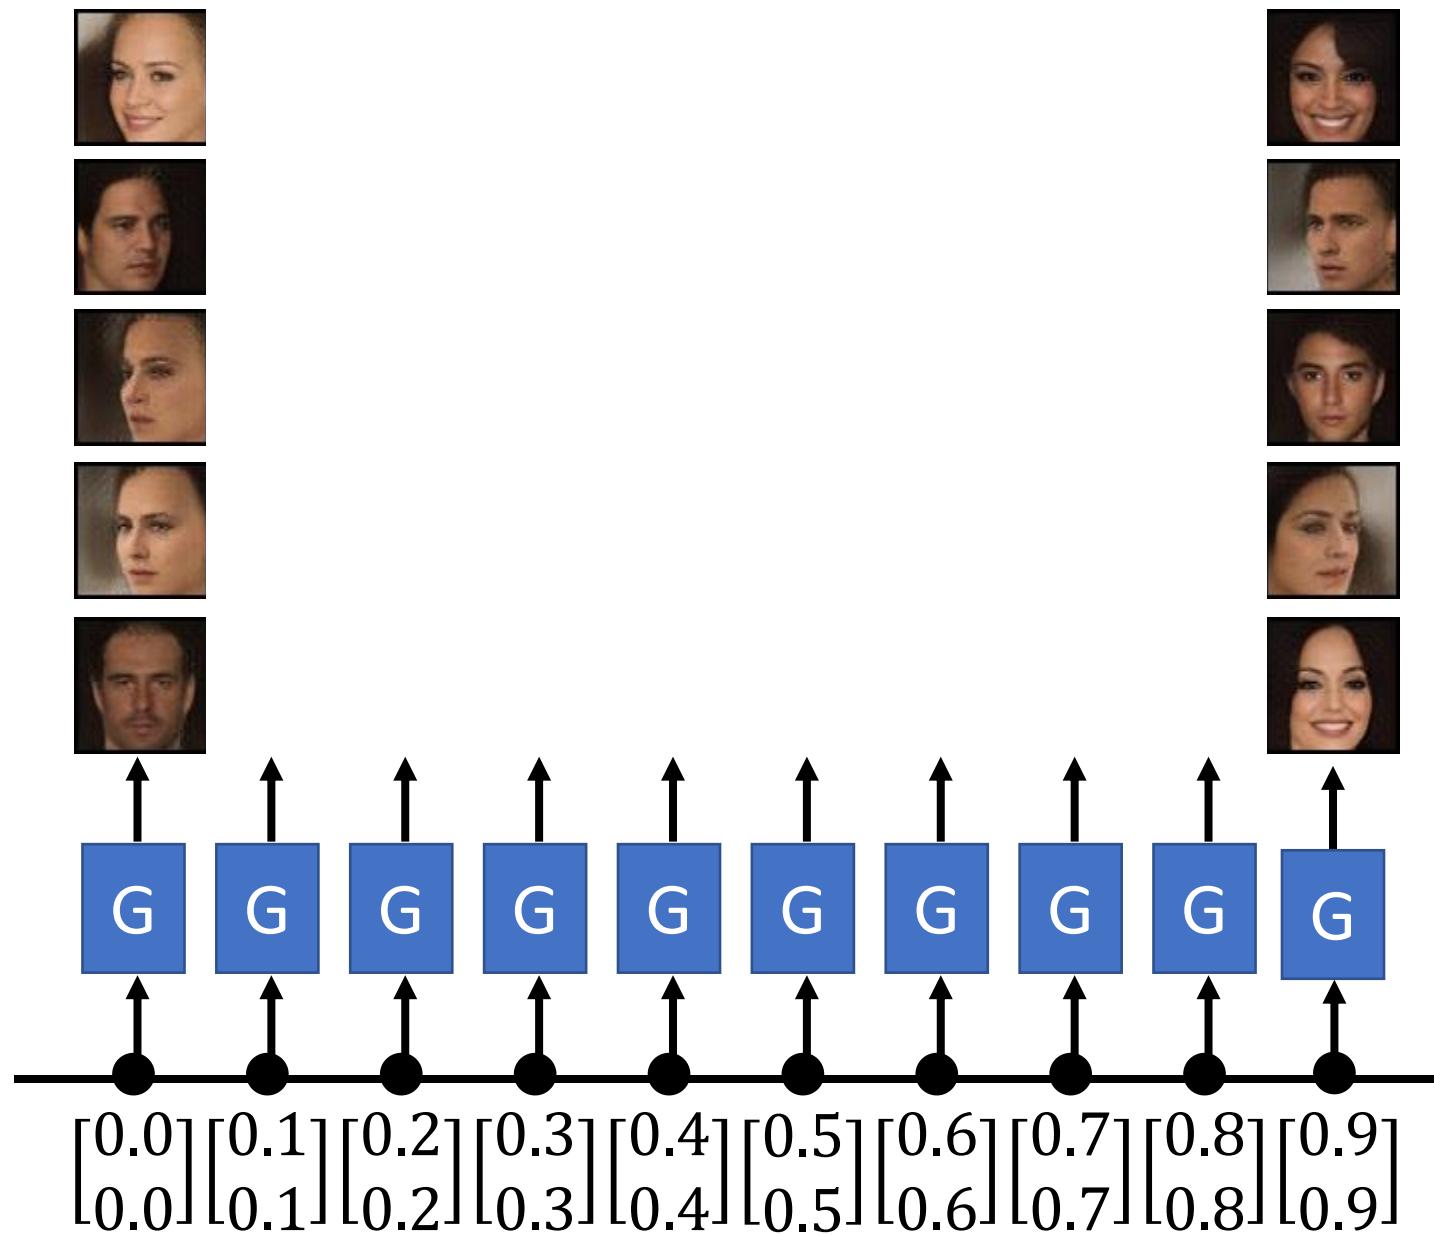
Source of video:

<https://www.gwern.net/Faces>



# Progressive GAN

<https://arxiv.org/abs/1710.10196>





# The first GAN |

<https://arxiv.org/abs/1406.2661> (Ian J.  
Goodfellow) 29



# Today ..... BigGAN |

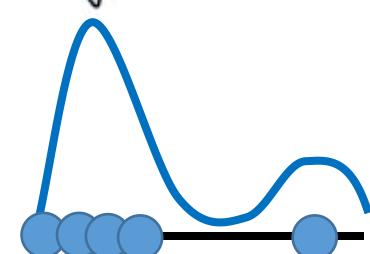
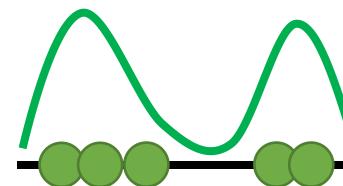
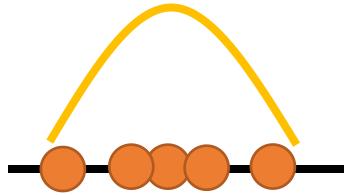
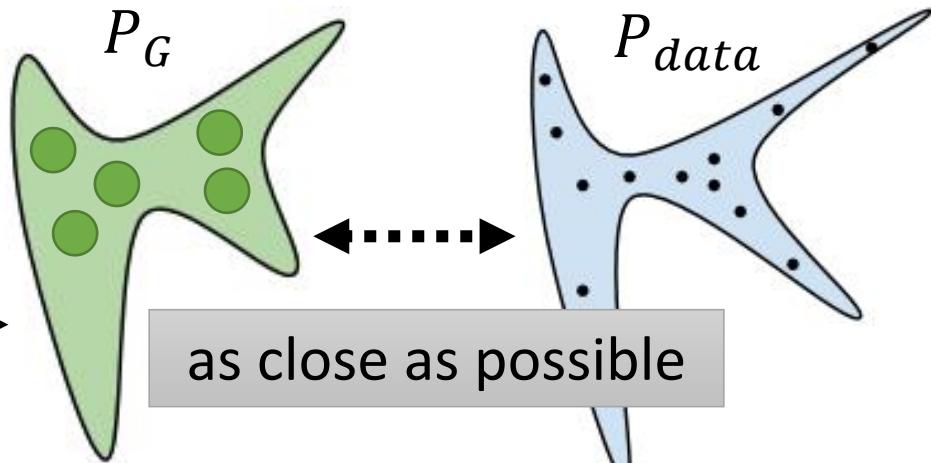
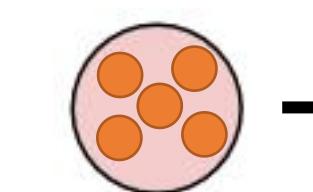
<https://arxiv.org/abs/1809.11096>

# Theory behind GAN

c.f.  $w^*, b^* = \operatorname{argmin}_{w,b} L$

# Our Objective

Normal  
Distribution



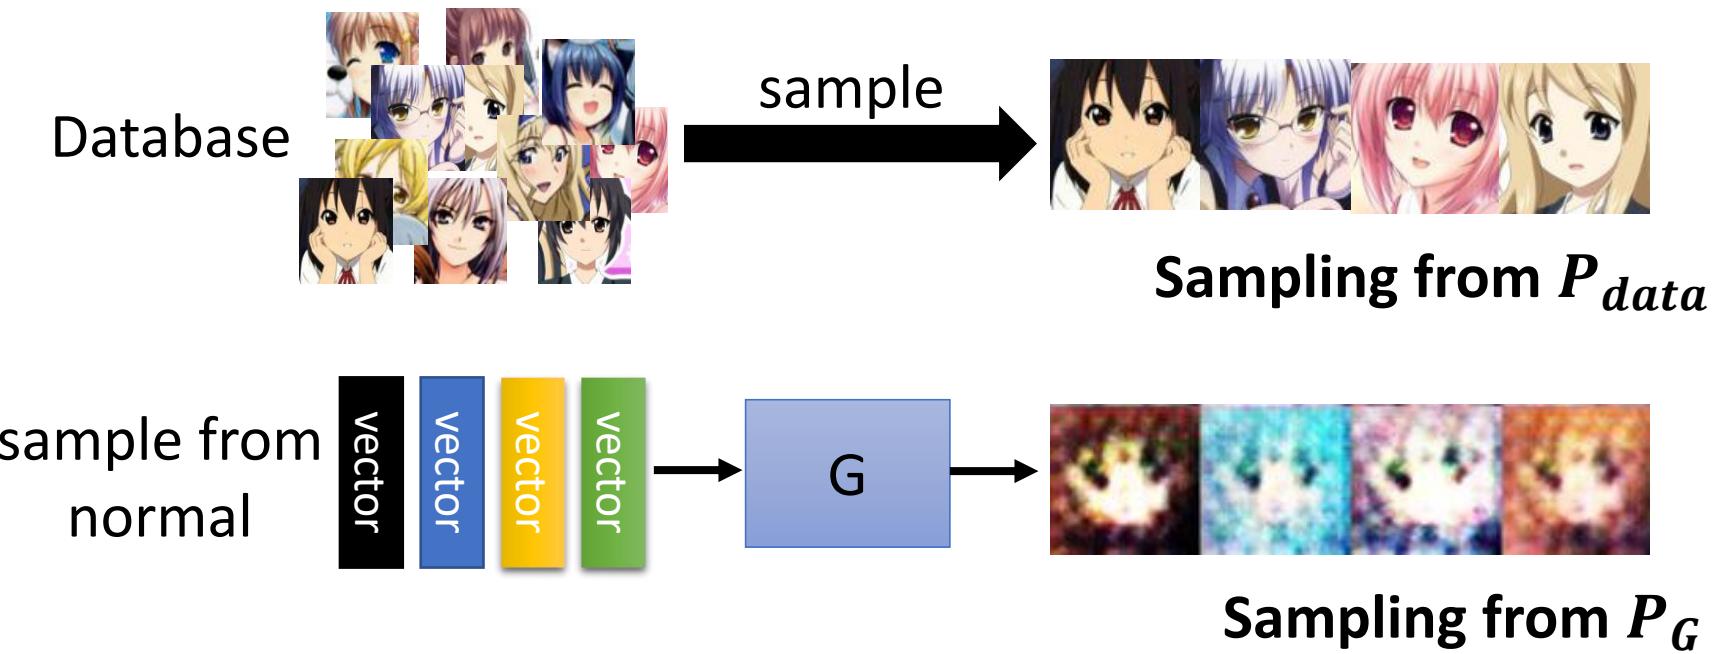
$$G^* = \operatorname{argmin}_G \underline{\operatorname{Div}}(P_G, P_{data})$$

Divergence between distributions  $P_G$  and  $P_{data}$   
How to compute the divergence?

# Sampling is good enough .....

$$G^* = \underset{G}{\operatorname{argmin}} \operatorname{Div}(P_G, P_{data})$$

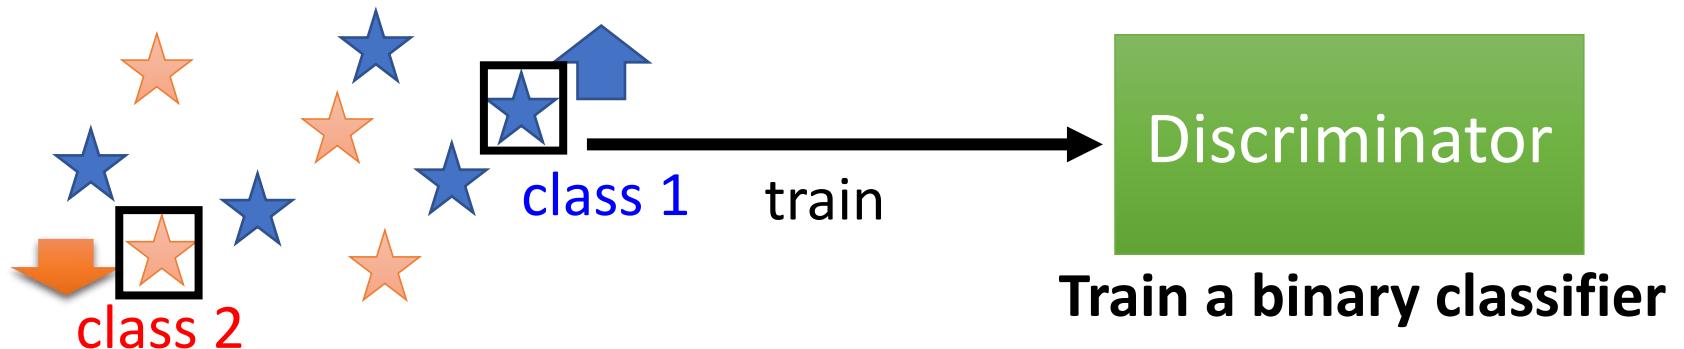
Although we do not know the distributions of  $P_G$  and  $P_{data}$ , we can sample from them.



# Discriminator

$$G^* = \operatorname{argmin}_G \text{Div}(P_G, P_{\text{data}})$$

★ : data sampled from  $P_{\text{data}}$       ★ : data sampled from  $P_G$



Training:  $D^* = \operatorname{argmax}_D V(D, G)$

The value is related to JS divergence.

## Objective Function for D

$$V(G, D) = E_{y \sim P_{\text{data}}} [\log D(y)] + E_{y \sim P_G} [\log(1 - D(y))]$$

$D^* = \operatorname{argmax}_D V(D, G)$   
negative cross entropy

=

Training classifier:  
minimize cross entropy

# Discriminator

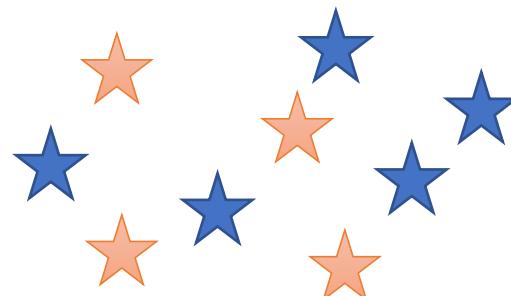
$$G^* = \underset{G}{\operatorname{argmin}} \operatorname{Div}(P_G, P_{data})$$

★ : data sampled from  $P_{data}$

☆ : data sampled from  $P_G$

**Training:**

$$D^* = \underset{D}{\operatorname{argmax}} V(D, G)$$



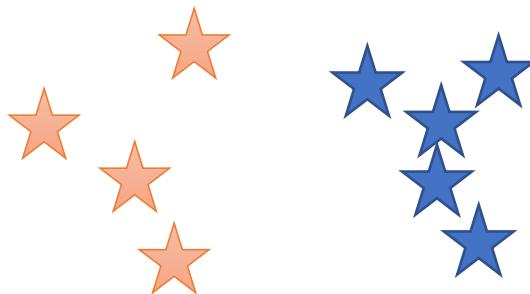
small divergence

train



hard to discriminate

$$\text{Small } \underset{D}{\operatorname{max}} V(D, G)$$



large divergence

train



easy to discriminate

$$G^* = \operatorname{argmin}_G \max_D V(G, D)$$

$$D^* = \operatorname{argmax}_D V(D, G)$$

The maximum objective value is related to JS divergence.

- Initialize generator and discriminator
- In each training iteration:

**Step 1**: Fix generator G, and update discriminator D

**Step 2**: Fix discriminator D, and update generator G

# Can we use other divergence?

| Name                     | $D_f(P\ Q)$   | Generator $f(u)$   |
|--------------------------|---|--|
| Total variation          | $\frac{1}{2} \int  p(x) - q(x)  dx$   | $\frac{1}{2} u - 1 $                                     |
| Kullback-Leibler         | $\int p(x) \log \frac{p(x)}{q(x)} dx$   | $u \log u$   |
| Reverse Kullback-Leibler | $\int q(x) \log \frac{q(x)}{p(x)} dx$   | $-\log u$  |
| Pearson $\chi^2$         | $\int \frac{(q(x)-p(x))^2}{p(x)} dx$  | $(u - 1)^2$  |
| Neyman $\chi^2$          | $\int \frac{(p(x)-q(x))^2}{q(x)} dx$  | $\frac{(1-u)^2}{u}$                                      |
| Squared Hellinger        | $\int \left( \sqrt{p(x)} - \sqrt{q(x)} \right)^2 dx$  | $(\sqrt{u} - 1)^2$                                       |
| Jeffrey                  | $\int (p(x) - q(x)) \log \left( \frac{p(x)}{q(x)} \right) dx$   | $(u - 1) \log u$   |
| Jensen-Shannon           | $\frac{1}{2} \int p(x) \log \frac{2p(x)}{p(x)+q(x)} + q(x) \log \frac{2q(x)}{p(x)+q(x)} dx$                   | $-(u + 1) \log \frac{1+u}{2} + u \log u$                 |
| Jensen-Shannon-weighted  | $\int p(x)\pi \log \frac{p(x)}{\pi p(x)+(1-\pi)q(x)} + (1-\pi)q(x) \log \frac{q(x)}{\pi p(x)+(1-\pi)q(x)} dx$ | $\pi u \log u - (1 - \pi + \pi u) \log(1 - \pi + \pi u)$ |
| GAN                      | $\int p(x) \log \frac{2p(x)}{p(x)+q(x)} + q(x) \log \frac{2q(x)}{p(x)+q(x)} dx - \log(4)$                     | $u \log u - (u + 1) \log(u + 1)$                         |

| Name                    | Conjugate $f^*(t)$                             |
|-------------------------|--|
| Total variation         | $t$  |
| Kullback-Leibler (KL)   | $\exp(t - 1)$                                  |
| Reverse KL              | $-1 - \log(-t)$                                |
| Pearson $\chi^2$        | $\frac{1}{4}t^2 + t$                           |
| Neyman $\chi^2$         | $2 - 2\sqrt{1-t}$                              |
| Squared Hellinger       | $\frac{t}{1-t}$                                |
| Jeffrey                 | $W(e^{1-t}) + \frac{1}{W(e^{1-t})} + t - 2$    |
| Jensen-Shannon          | $-\log(2 - \exp(t))$                           |
| Jensen-Shannon-weighted | $(1 - \pi) \log \frac{1-\pi}{1-\pi e^{t/\pi}}$ |
| GAN                     | $-\log(1 - \exp(t))$                           |

Using the divergence  
you like ☺

<https://arxiv.org/abs/1606.00709>

GAN is difficult to train .....

**NO PAIN**  
**NO GAN**

# Tips for GAN

# JS divergence is not suitable

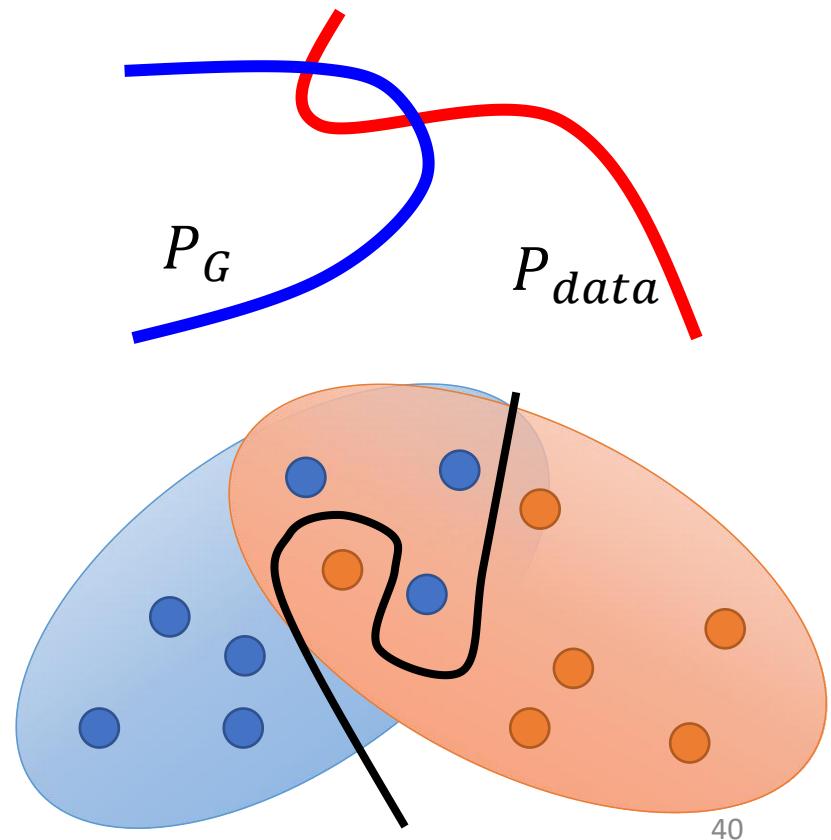
- In most cases,  $P_G$  and  $P_{data}$  are not overlapped.
- 1. The nature of data

Both  $P_{data}$  and  $P_G$  are low-dim manifold in high-dim space.  
The overlap can be ignored.

- 2. Sampling

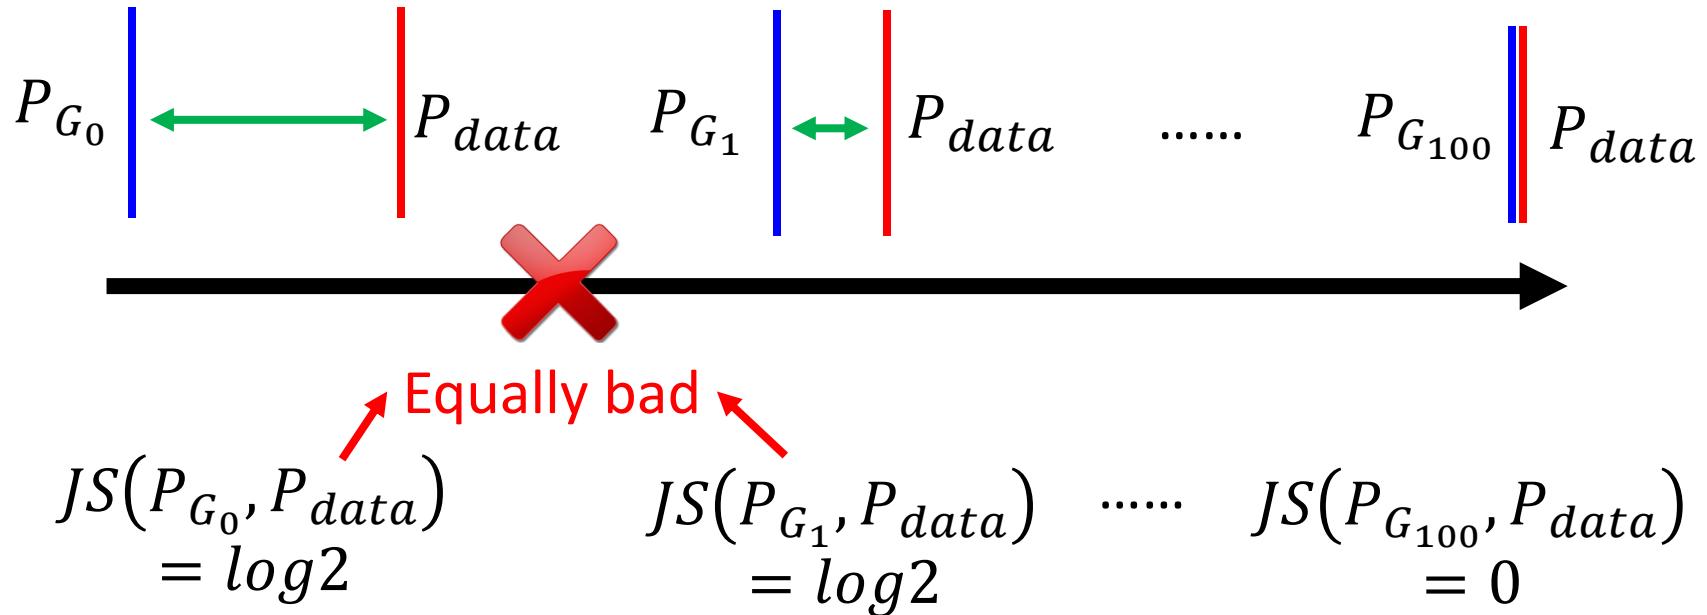
Even though  $P_{data}$  and  $P_G$  have overlap.

If you do not have enough sampling .....



## What is the problem of JS divergence?

JS divergence is always  $\log 2$  if two distributions do not overlap.

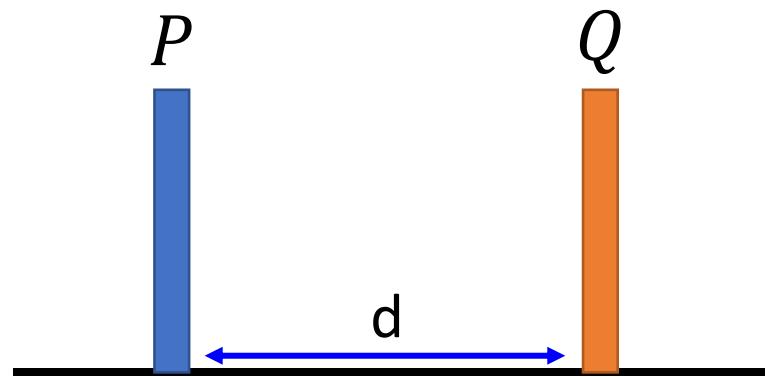


Intuition: If two distributions do not overlap, binary classifier achieves 100% accuracy.

The accuracy (or loss) means nothing during GAN training.

# Wasserstein distance

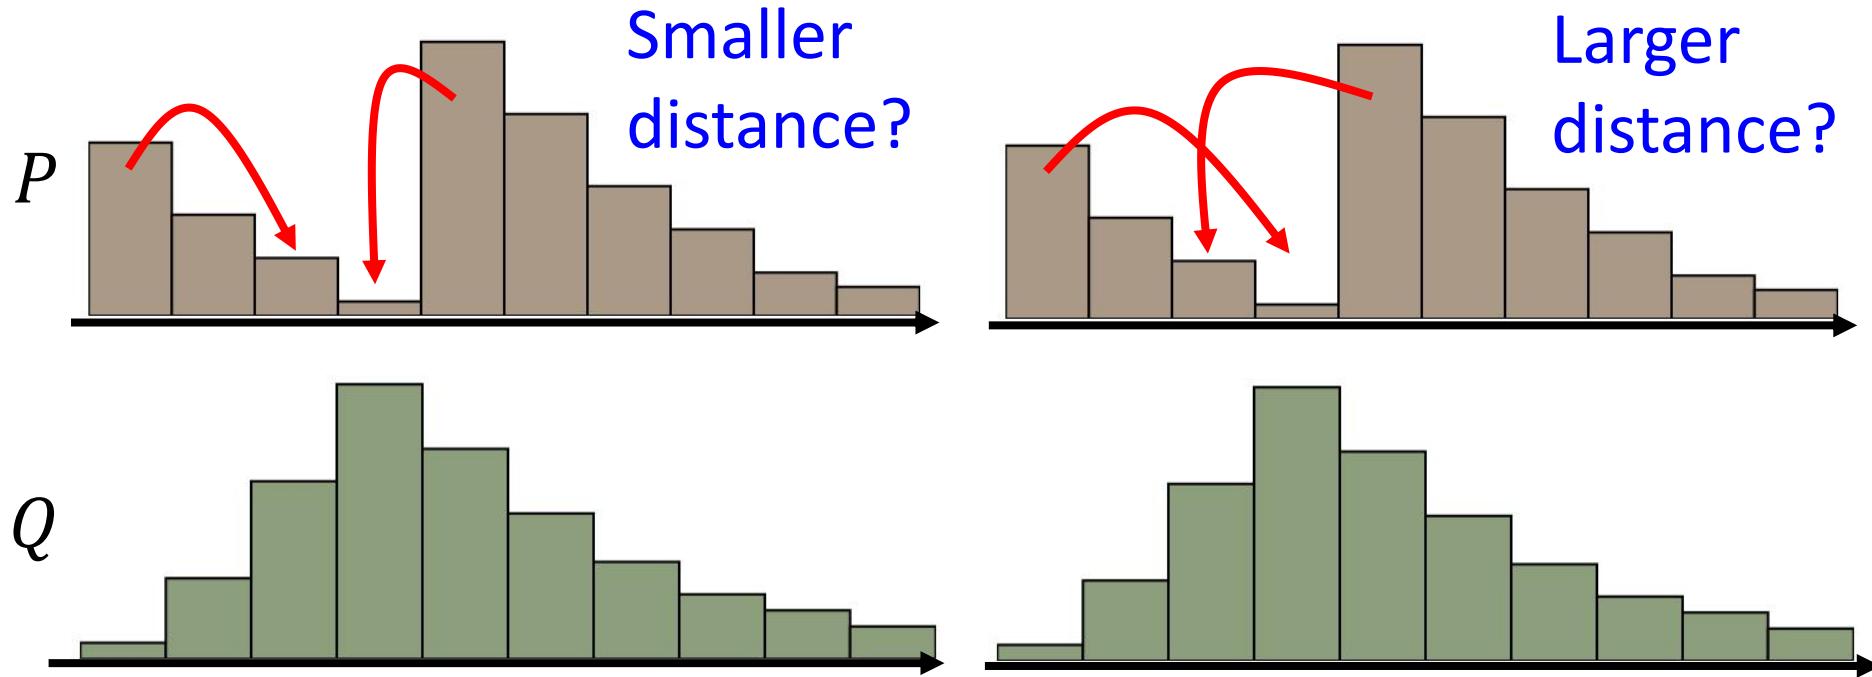
- Considering one distribution P as a pile of earth, and another distribution Q as the target
- The average distance the earth mover has to move the earth.



$$W(P, Q) = d$$



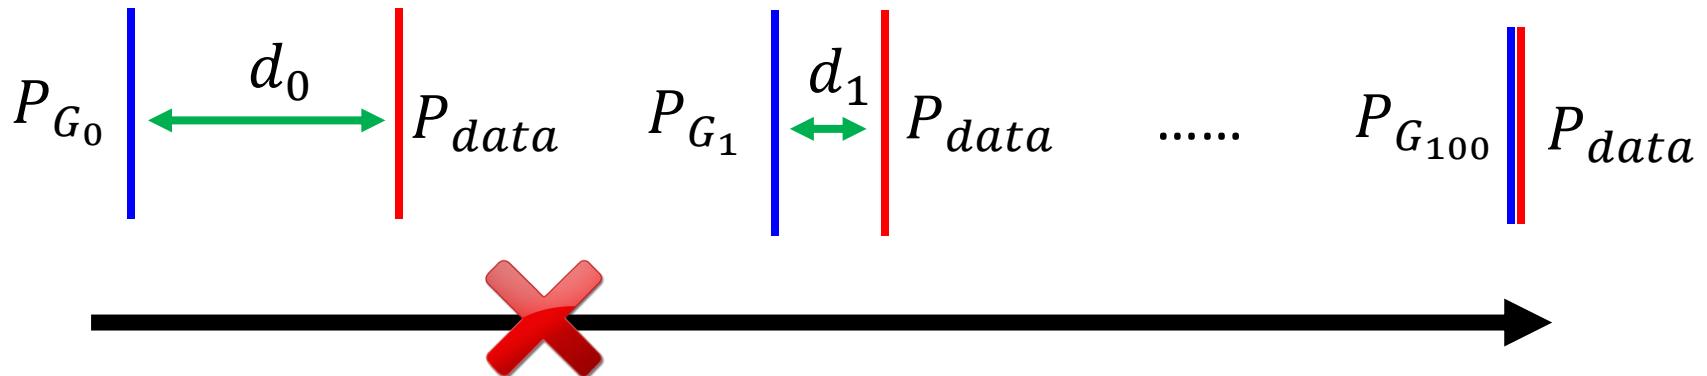
# Wasserstein distance



There are many possible “moving plans”.

Using the “moving plan” with the smallest average distance to define the Wasserstein distance.

# What is the problem of JS divergence?



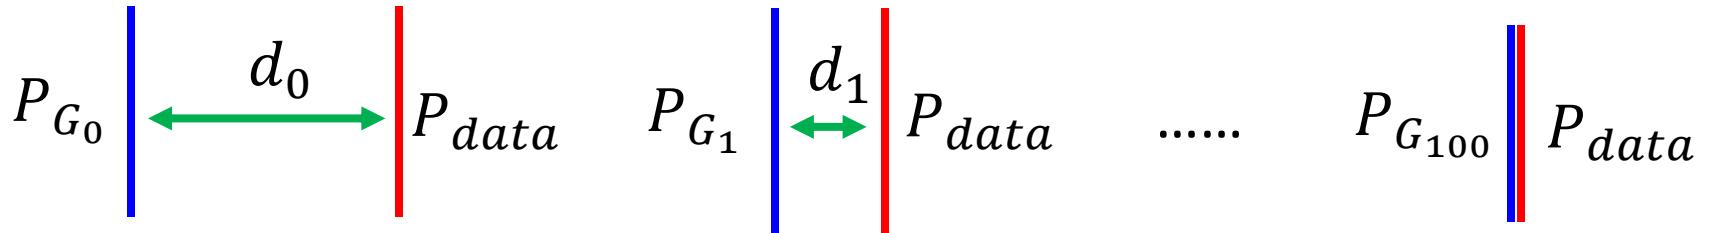
$$JS(P_{G_0}, P_{data}) = \log 2$$
$$JS(P_{G_1}, P_{data}) = \log 2$$
$$\dots$$
$$JS(P_{G_{100}}, P_{data}) = 0$$

$$W(P_{G_0}, P_{data}) = d_0$$
$$W(P_{G_1}, P_{data}) = d_1$$
$$\dots$$
$$W(P_{G_{100}}, P_{data}) = 0$$

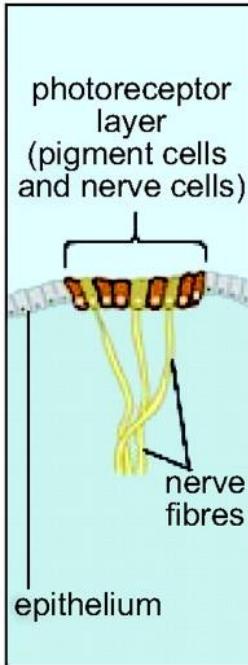
Better!



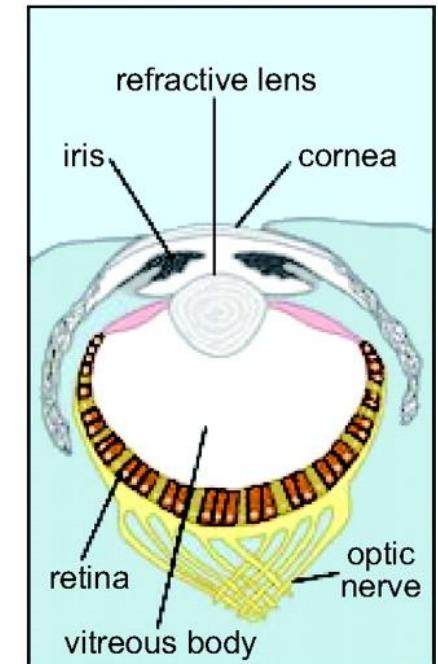
# What is the problem of JS divergence?



pigment spot  
(limpet, *Patella*)



Complex eye  
(octopus)



# WGAN

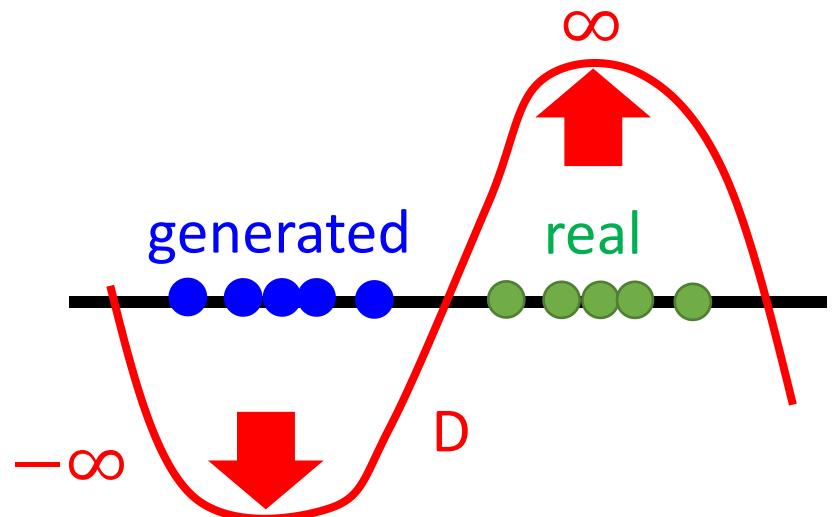
Evaluate Wasserstein distance between  $P_{data}$  and  $P_G$

$$\max_{\substack{D \in 1-\text{Lipschitz}}} \{E_{y \sim P_{data}}[D(y)] - E_{y \sim P_G}[D(y)]\}$$

D has to be smooth enough. How to fulfill this constraint?

Without the constraint, the training of D will not converge.

Keeping the D smooth forces  $D(y)$  become  $\infty$  and  $-\infty$

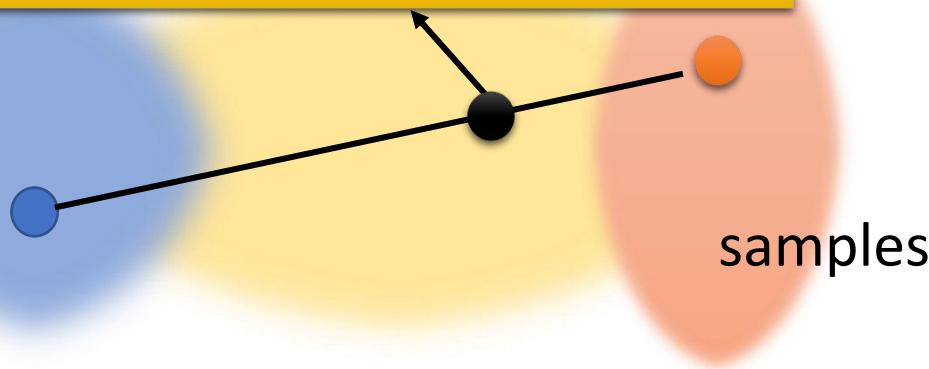


$$\max_{D \in 1\text{-Lipschitz}} \{E_{y \sim P_{data}}[D(y)] - E_{y \sim P_G}[D(y)]\}$$

- Original WGAN → Weight
  - Force the parameters w between c and -c
  - After parameter update, if  $w > c$ ,  $w = c$ ; if  $w < -c$ ,  $w = -c$
- Improved WGAN → Gradient Penalty

Keep the gradient close to 1

<https://arxiv.org/abs/1704.00028>

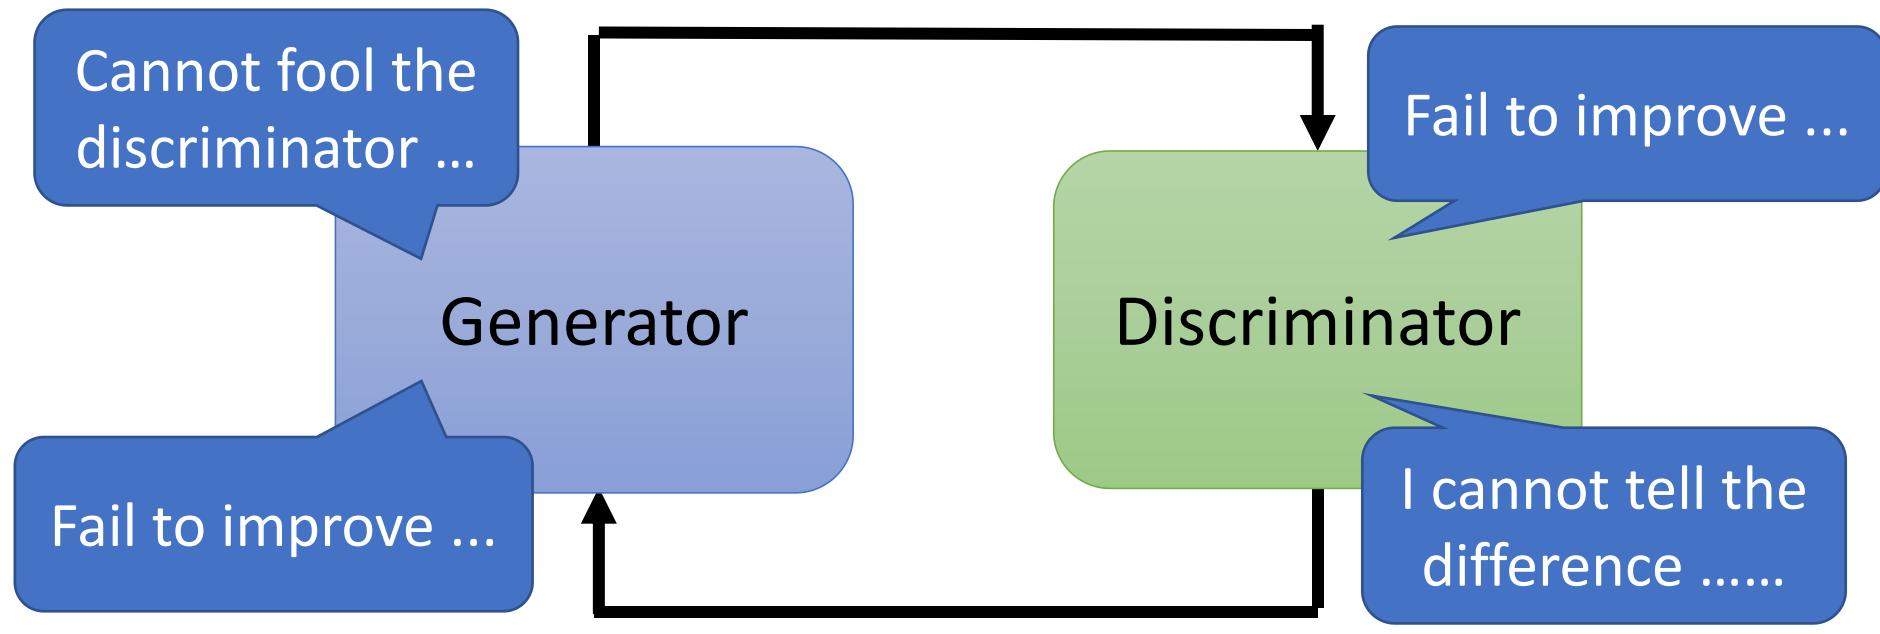


- Spectral Normalization → Keep gradient norm smaller than 1 everywhere
  - <https://arxiv.org/abs/1802.05957>

# GAN is still challenging ...

- Generator and Discriminator needs to match each other

Generate fake images to fool discriminator

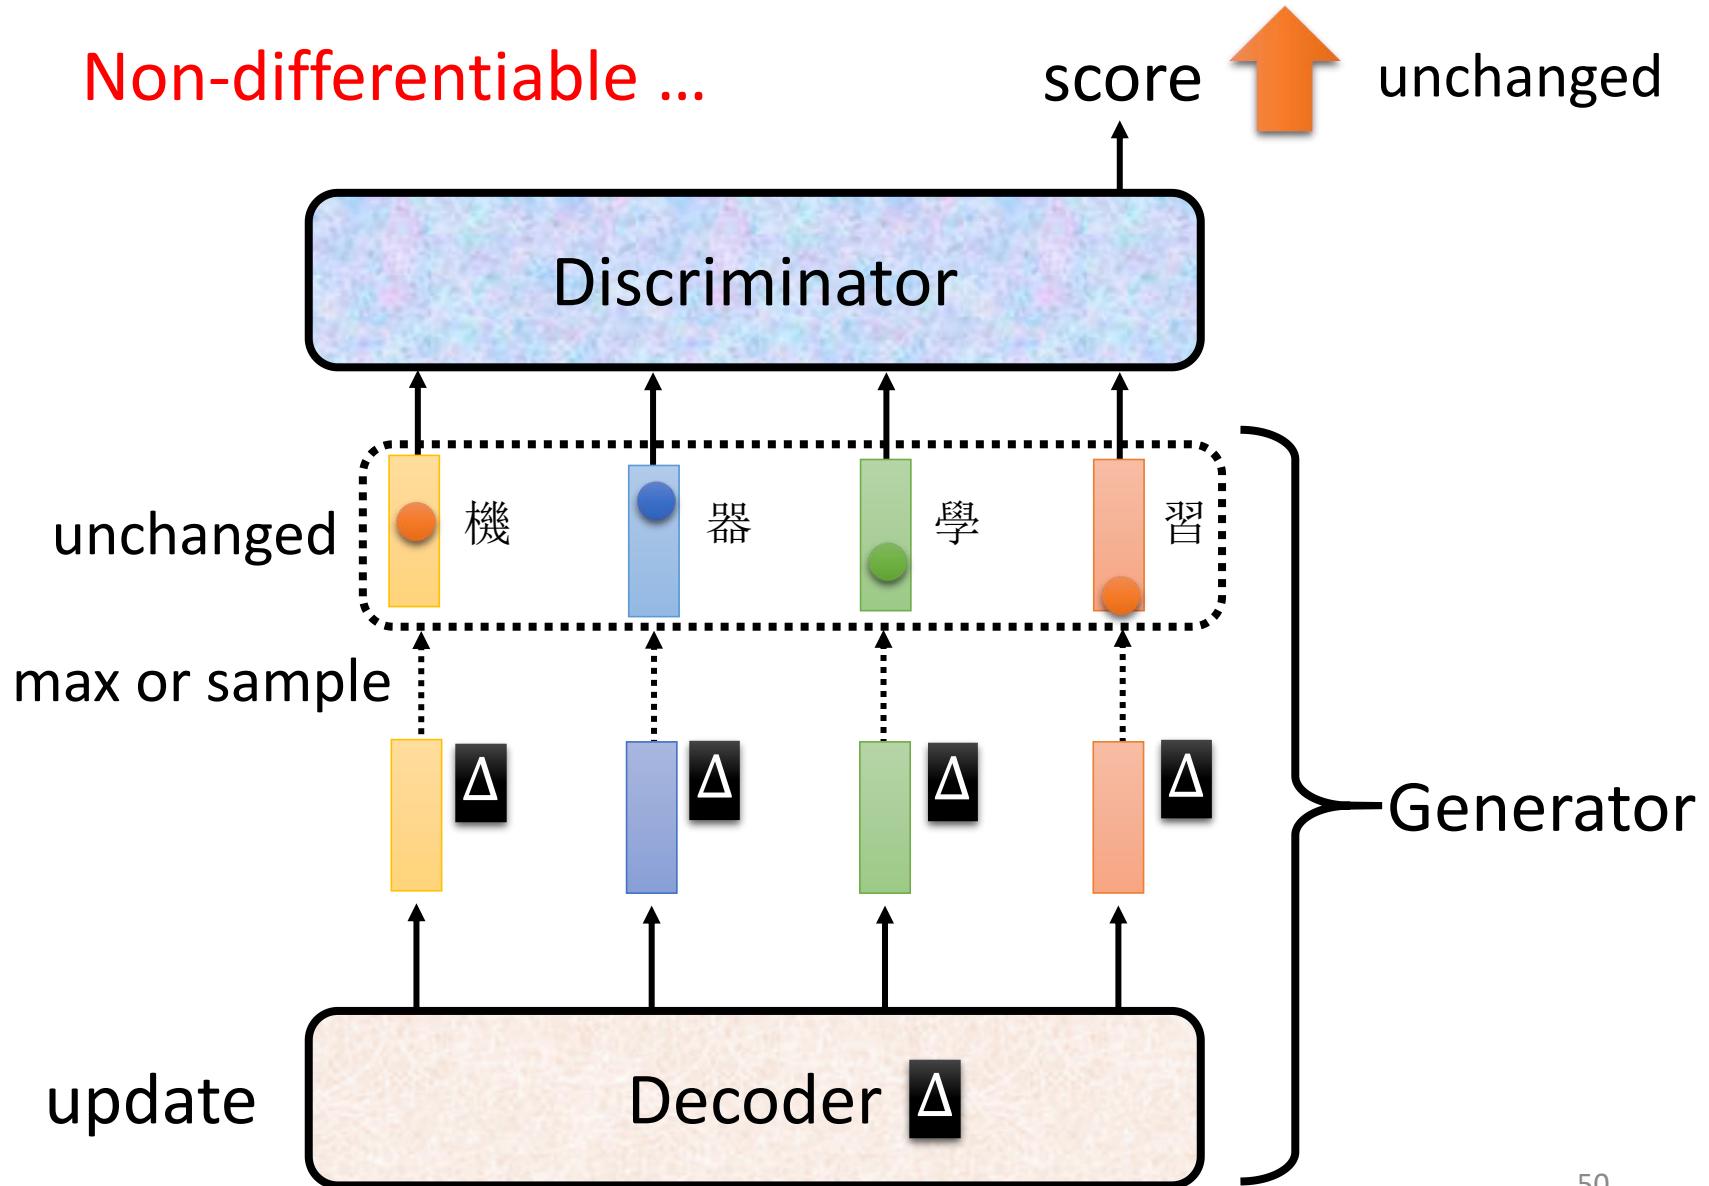


Tell the difference between real and fake

# More Tips

- Tips from Soumith
  - <https://github.com/soumith/ganhacks>
- Tips in DCGAN: Guideline for network architecture design for image generation
  - <https://arxiv.org/abs/1511.06434>
- Improved techniques for training GANs
  - <https://arxiv.org/abs/1606.03498>
- Tips from BigGAN
  - <https://arxiv.org/abs/1809.11096>

# GAN for Sequence Generation

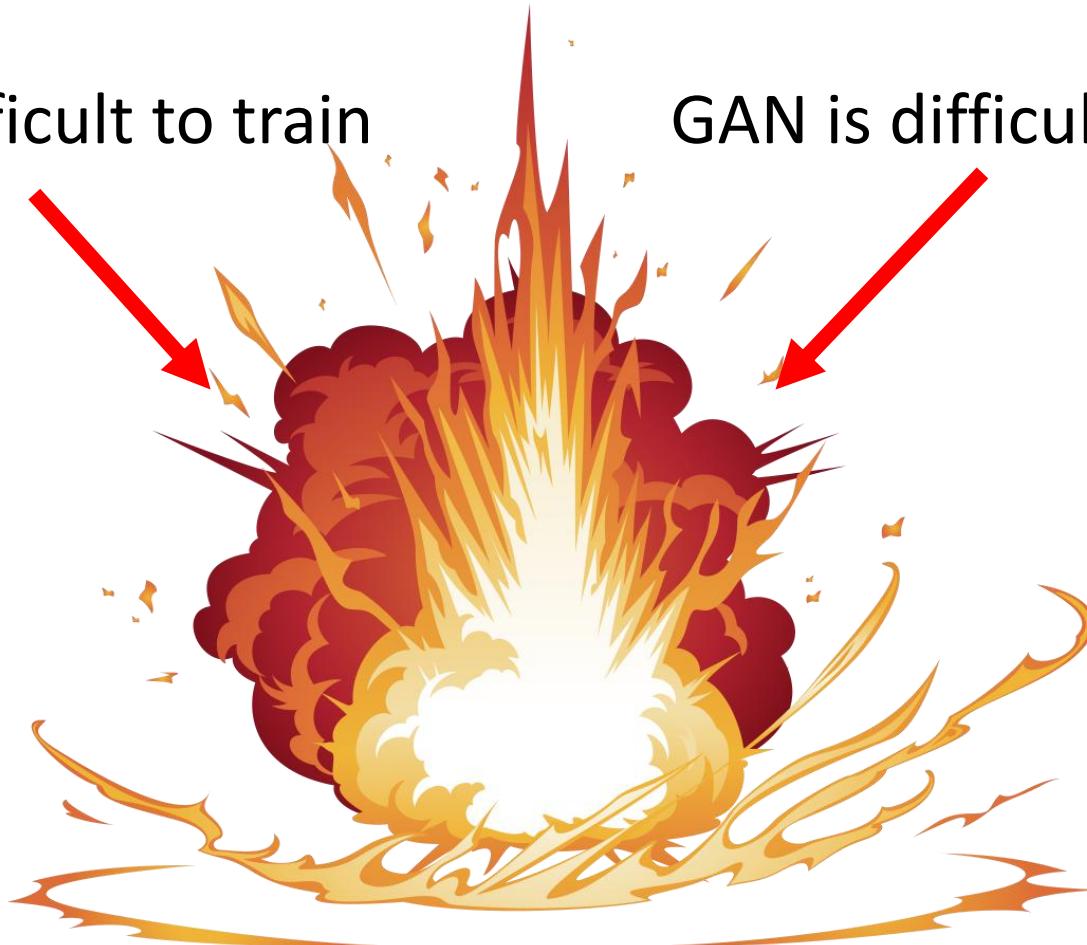


# GAN for Sequence Generation

Reinforcement learning (RL) is involved .....

'RL is difficult to train

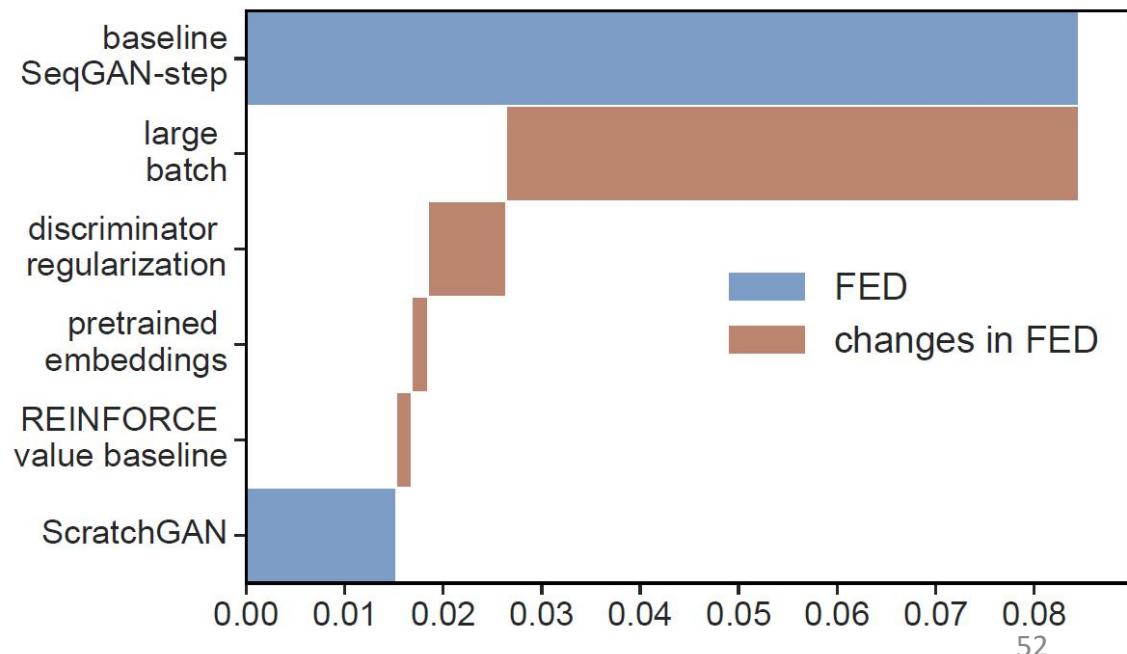
GAN is difficult to train



Sequence Generation GAN (RL+GAN)

# GAN for Sequence Generation

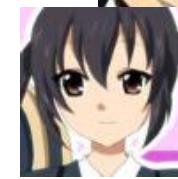
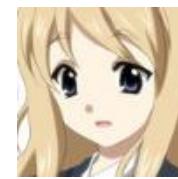
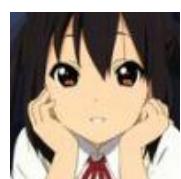
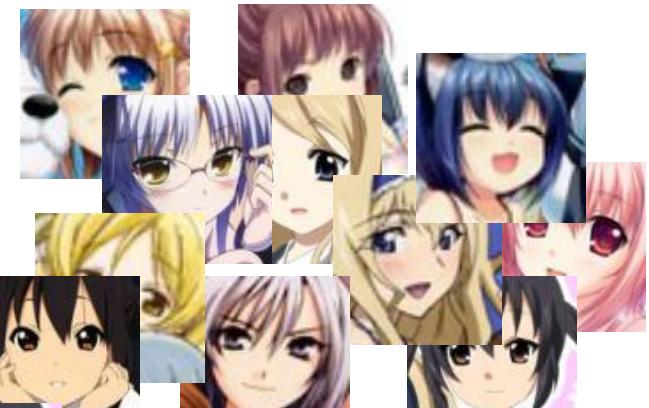
- Usually, the generator are fine-tuned from a model learned by other approaches.
- However, with enough hyperparameter-tuning and tips, ScarchGAN can train from scratch.



Training language  
GANs from Scratch

[https://arxiv.org/abs/  
1905.09922](https://arxiv.org/abs/1905.09922)

# Possible Solution?



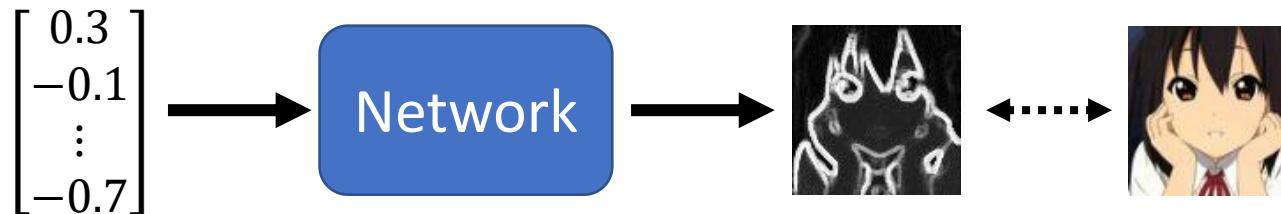
$$\begin{bmatrix} 0.3 \\ -0.1 \\ \vdots \\ -0.7 \end{bmatrix}$$

$$\begin{bmatrix} 0.1 \\ -0.1 \\ \vdots \\ 0.7 \end{bmatrix}$$

$$\begin{bmatrix} -0.3 \\ 0.1 \\ \vdots \\ 0.9 \end{bmatrix}$$

$$\begin{bmatrix} 0.7 \\ 0.1 \\ \vdots \\ -0.9 \end{bmatrix}$$

$$\begin{bmatrix} -0.1 \\ 0.8 \\ \vdots \\ 0.8 \end{bmatrix}$$

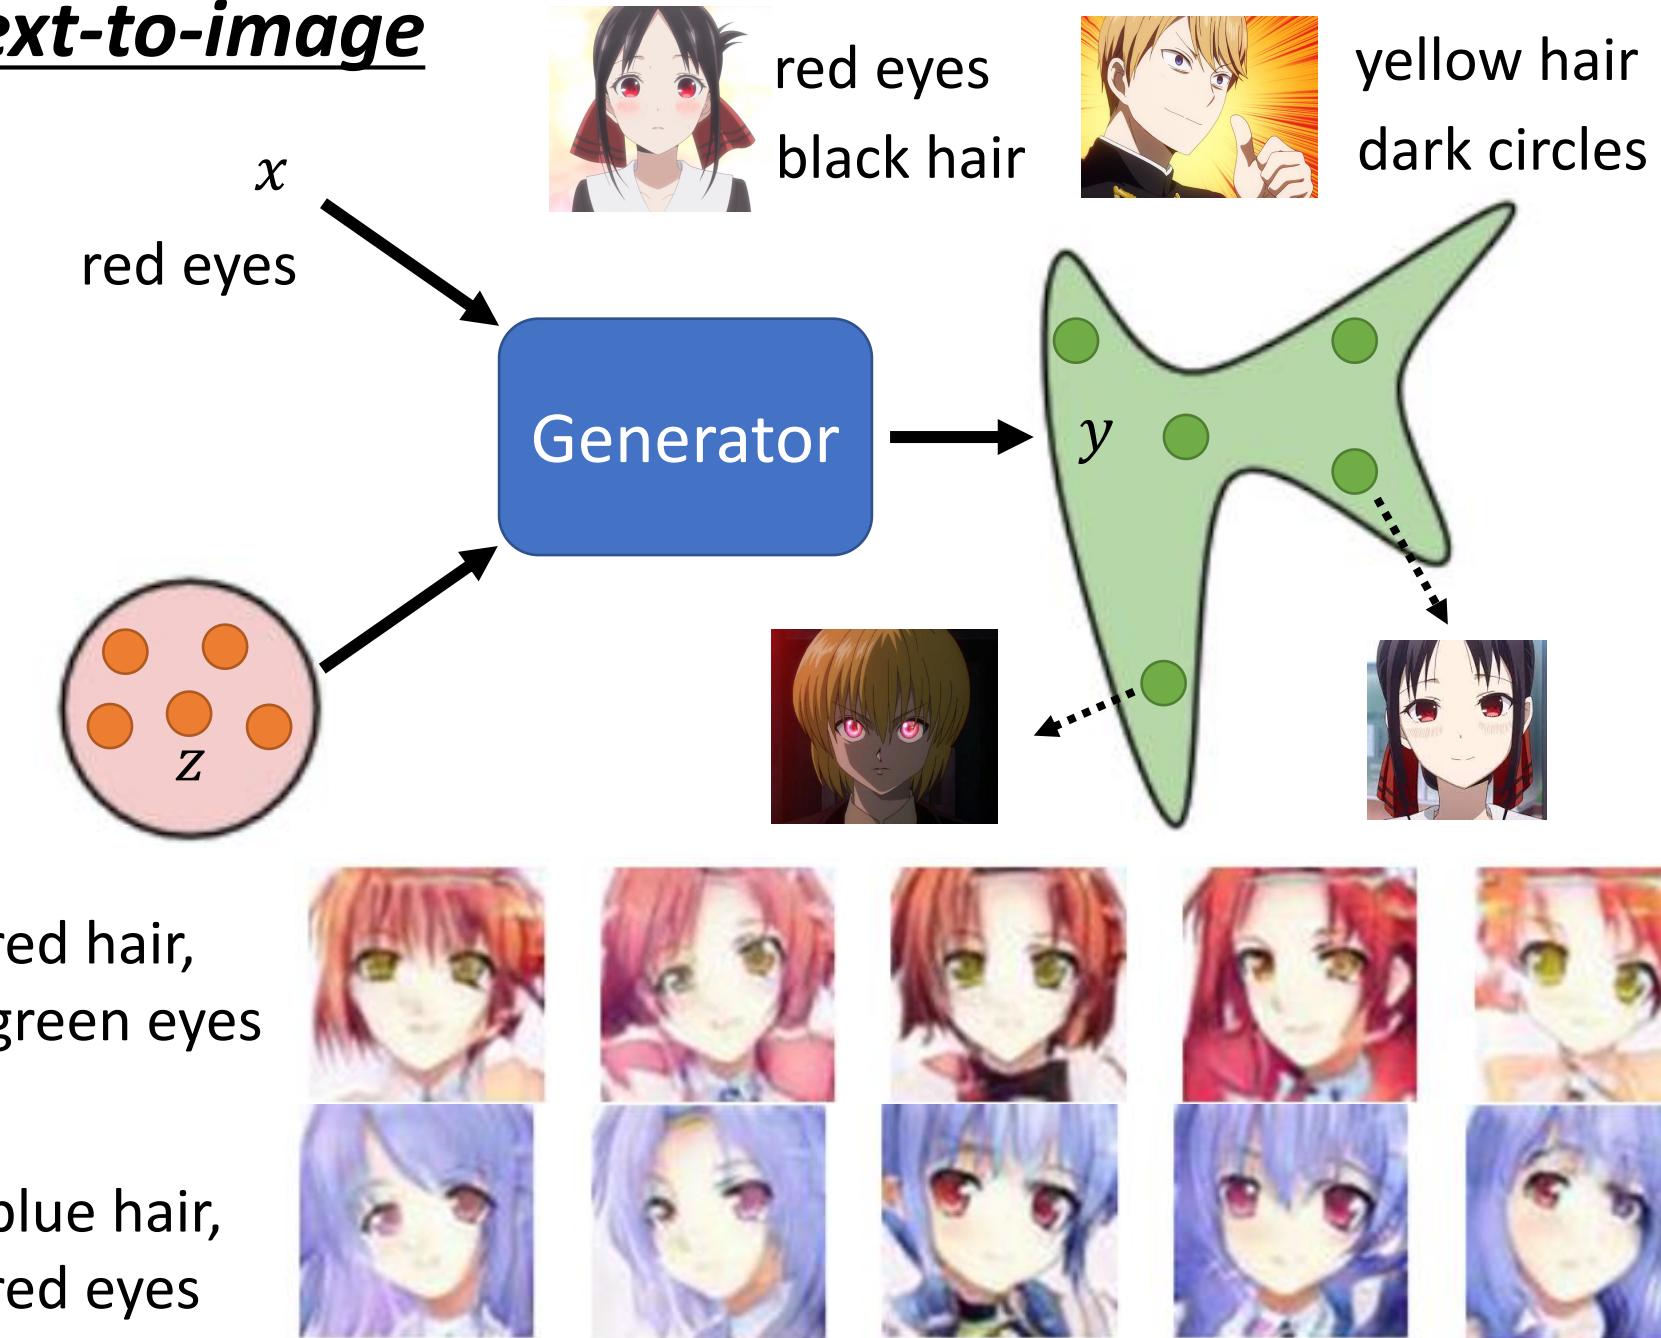


Using typical learning approaches?

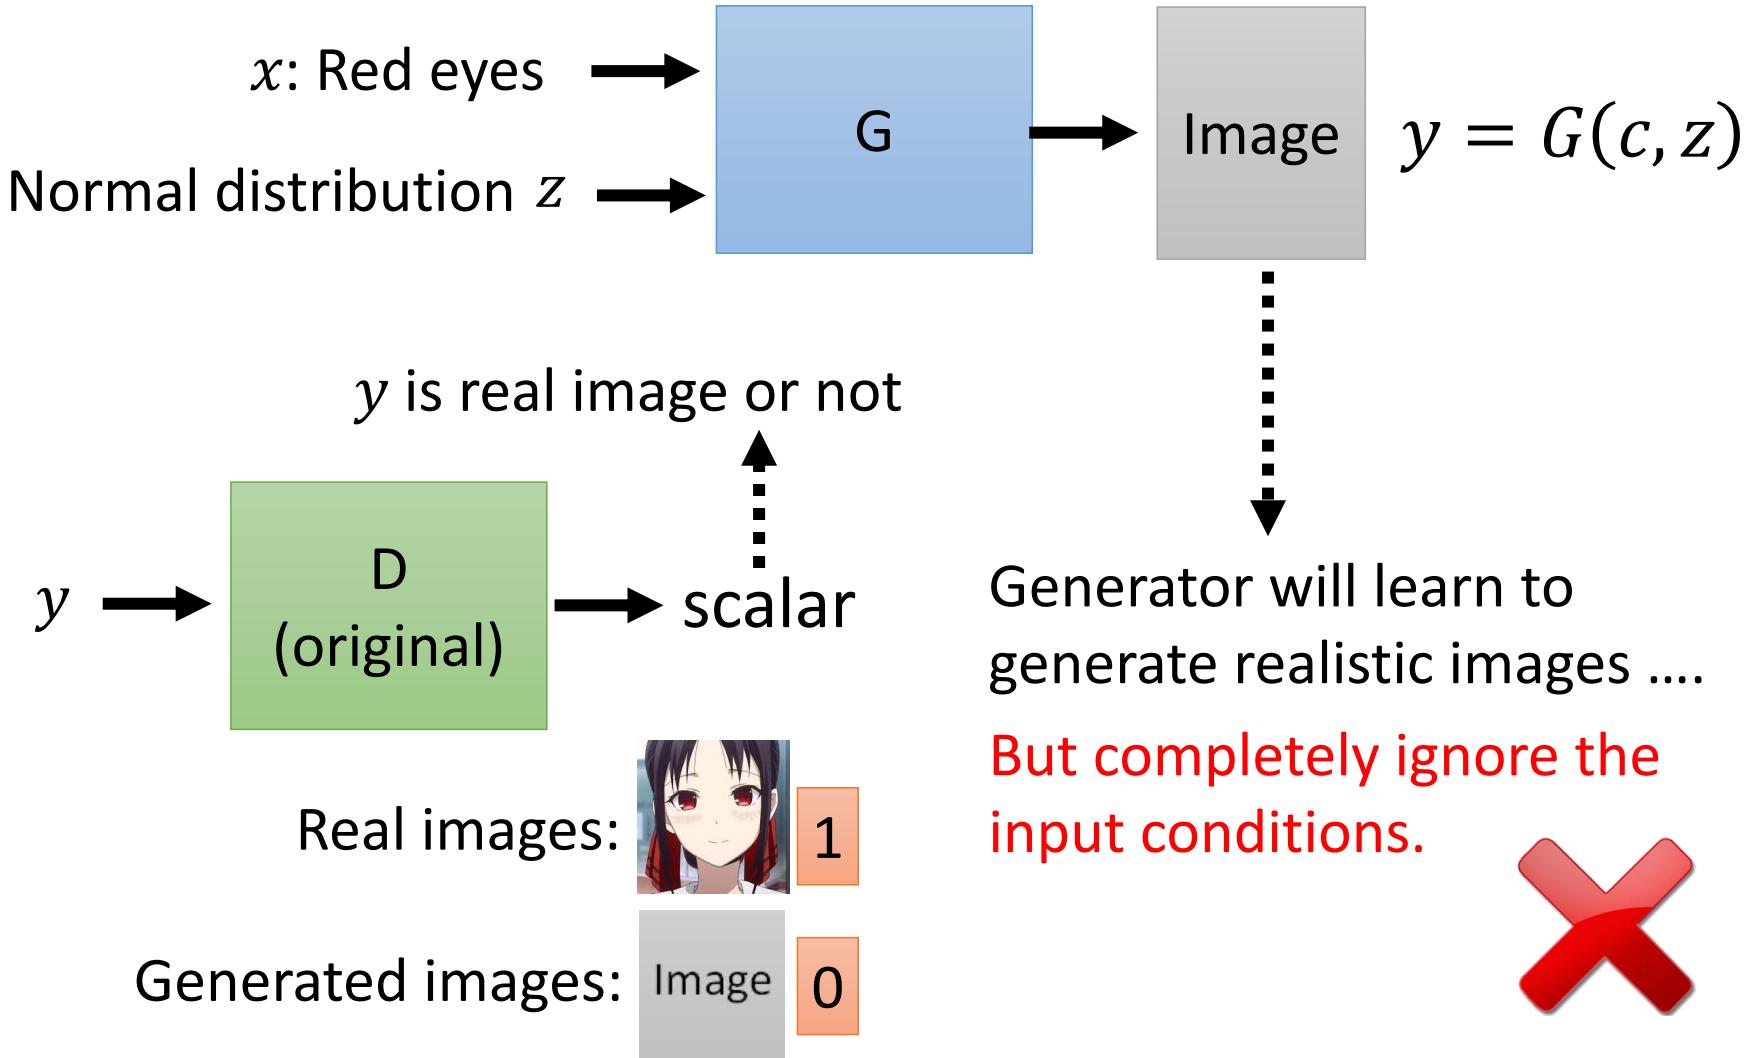
Generative Latent Optimization (GLO), <https://arxiv.org/abs/1707.05776>  
Gradient Origin Networks, <https://arxiv.org/abs/2007.02798>

# Conditional Generation

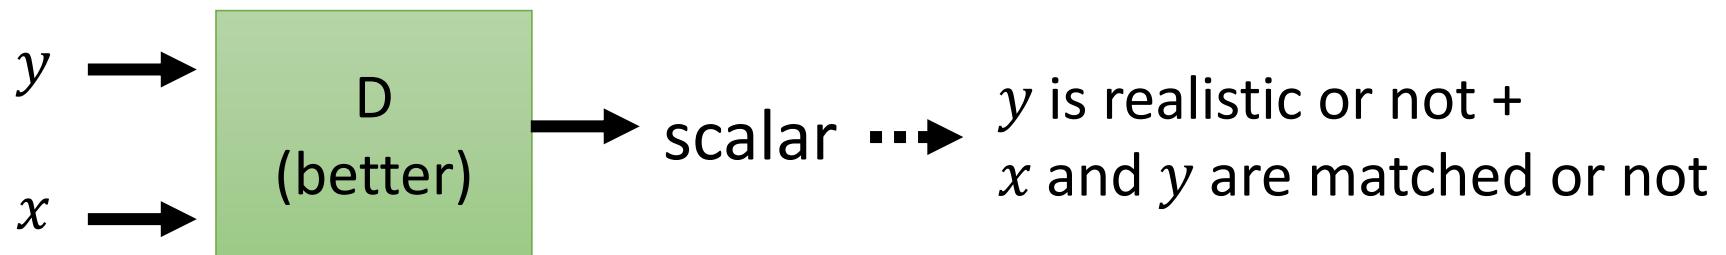
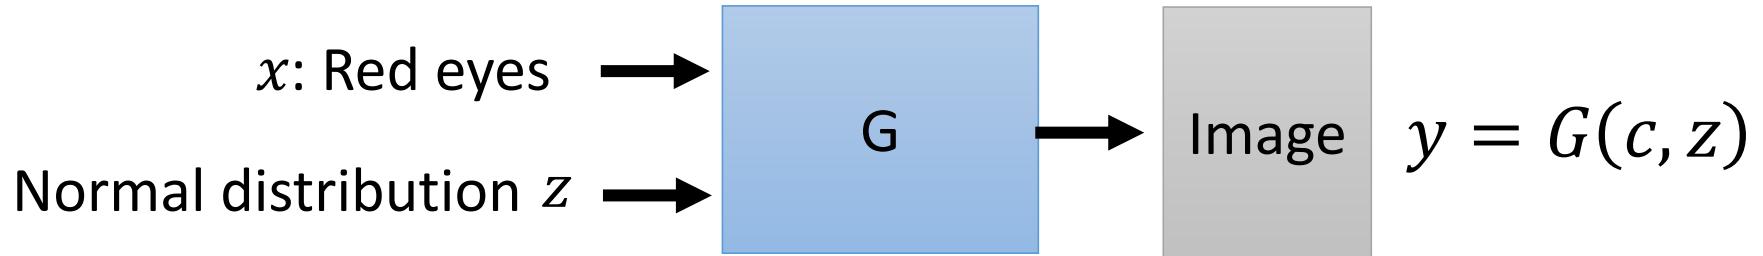
# Text-to-image



# Conditional GAN



# Conditional GAN



True text-image pairs: (red eyes,

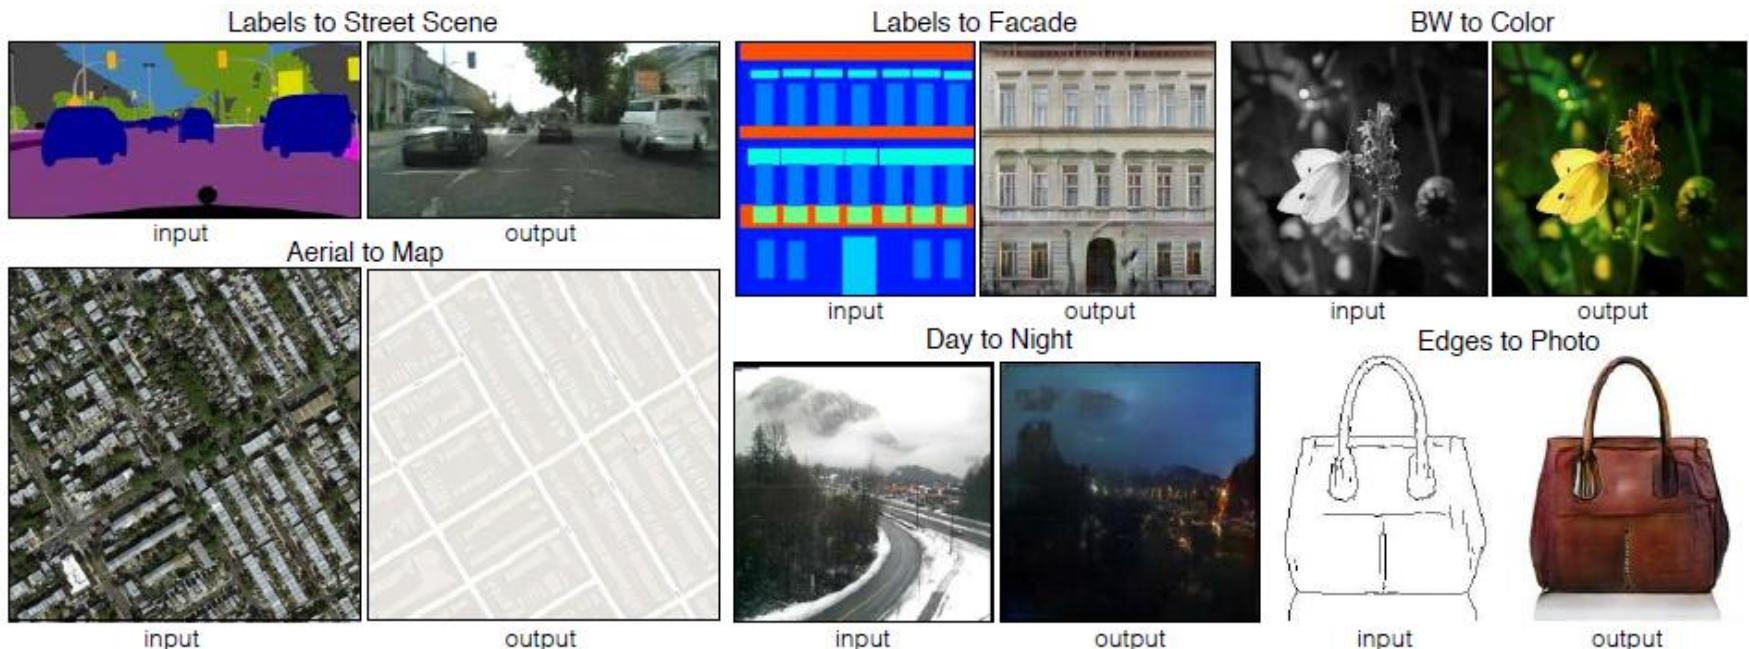
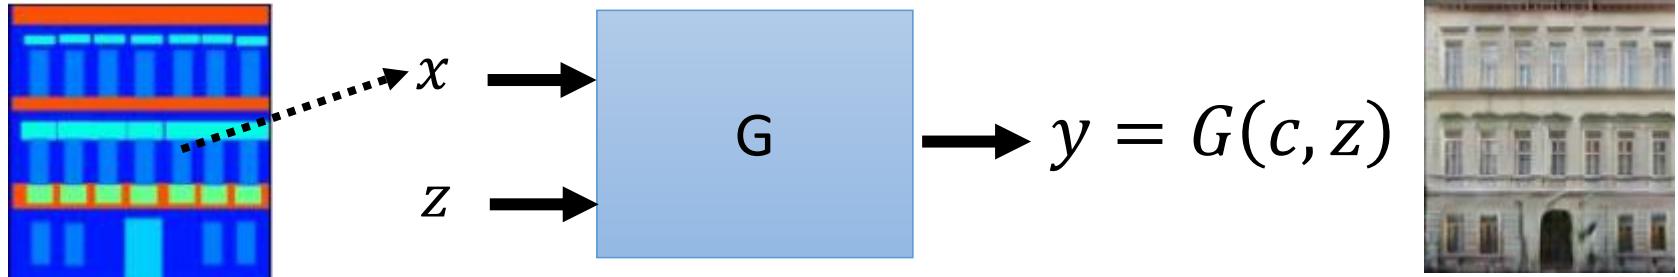


) 1

(red eyes, ) 0

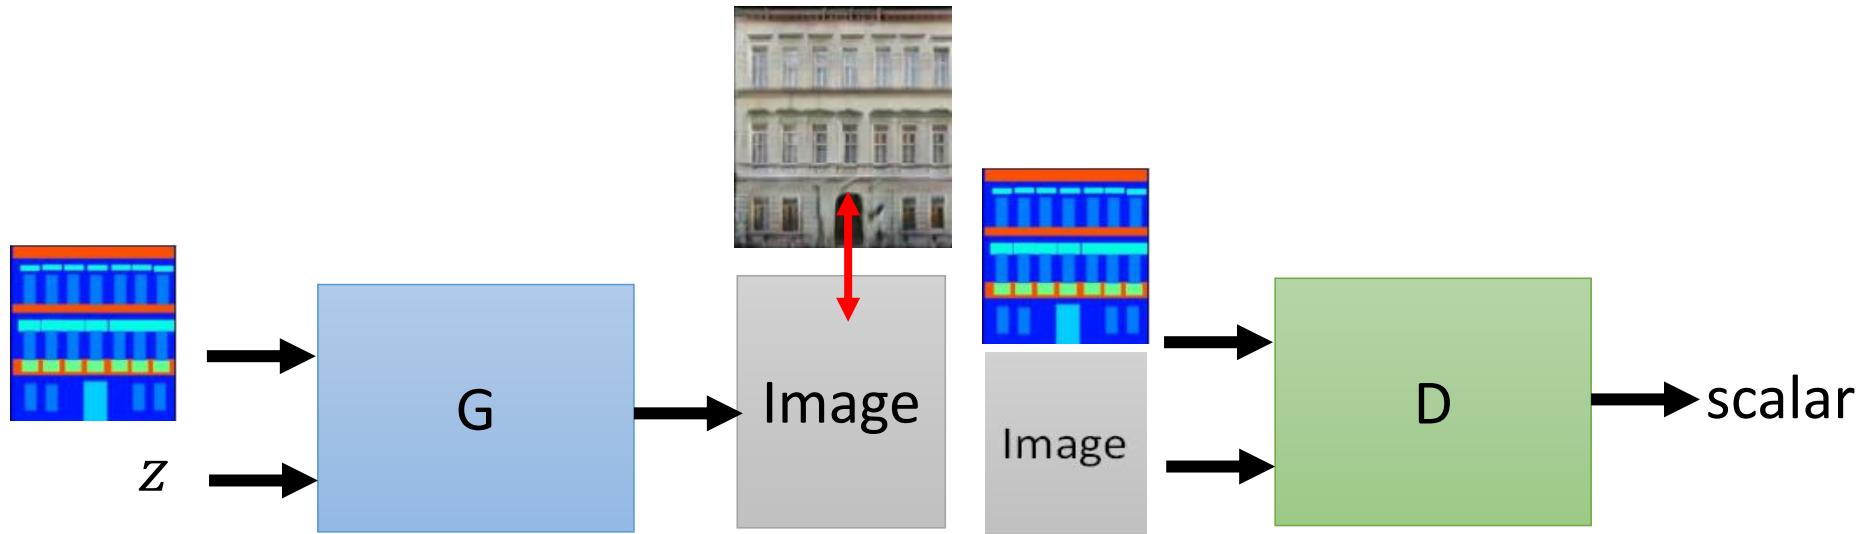
(red eyes, ) 0

# Conditional GAN

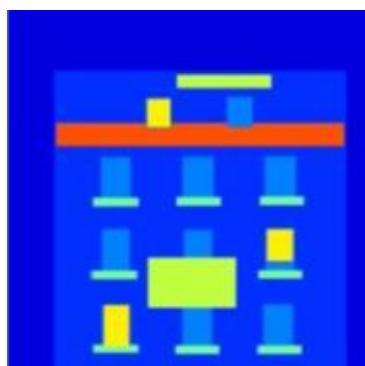


**Image translation, or pix2pix**

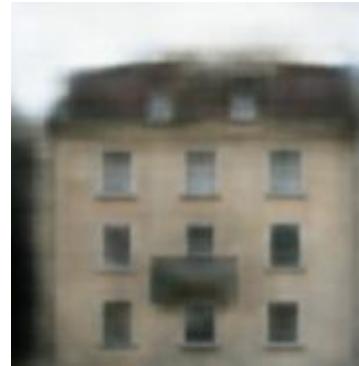
# Conditional GAN



Testing:



input



supervised

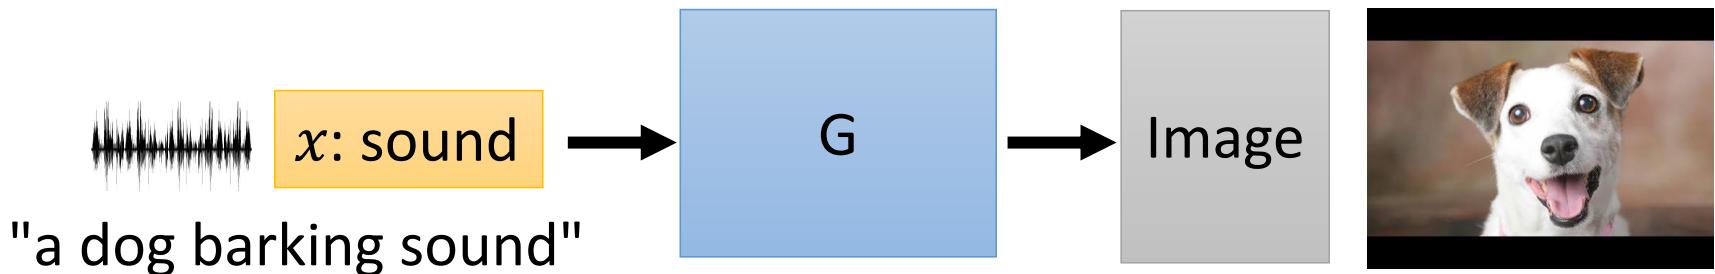


GAN



GAN + supervised

# Conditional GAN



*Training Data  
Collection*



video

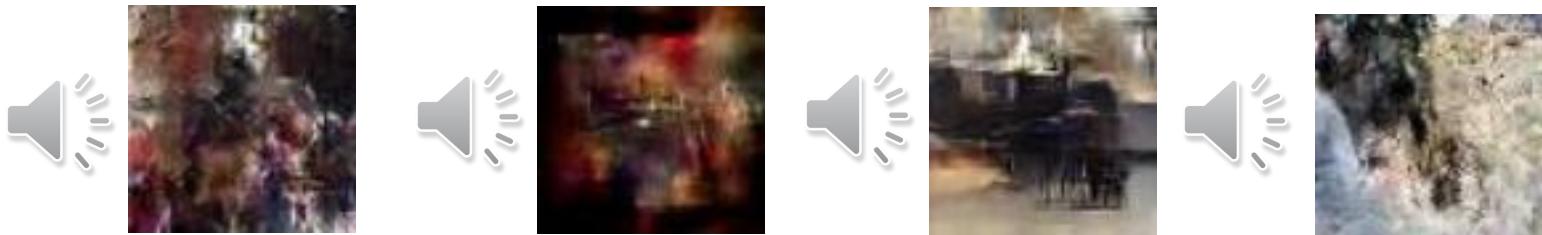
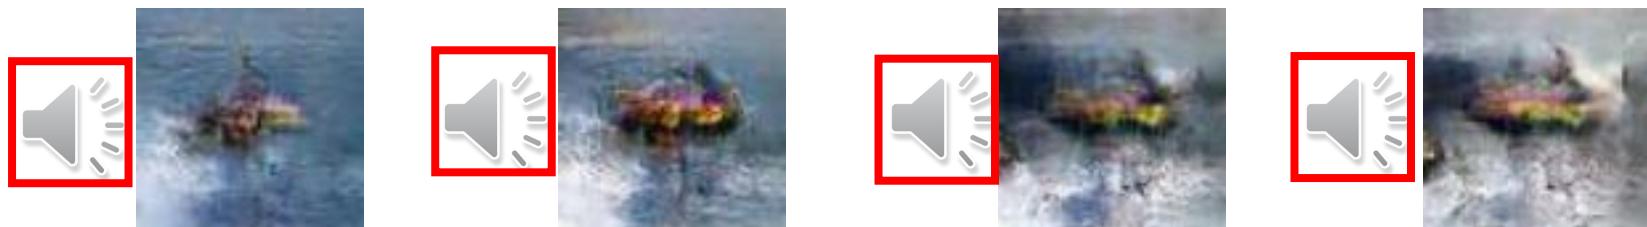


# Conditional GAN

The images are generated by  
Chia-Hung Wan and Shun-  
Peng Chiang  
[https://chiahung1483.github.io/  
audio\\_to\\_scene/index.html](https://chiahung1483.github.io/audio_to_scene/index.html)

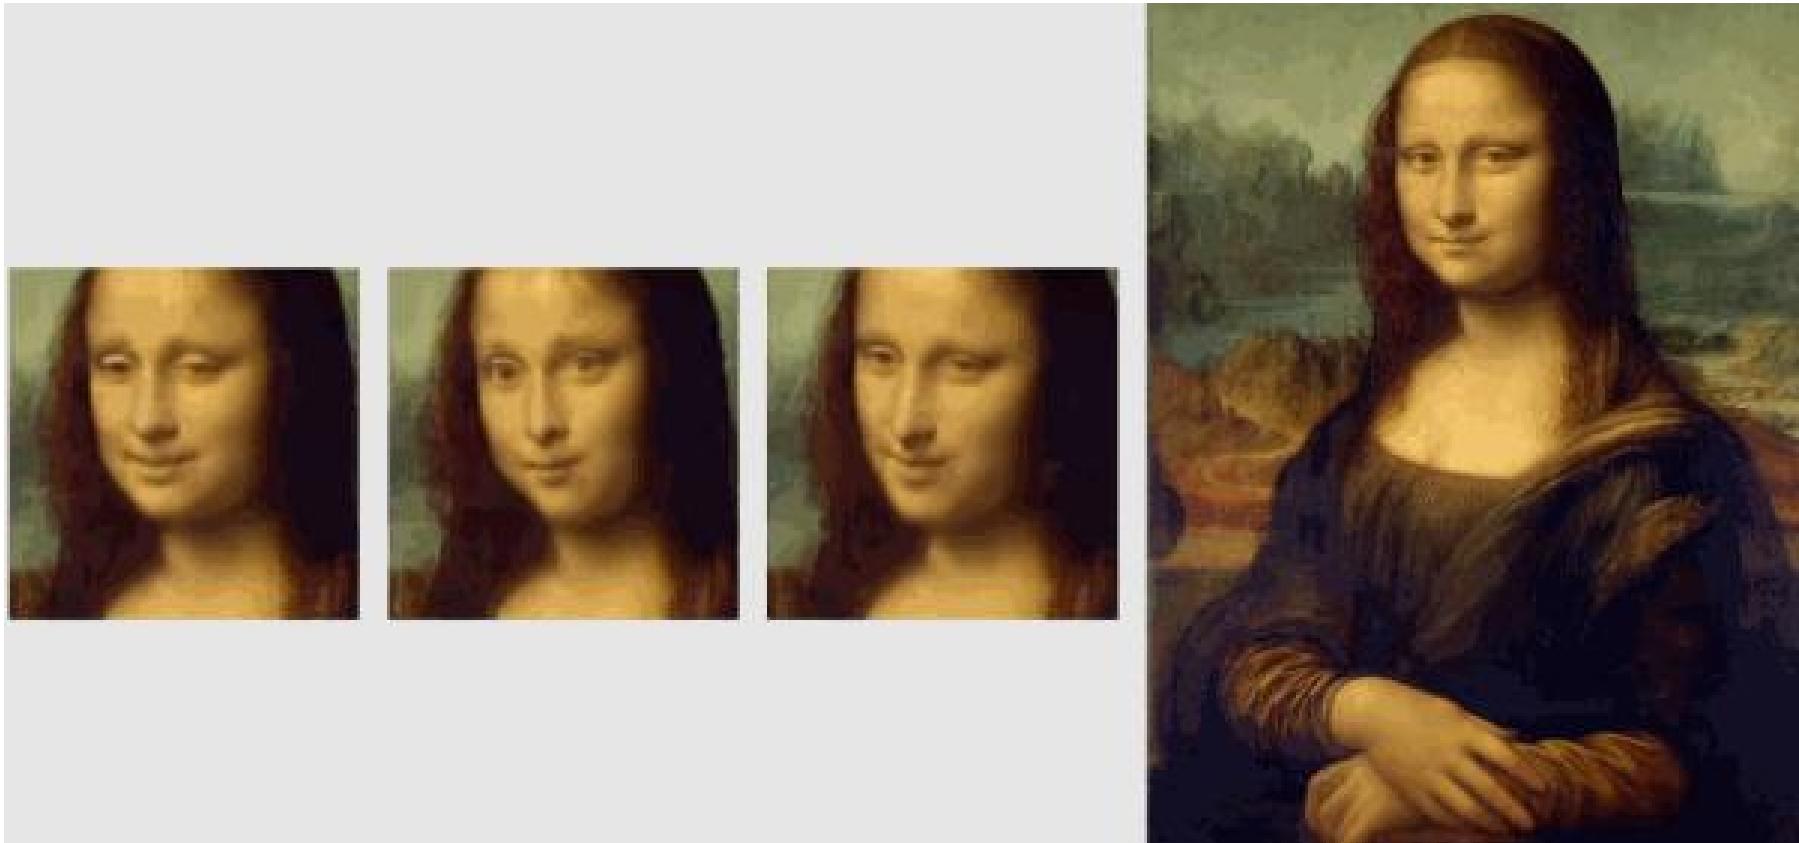
- Sound-to-image

Louder



# Conditional GAN

## Talking Head Generation



<https://arxiv.org/abs/1905.08233>

# Conditional GAN

Multi-label Image Classifier = Conditional Generator

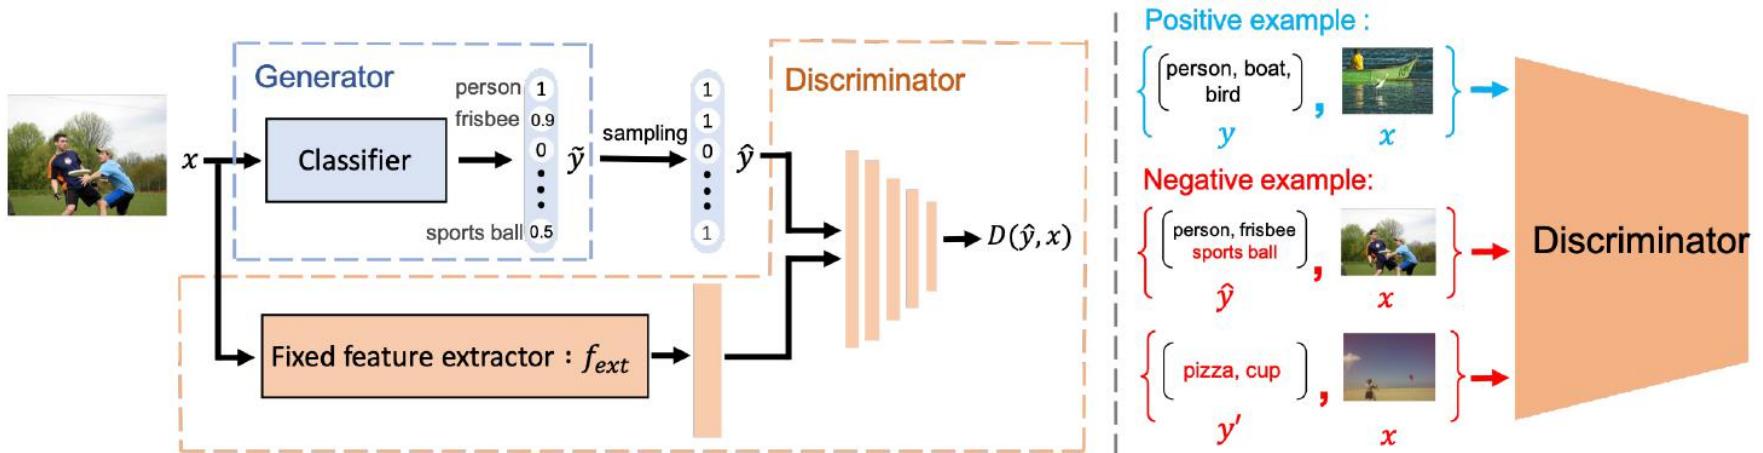


Input condition

person, sports ball,  
baseball bat, baseball glove



Generated output



Positive example :

$\left\{ \begin{array}{l} \text{person, boat,} \\ \text{bird} \end{array} \right. , \text{ } \left. \begin{array}{l} \text{y} \\ \text{x} \end{array} \right\}$

Negative example:

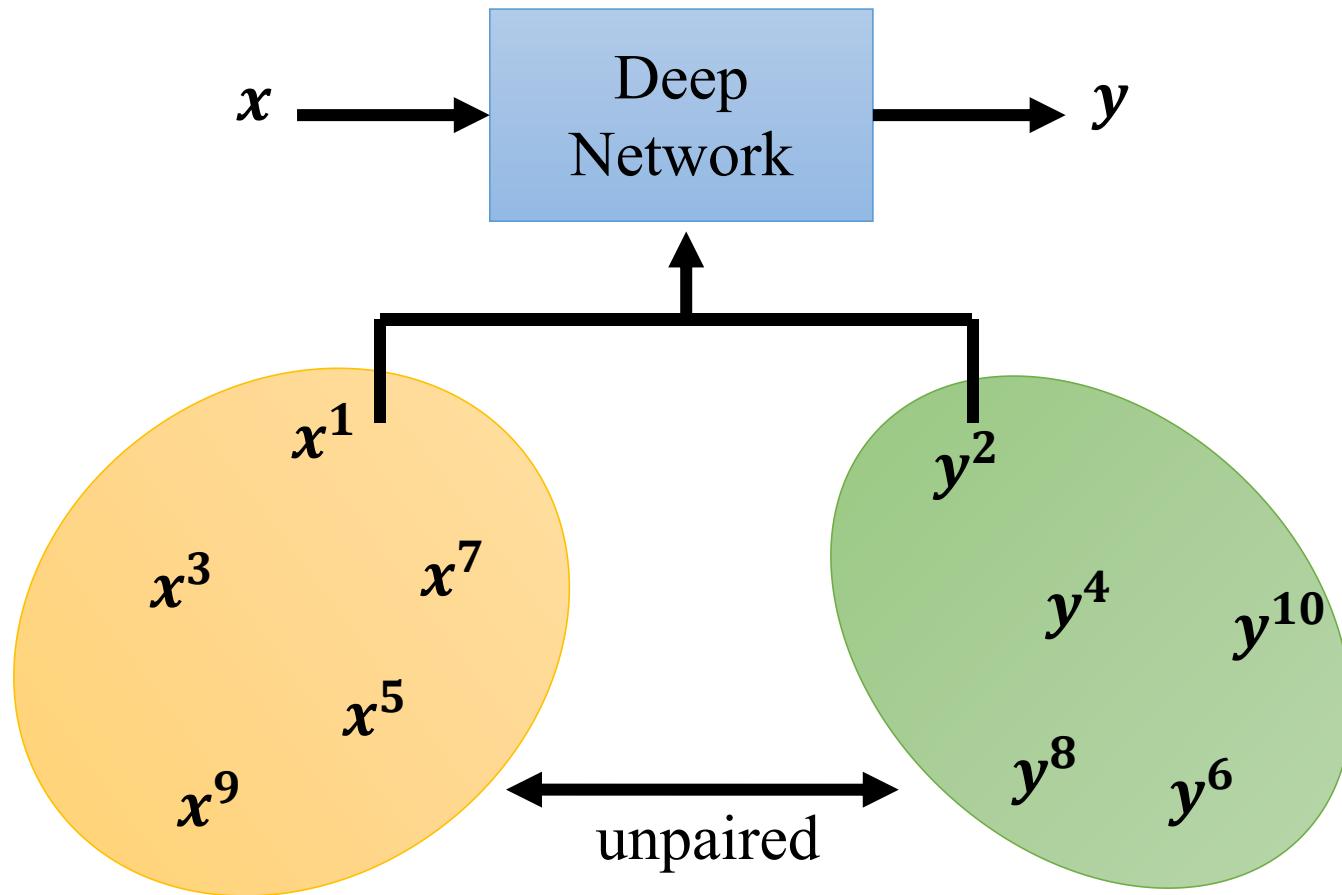
$\left\{ \begin{array}{l} \text{person, frisbee,} \\ \text{sports ball} \end{array} \right. , \text{ } \left. \begin{array}{l} \text{y} \\ \text{x} \end{array} \right\}$

$\left\{ \begin{array}{l} \text{pizza, cup} \end{array} \right. , \text{ } \left. \begin{array}{l} \text{y'} \\ \text{x} \end{array} \right\}$

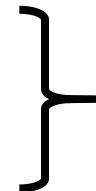
Discriminator

# Learning from Unpaired Data

# Learning from Unpaired Data

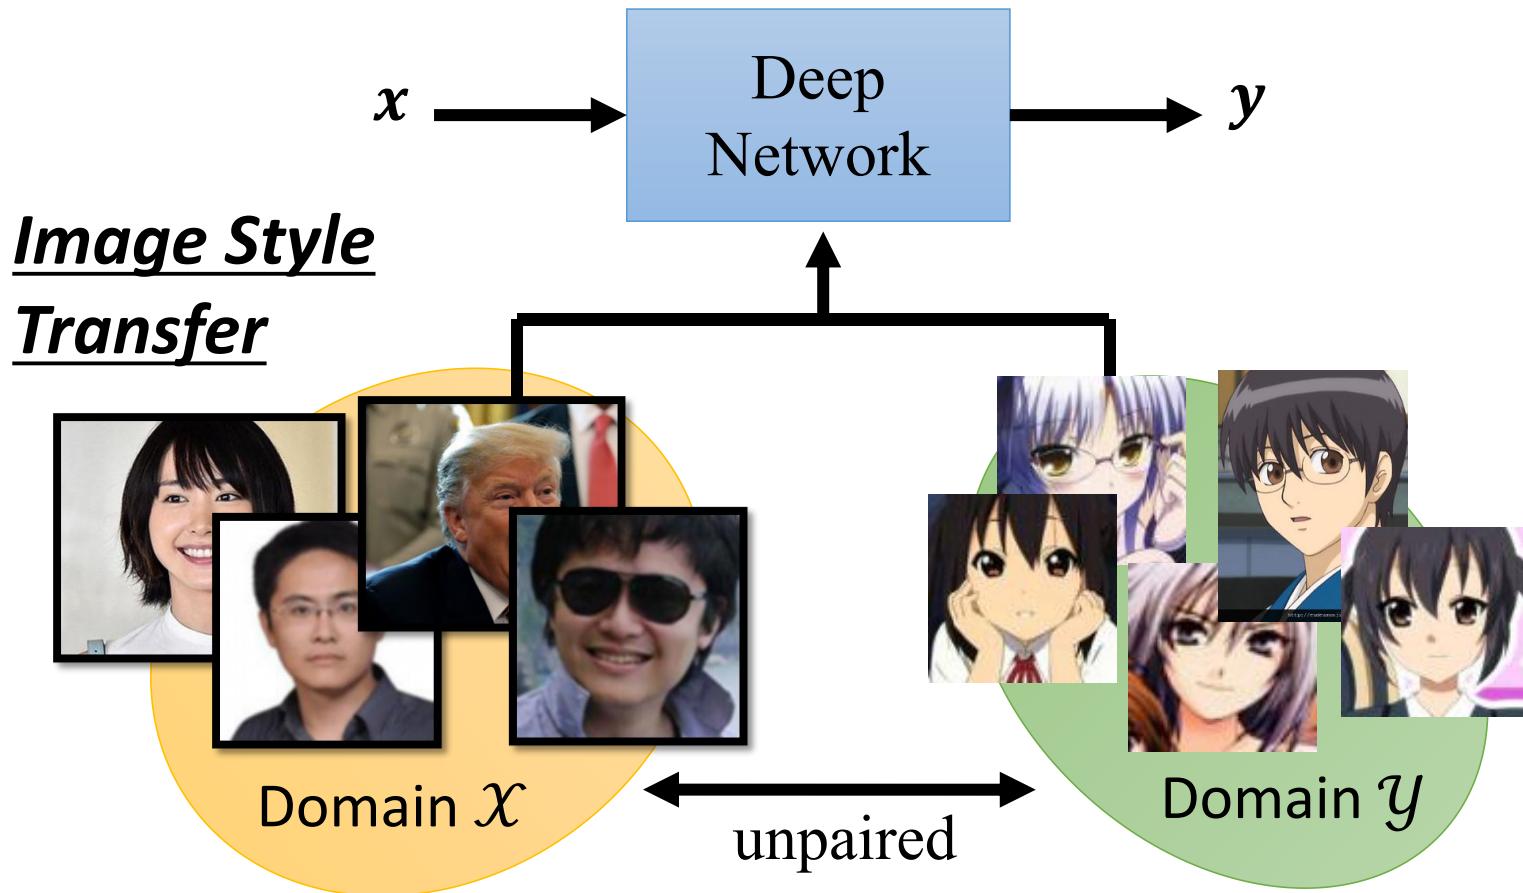


HW3: pseudo labeling  
HW5: back translation



Still need **some**  
paired data

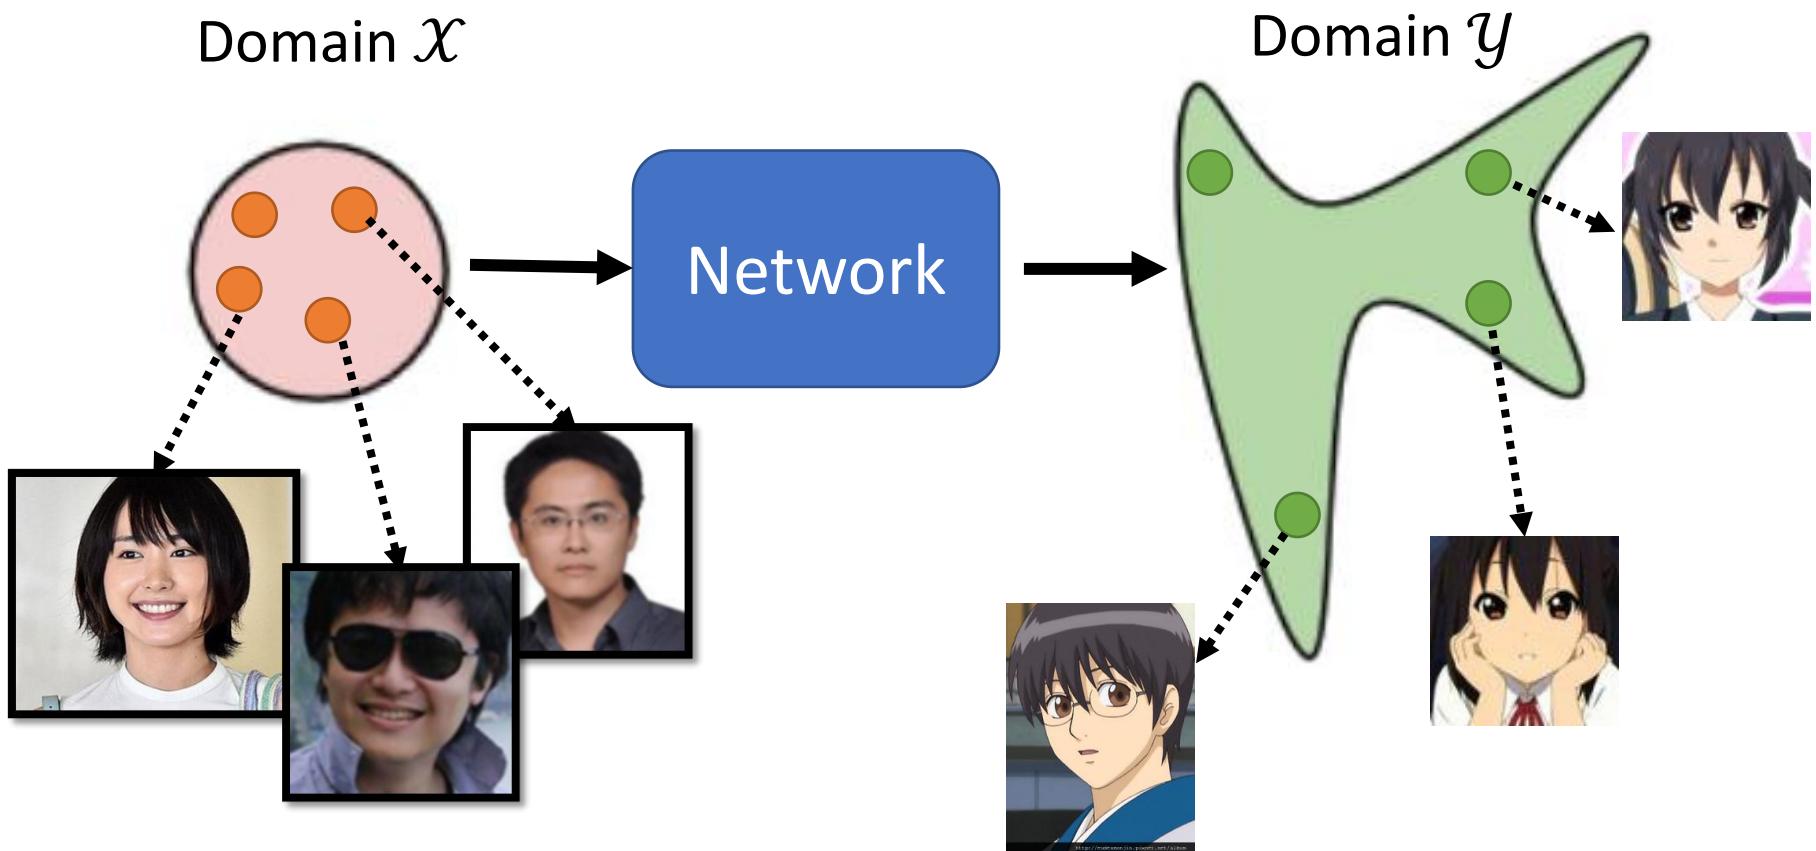
# Learning from Unpaired Data



Can we learn the mapping without any paired data?

**Unsupervised Conditional Generation**

# Learning from Unpaired Data



# Cycle GAN

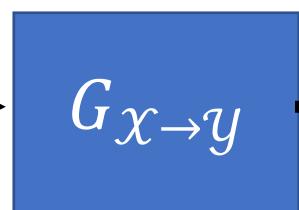
Domain  $\mathcal{X}$



Domain  $\mathcal{Y}$



Domain  $\mathcal{X}$



Become similar  
to domain  $\mathcal{Y}$



Domain  $\mathcal{Y}$



scalar



Input image  
belongs to  
domain  $\mathcal{Y}$  or not

# Cycle GAN

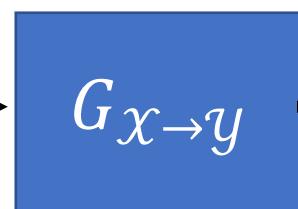
Domain  $\mathcal{X}$



Domain  $\mathcal{Y}$



Domain  $\mathcal{X}$

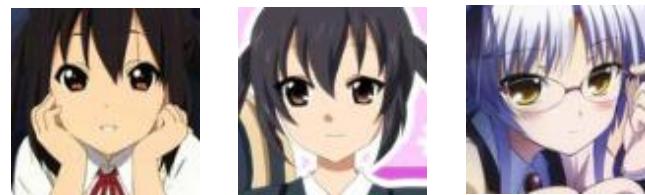


ignore input

Become similar  
to domain  $\mathcal{Y}$

$$D_y$$

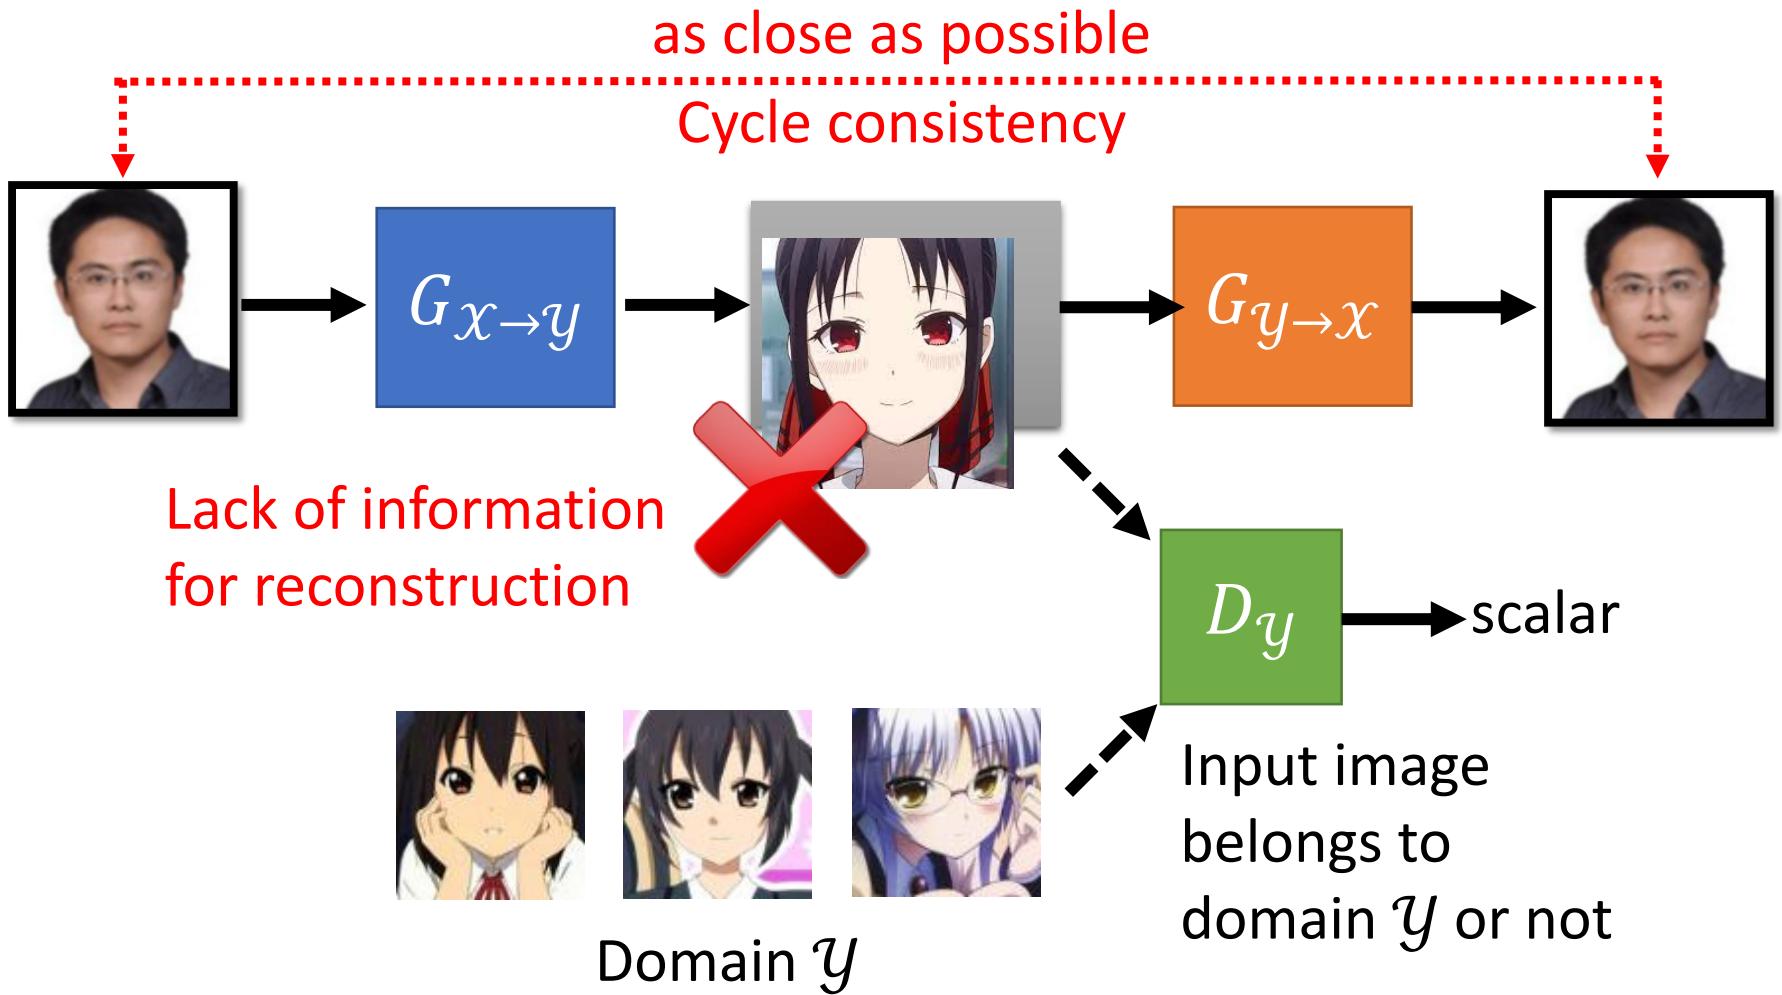
scalar



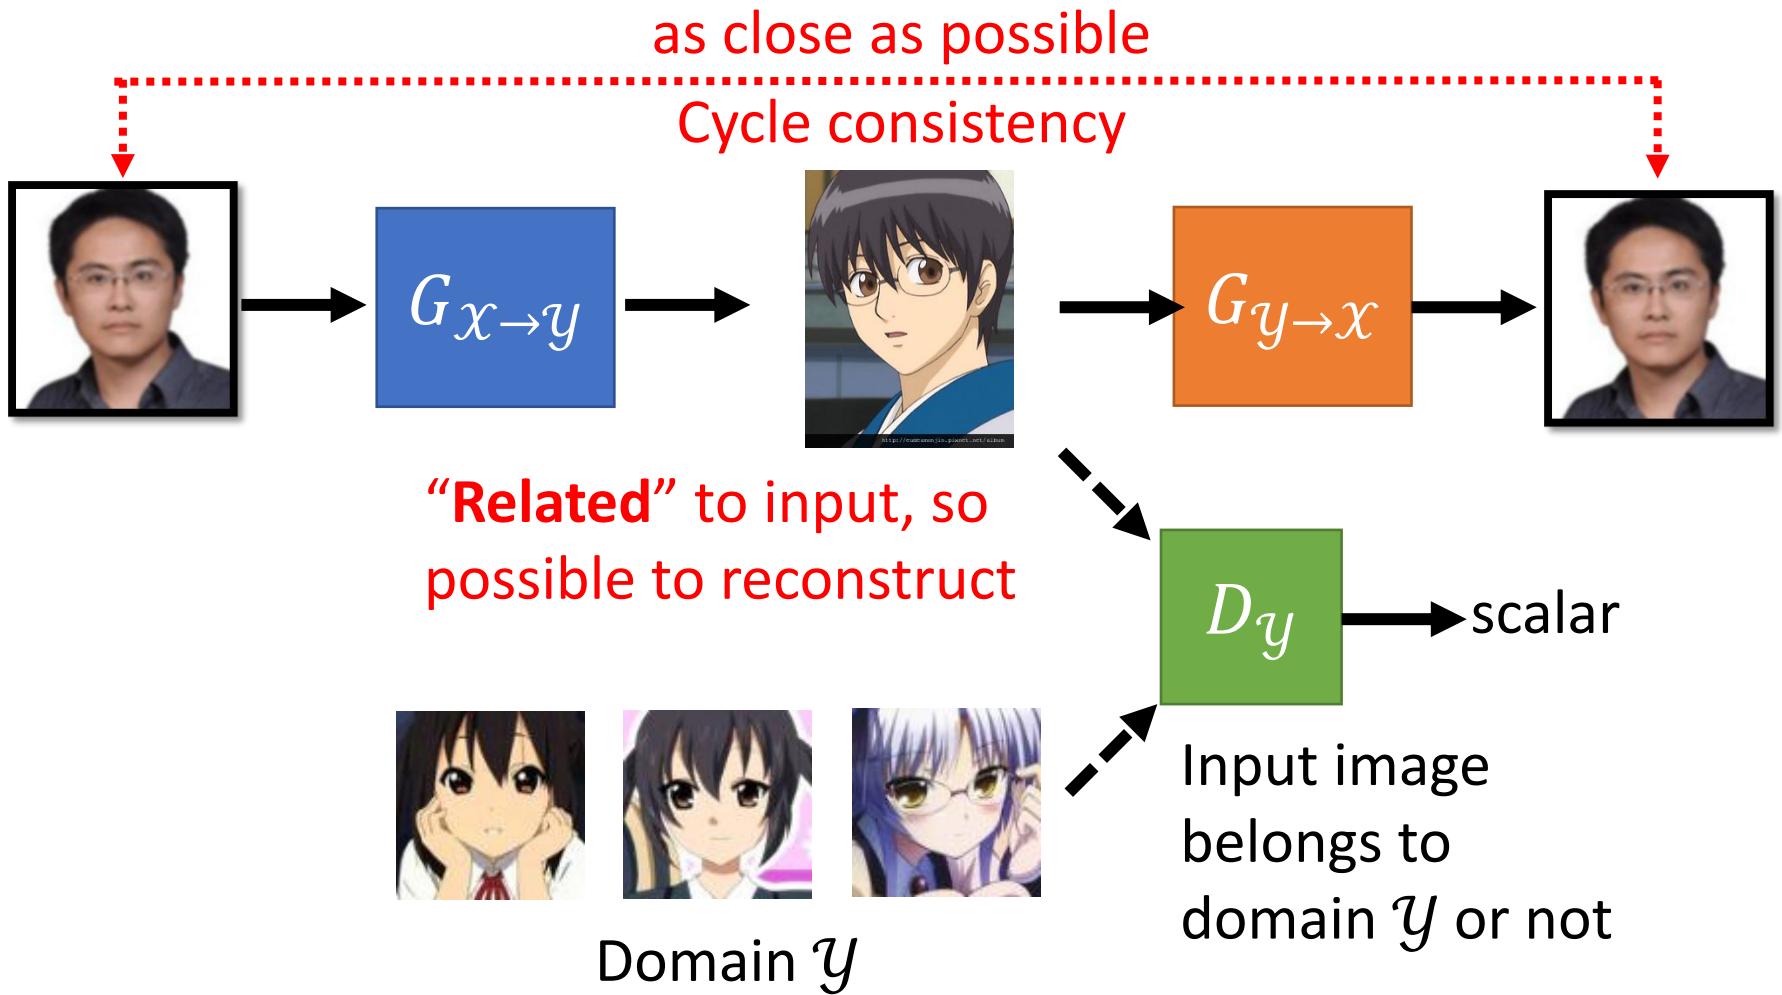
Domain  $\mathcal{Y}$

Input image  
belongs to  
domain  $\mathcal{Y}$  or not

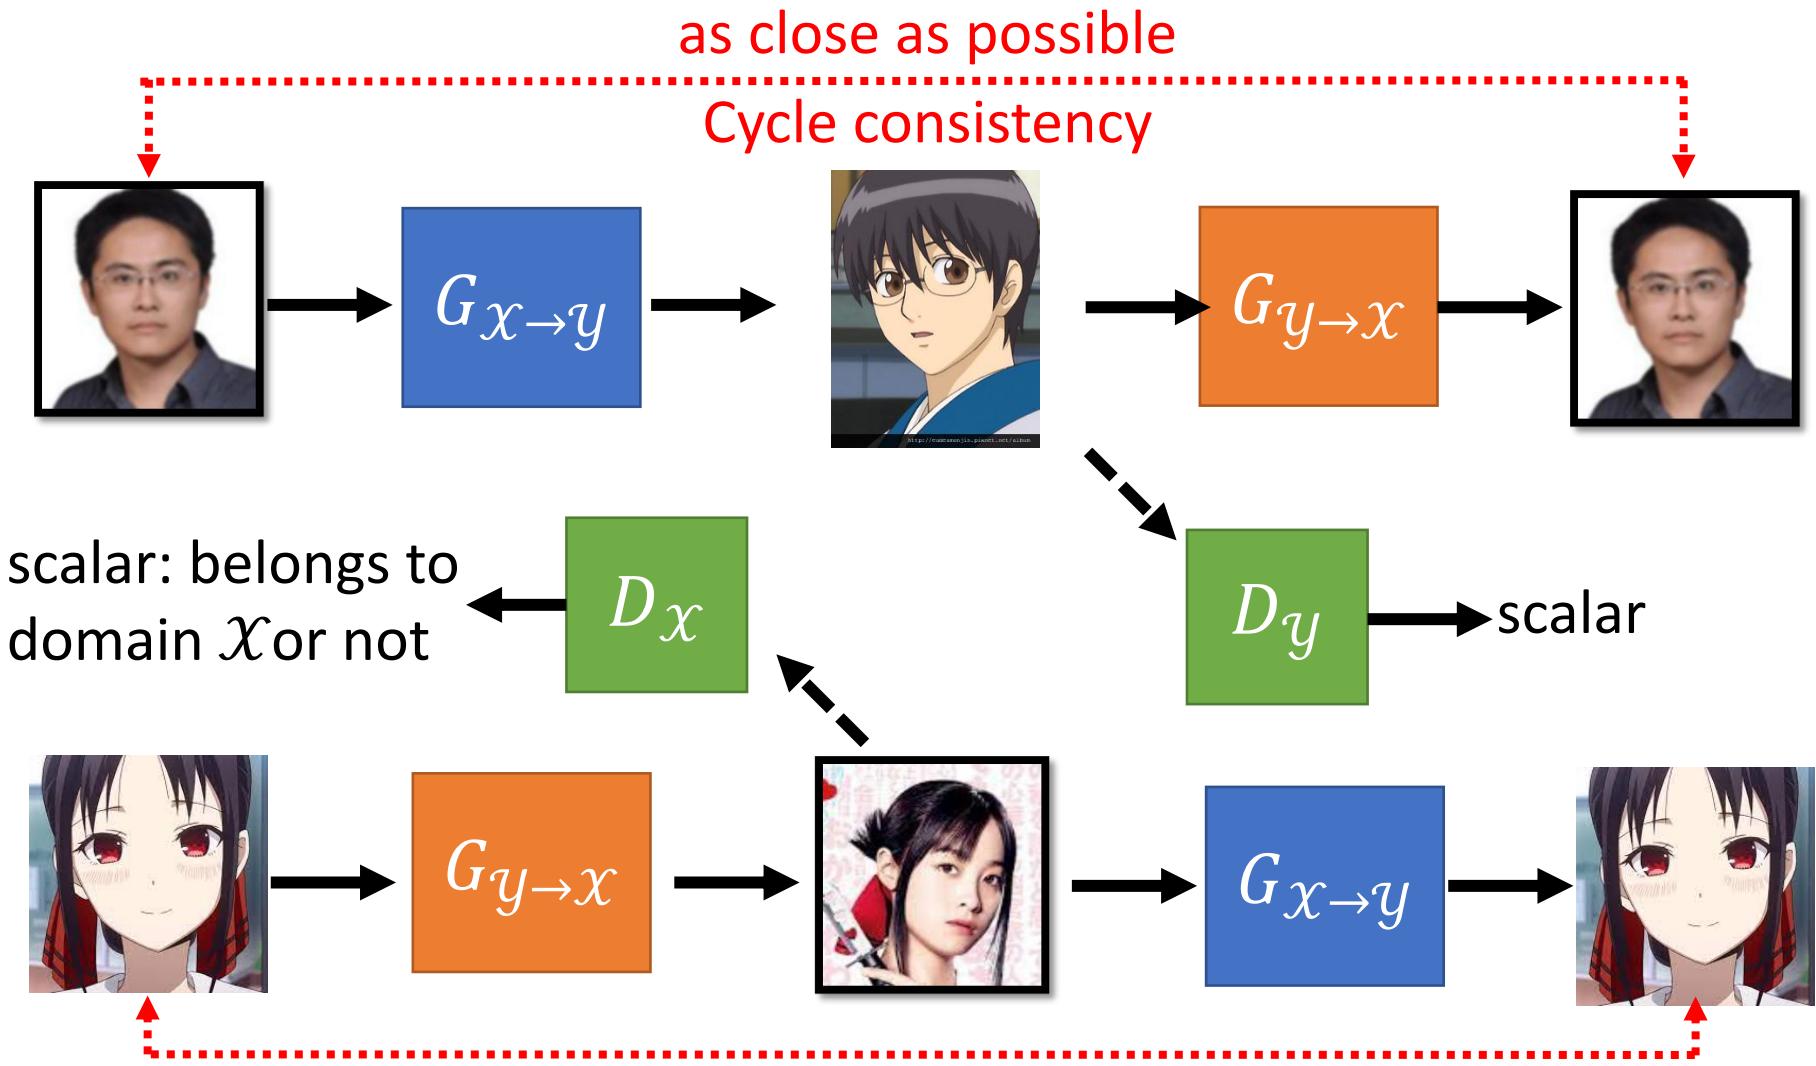
# Cycle GAN



# Cycle GAN

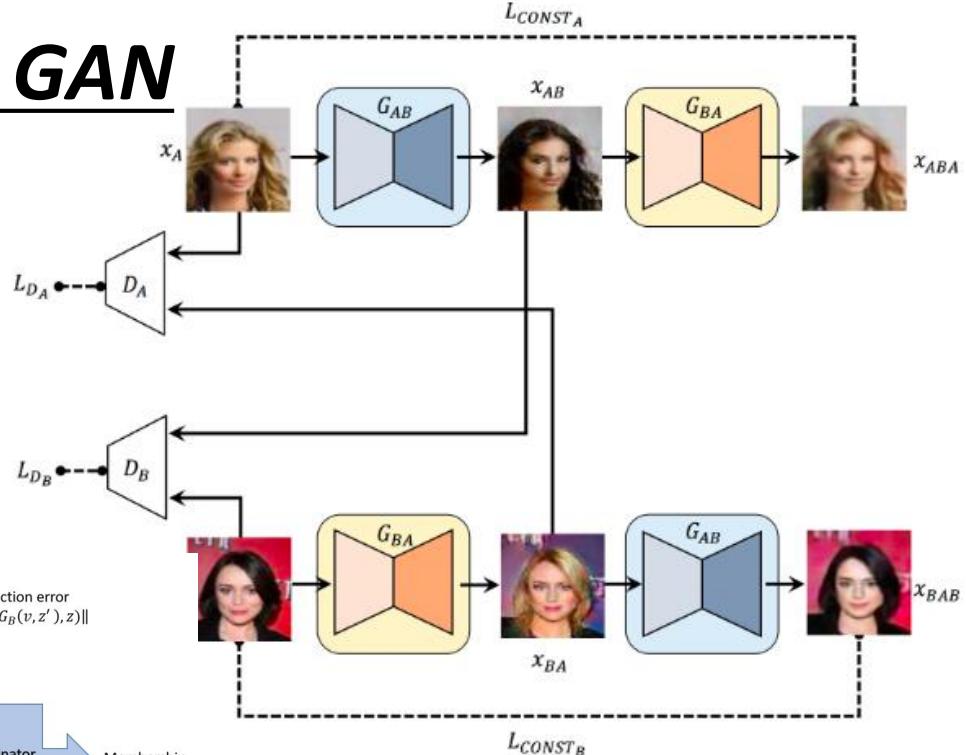


# Cycle GAN



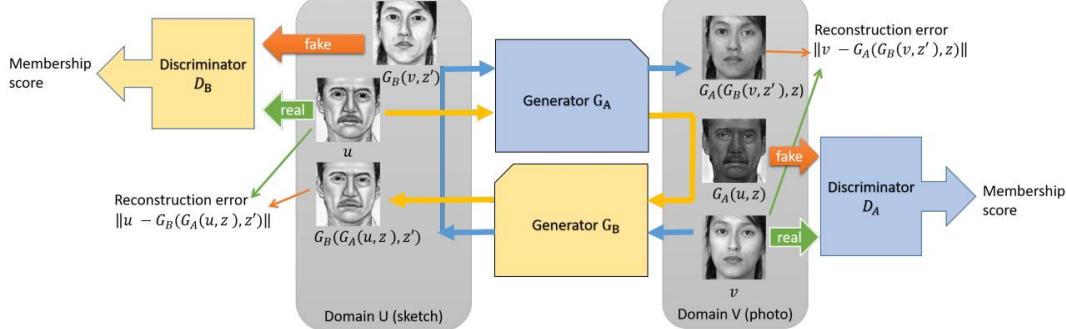
# Disco GAN

<https://arxiv.org/abs/1703.05192>



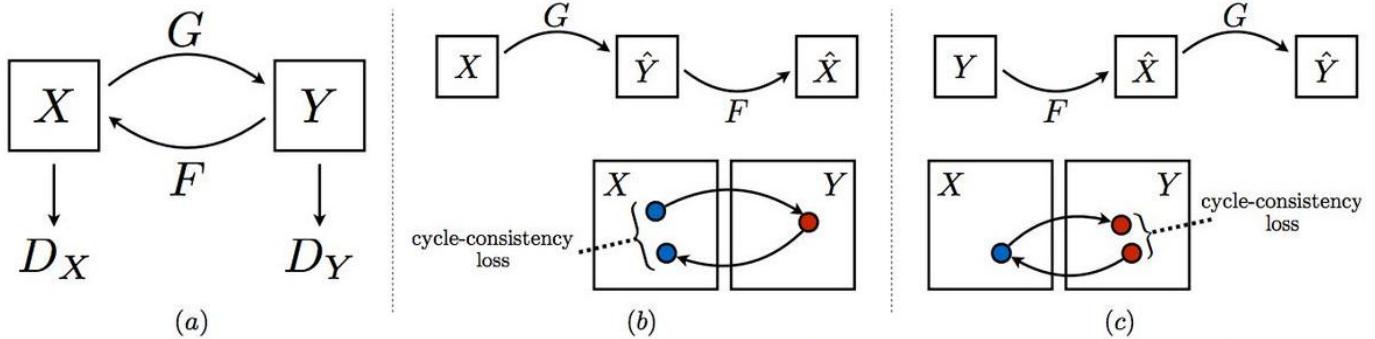
# Dual GAN

<https://arxiv.org/abs/1704.02510>



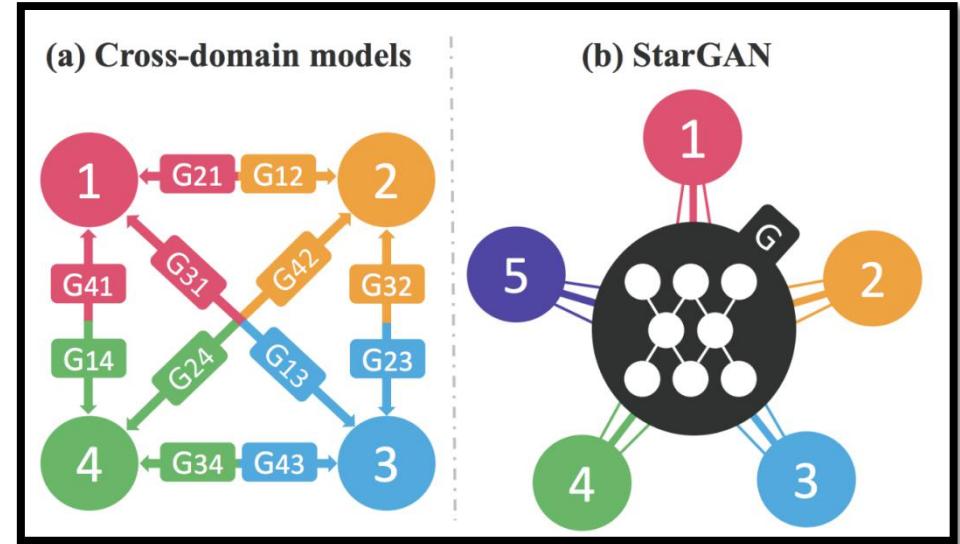
# Cycle GAN

<https://arxiv.org/abs/1703.10593>

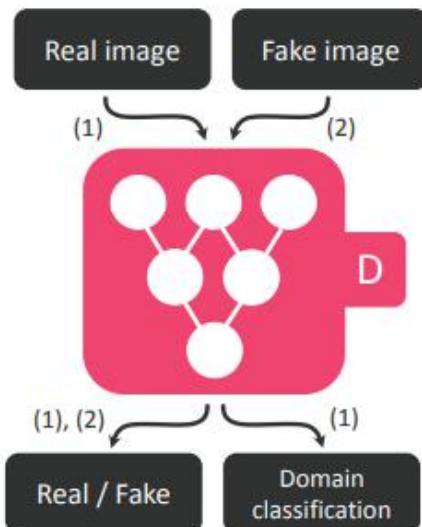


# StarGAN

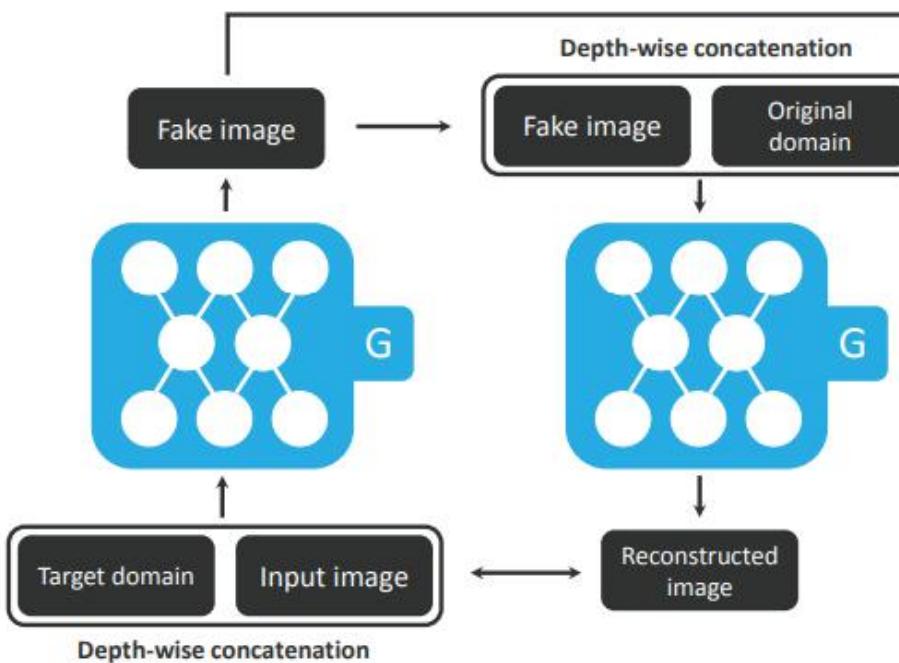
<https://arxiv.org/abs/1711.09020>



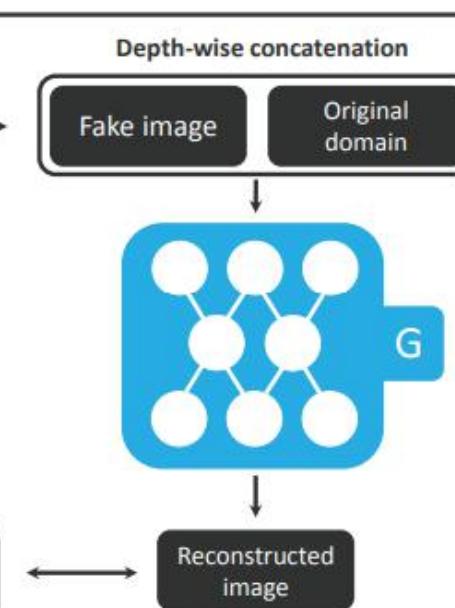
(a) Training the discriminator



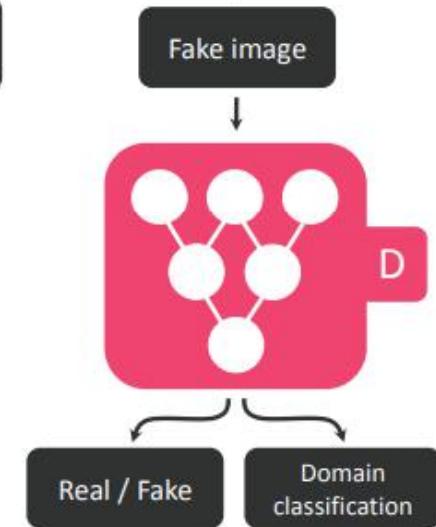
(b) Original-to-target domain



(c) Target-to-original domain



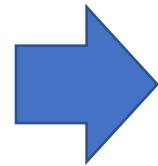
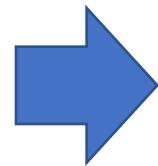
(d) Fooling the discriminator



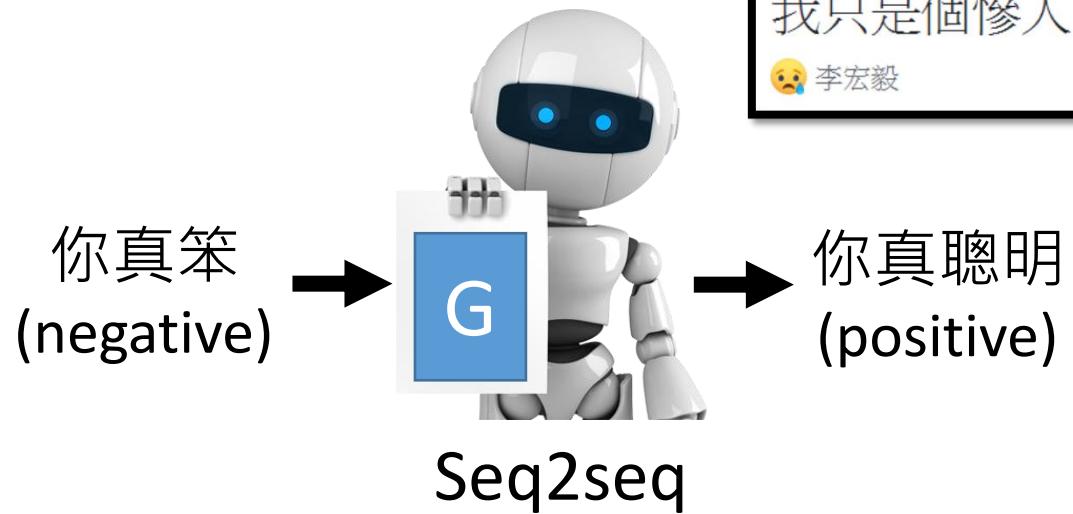
# Q&A

# SELFIE2ANIME

<https://selfie2anime.com/>  
<https://arxiv.org/abs/1907.10830>

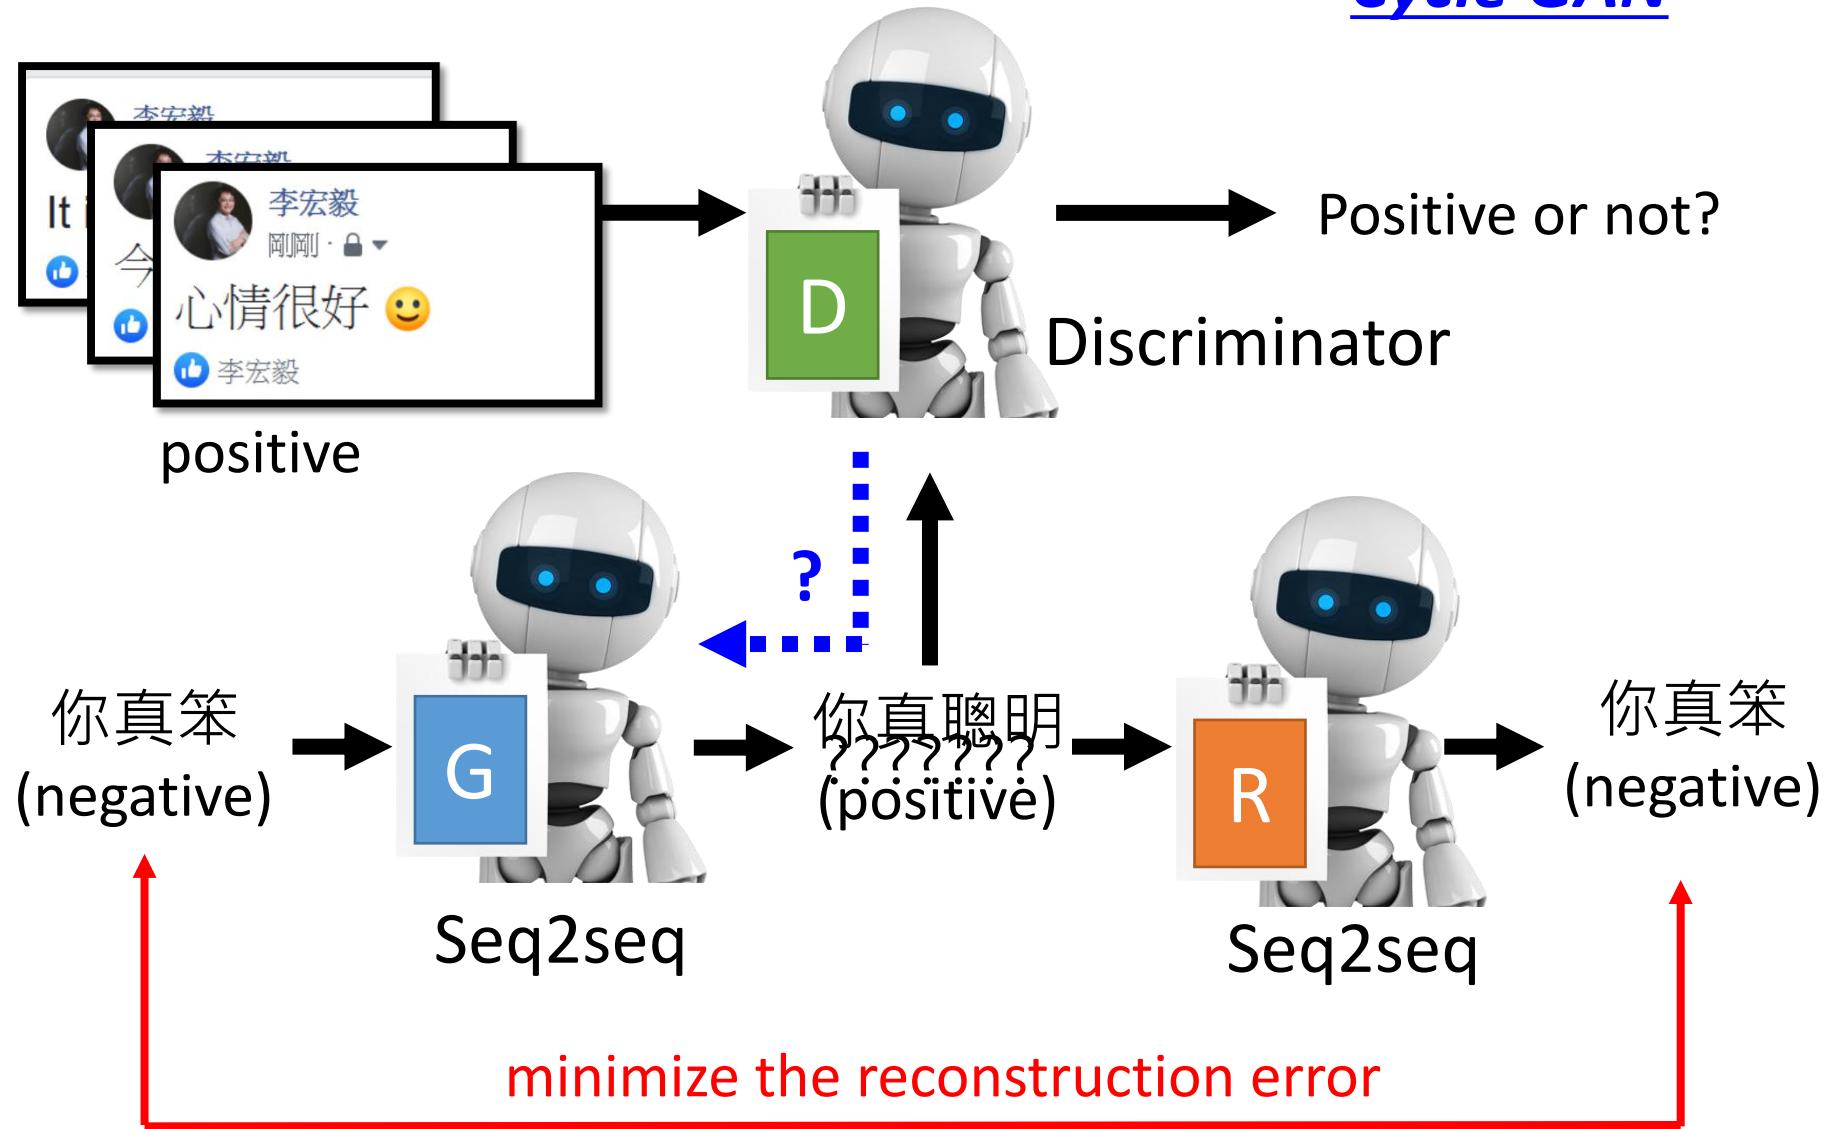


# Text Style Transfer

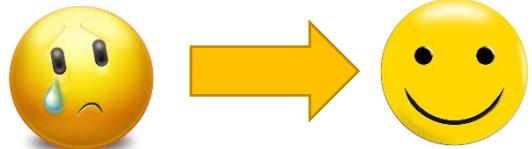


# Text Style Transfer

Cycle GAN



# Text Style Transfer



- From **negative** sentence to **positive** one

胃疼，沒睡醒，各種不舒服

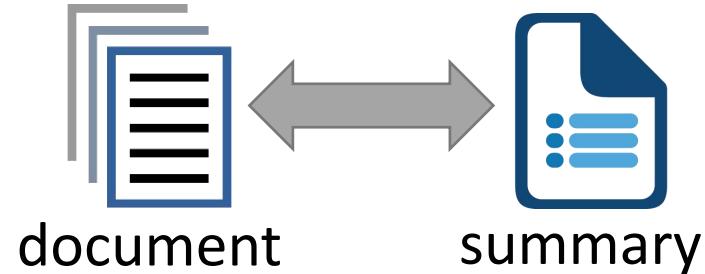
我都想去上班了，真夠賤的！

暈死了，吃燒烤、竟然遇到個變態狂

我肚子痛的厲害

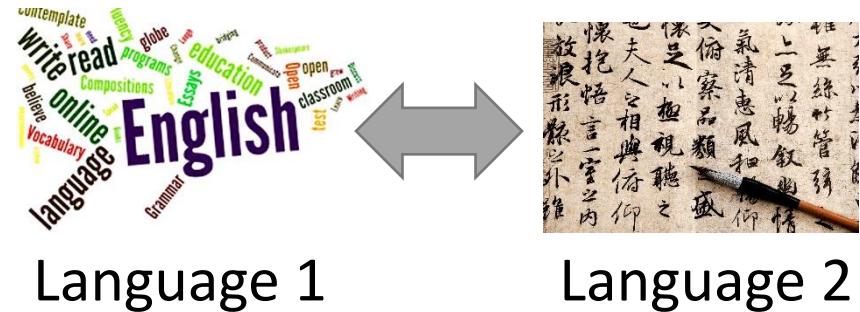
## Unsupervised Abstractive Summarization

<https://arxiv.org/abs/1810.02851>



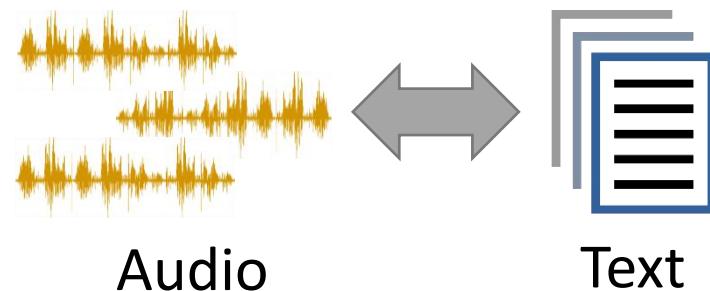
## Unsupervised Translation

<https://arxiv.org/abs/1710.04087>  
<https://arxiv.org/abs/1710.11041>



## Unsupervised ASR

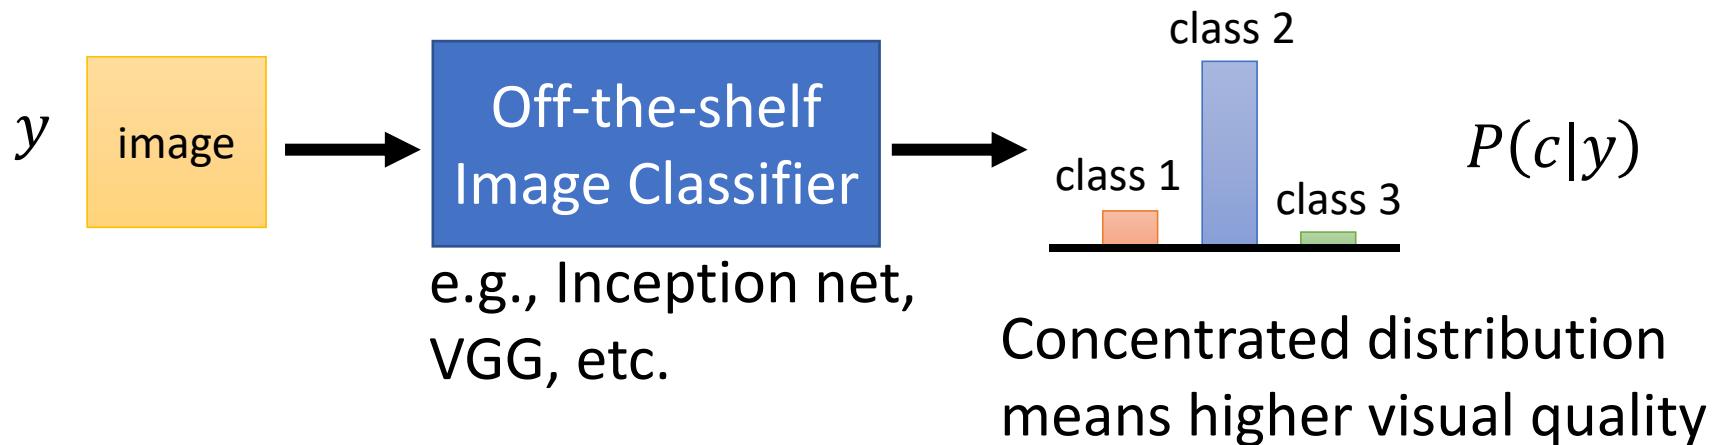
<https://arxiv.org/abs/1804.00316>  
<https://arxiv.org/abs/1812.09323>  
<https://arxiv.org/abs/1904.04100>



# Evaluation of Generation

# Quality of Image

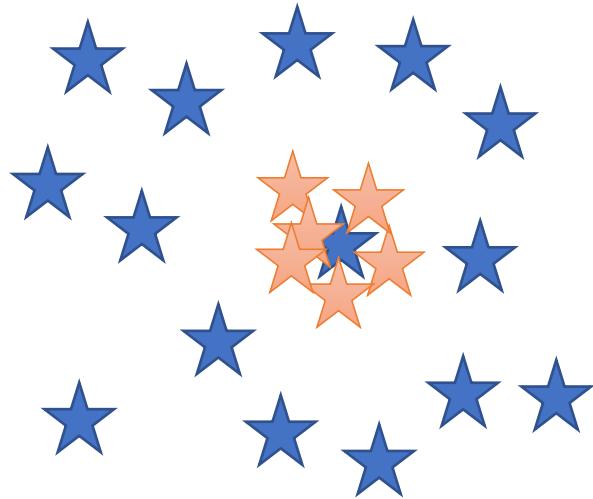
- Human evaluation is expensive (and sometimes unfair/unstable).
- How to evaluate the quality of the generated images automatically?



# Diversity - Mode Collapse

★ : real data

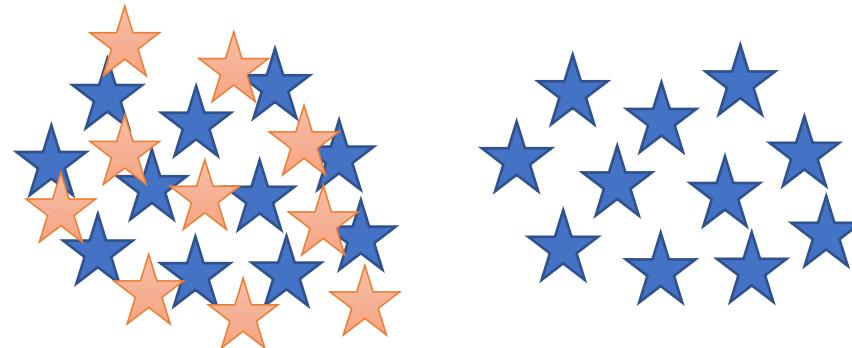
★ : generated data



# Diversity - Mode Dropping

★ : real data

★ : generated data



Generator  
at iteration t

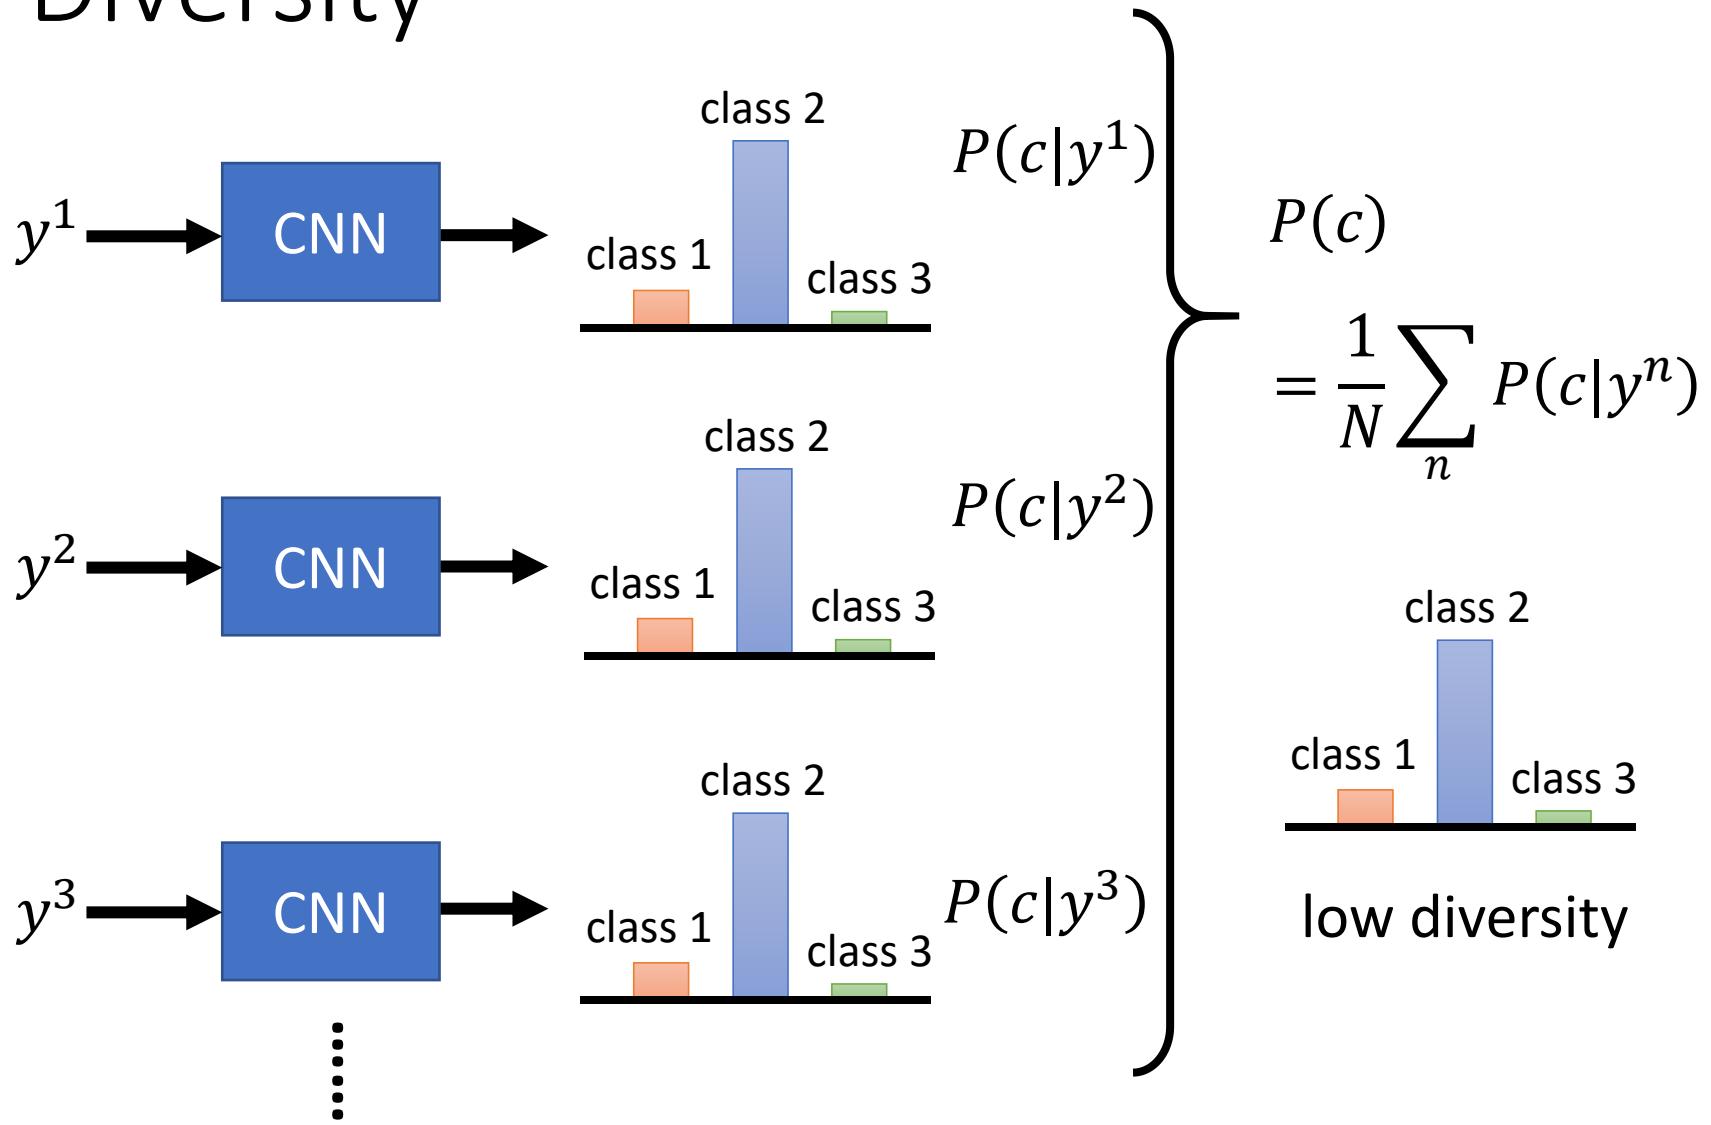


Generator  
at iteration t+1

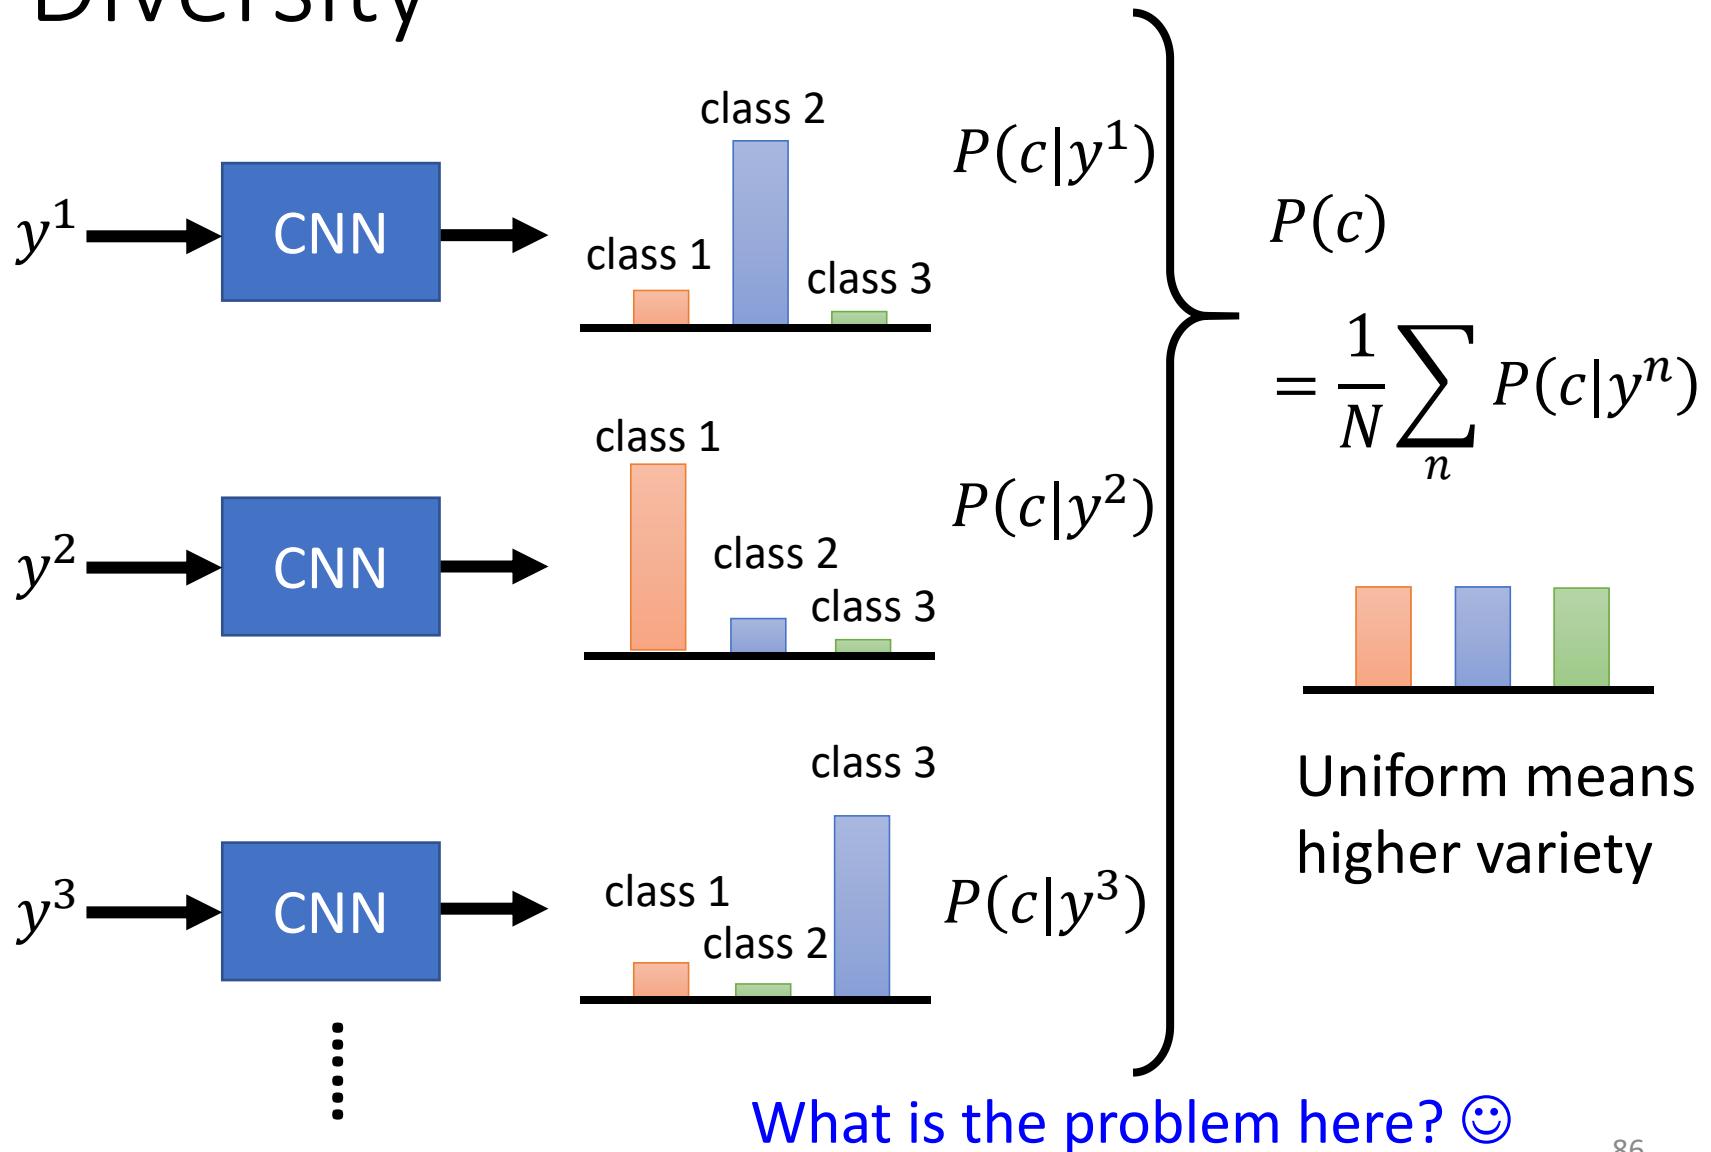


(BEGAN on CelebA)

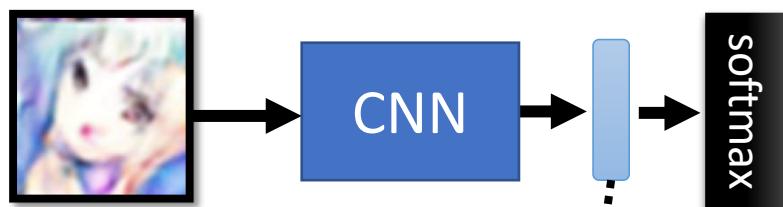
# Diversity



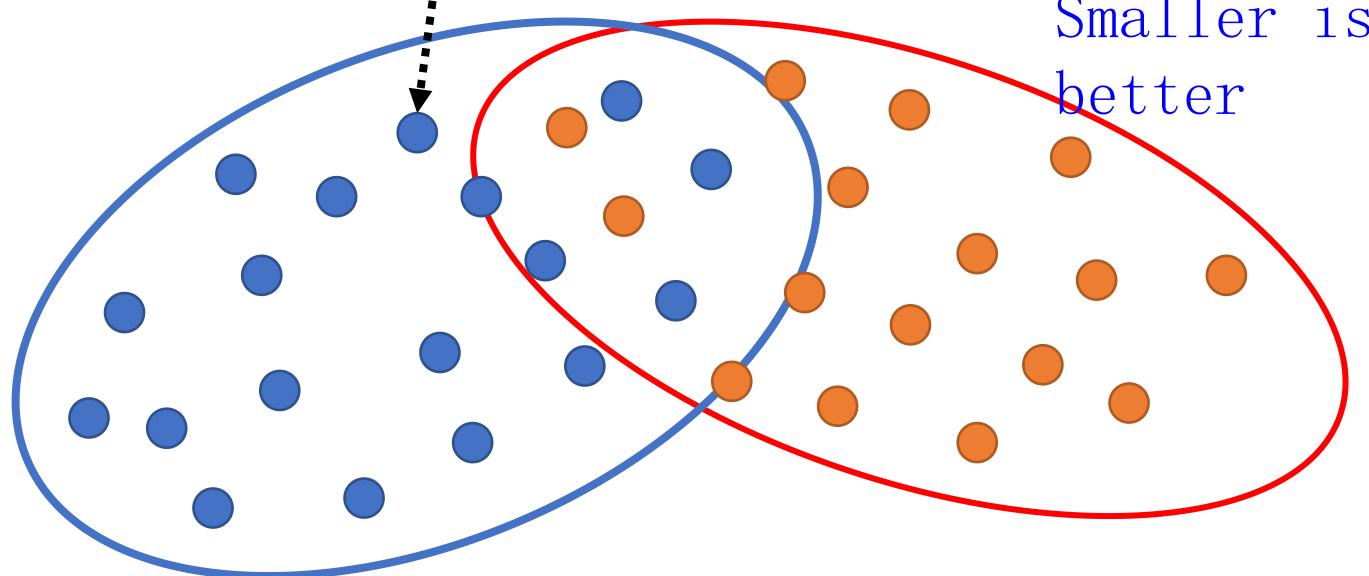
# Diversity



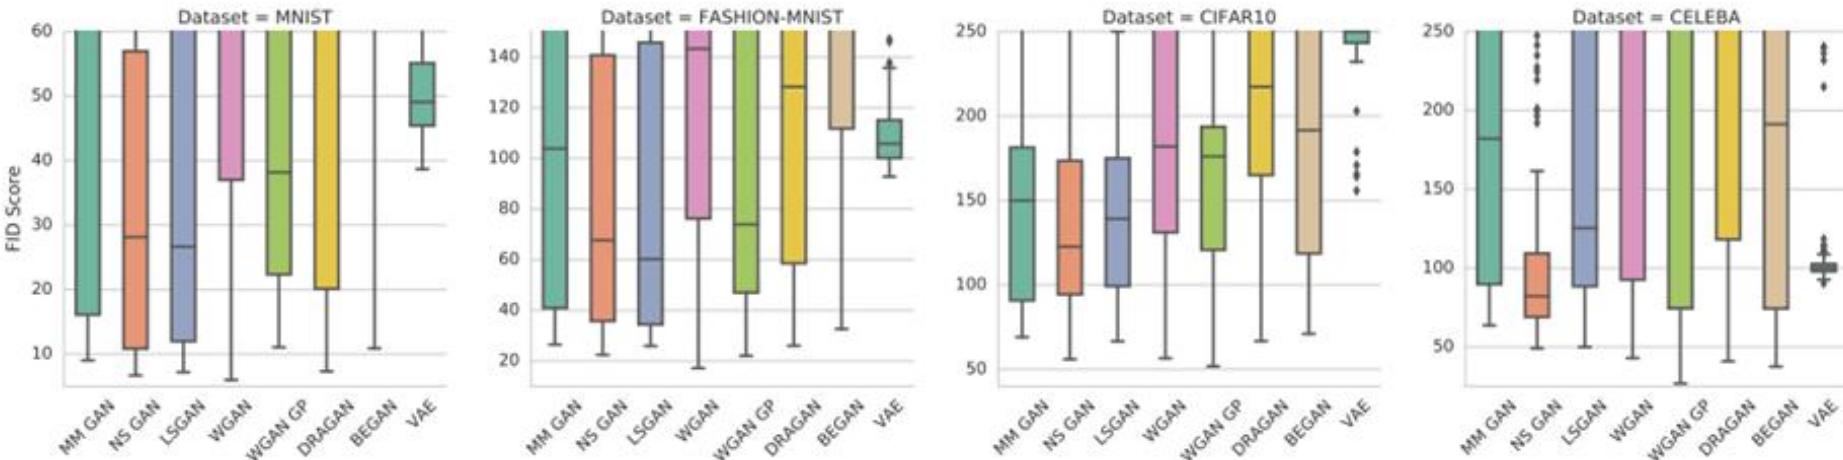
# Fréchet Inception Distance (FID)



red points: real images  
blue points: generated images  
**FID** = Fréchet distance  
between the two **Gaussians** ???



| GAN     | DISCRIMINATOR LOSS  | GENERATOR LOSS  |
|---------|---|---|
| MM GAN  | $\mathcal{L}_D^{GAN} = -\mathbb{E}_{x \sim p_d} [\log(D(x))] + \mathbb{E}_{\hat{x} \sim p_g} [\log(1 - D(\hat{x}))]$                                  | $\mathcal{L}_G^{GAN} = -\mathcal{L}_D^{GAN}$  |
| NS GAN  | $\mathcal{L}_D^{NSGAN} = \mathcal{L}_D^{GAN}$   | $\mathcal{L}_G^{NSGAN} = \mathbb{E}_{\hat{x} \sim p_g} [\log(D(\hat{x}))]$                    |
| WGAN    | $\mathcal{L}_D^{WGAN} = -\mathbb{E}_{x \sim p_d} [D(x)] + \mathbb{E}_{\hat{x} \sim p_g} [D(\hat{x})]$   | $\mathcal{L}_G^{WGAN} = -\mathcal{L}_D^{WGAN}$  |
| WGAN GP | $\mathcal{L}_D^{WGAN} = \mathcal{L}_D^{WGAN} + \lambda \mathbb{E}_{\hat{x} \sim p_g} [(  \nabla D(\alpha x + (1 - \alpha)\hat{x})  _2 - 1)^2]$        | $\mathcal{L}_G^{WGAN} = -\mathbb{E}_{\hat{x} \sim p_g} [D(\hat{x})]$                          |
| LS GAN  | $\mathcal{L}_D^{LSGAN} = -\mathbb{E}_{x \sim p_d} [(D(x) - 1)^2] + \mathbb{E}_{\hat{x} \sim p_g} [D(\hat{x})^2]$                                      | $\mathcal{L}_G^{LSGAN} = -\mathbb{E}_{\hat{x} \sim p_g} [(D(\hat{x}) - 1)^2]$                 |
| DRAGAN  | $\mathcal{L}_D^{DRAGAN} = \mathcal{L}_D^{GAN} + \lambda \mathbb{E}_{\hat{x} \sim p_d + \mathcal{N}(0, c)} [(  \nabla D(\hat{x})  _2 - 1)^2]$          | $\mathcal{L}_G^{DRAGAN} = -\mathcal{L}_D^{NSGAN}$   |
| BEGAN   | $\mathcal{L}_D^{BEGAN} = \mathbb{E}_{x \sim p_d} [      x - AE(x)      _1] - k_t \mathbb{E}_{\hat{x} \sim p_g} [      \hat{x} - AE(\hat{x})      _1]$ | $\mathcal{L}_G^{BEGAN} = \mathbb{E}_{\hat{x} \sim p_g} [      \hat{x} - AE(\hat{x})      _1]$ |

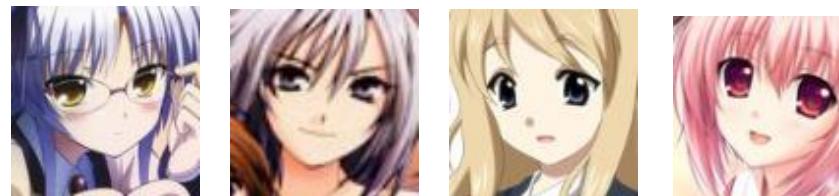


FIT: Smaller is better

Are GANs Created Equal? A Large-Scale Study  
<https://arxiv.org/abs/1711.10337>

# We don't want memory GAN.

Real Data



Generated  
Data



Same as real data ...

Generated  
Data



Simply flip real data ...

# To learn more about evaluation ...

|              | Measure  | Description   |
|--------------|--|---|
| Quantitative | 1. Average Log-likelihood [18, 22]                       | <ul style="list-style-type: none"> <li>• Log likelihood of explaining realworld held out/test data using a density estimated from the generated data (e.g. using KDE or Parzen window estimation). <math>L = \frac{1}{n} \sum_i \log P_{model}(\mathbf{x}_i)</math></li> </ul>  |
|              | 2. Coverage Metric [33]                                  | <ul style="list-style-type: none"> <li>• The probability mass of the true data "covered" by the model distribution <math>C := P_{data}(dP_{model} &gt; t)</math> with <math>t</math> such that <math>P_{model}(dP_{model} &gt; t) = 0.95</math></li> </ul>  |
|              | 3. Inception Score (IS) [3]                              | <ul style="list-style-type: none"> <li>• KLD between conditional and marginal label distributions over generated data. <math>\exp(\mathbb{E}_{\mathbf{x}}[\text{KL}(p(y \mathbf{x}) \  p(y))])</math></li> </ul>  |
|              | 4. Modified Inception Score (m-IS) [34]                  | <ul style="list-style-type: none"> <li>• Encourages diversity within images sampled from a particular category. <math>\exp(\mathbb{E}_{\mathbf{x}_i}[\mathbb{E}_{\mathbf{x}_j}([\text{KL}(P(y \mathbf{x}_i)    P(y \mathbf{x}_j))])])</math></li> </ul>   |
|              | 5. Mode Score (MS) [35]                                  | <ul style="list-style-type: none"> <li>• Similar to IS but also takes into account the prior distribution of the labels over real data. <math>\exp(\mathbb{E}_{\mathbf{x}}[\text{KL}(p(y \mathbf{x}) \  p(y^{train}))] - \text{KL}(p(y) \  p(y^{train})))</math></li> </ul>   |
|              | 6. AM Score [36]   | <ul style="list-style-type: none"> <li>• Takes into account the KLD between distributions of training labels vs. predicted labels, as well as the entropy of predictions. <math>\text{KL}(p(y^{train}) \  p(y)) + \mathbb{E}_{\mathbf{x}}[H(y \mathbf{x})]</math></li> </ul>  |
|              | 7. Fréchet Inception Distance (FID) [37]                 | <ul style="list-style-type: none"> <li>• Wasserstein-2 distance between multi-variate Gaussians fitted to data embedded into a feature space <math>FID(r, g) = \ \mu_r - \mu_g\ _2^2 + Tr(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{\frac{1}{2}})</math></li> </ul>   |
|              | 8. Maximum Mean Discrepancy (MMD) [38]                   | <ul style="list-style-type: none"> <li>• Measures the dissimilarity between two probability distributions <math>P_r</math> and <math>P_g</math> using samples drawn independently from each distribution. <math>M_k(P_r, P_g) = \mathbb{E}_{\mathbf{x}, \mathbf{x}' \sim P_r}[k(\mathbf{x}, \mathbf{x}')] - 2\mathbb{E}_{\mathbf{x} \sim P_r, \mathbf{y} \sim P_g}[k(\mathbf{x}, \mathbf{y})] + \mathbb{E}_{\mathbf{y}, \mathbf{y}' \sim P_g}[k(\mathbf{y}, \mathbf{y}')]</math></li> </ul> |
|              | 9. The Wasserstein Critic [39]                           | <ul style="list-style-type: none"> <li>• The critic (e.g. an NN) is trained to produce high values at real samples and low values at generated samples <math>\hat{W}(\mathbf{x}_{test}, \mathbf{x}_g) = \frac{1}{N} \sum_{i=1}^N \hat{f}(\mathbf{x}_{test}[i]) - \frac{1}{N} \sum_{i=1}^N \hat{f}(\mathbf{x}_g[i])</math></li> </ul>  |
|              | 10. Birthday Paradox Test [27]                           | <ul style="list-style-type: none"> <li>• Measures the support size of a discrete (continuous) distribution by counting the duplicates (near duplicates)</li> </ul>  |
|              | 11. Classifier Two Sample Test (C2ST) [40]               | <ul style="list-style-type: none"> <li>• Answers whether two samples are drawn from the same distribution (e.g. by training a binary classifier)</li> </ul>   |
|              | 12. Classification Performance [1, 15]                   | <ul style="list-style-type: none"> <li>• An indirect technique for evaluating the quality of unsupervised representations (e.g. feature extraction; FCN score). See also the GAN Quality Index (GQI) [41].</li> </ul>   |
|              | 13. Boundary Distortion [42]                             | <ul style="list-style-type: none"> <li>• Measures diversity of generated samples and covariate shift using classification methods.</li> </ul>   |
|              | 14. Number of Statistically-Different Bins (NDB) [43]    | <ul style="list-style-type: none"> <li>• Given two sets of samples from the same distribution, the number of samples that fall into a given bin should be the same up to sampling noise</li> </ul>  |
|              | 15. Image Retrieval Performance [44]                     | <ul style="list-style-type: none"> <li>• Measures the distributions of distances to the nearest neighbors of some query images (i.e. diversity)</li> </ul>  |
|              | 16. Generative Adversarial Metric (GAM) [31]             | <ul style="list-style-type: none"> <li>• Compares two GANs by having them engaged in a battle against each other by swapping discriminators or generators. <math>p(\mathbf{x} y=1; M_1)/p(\mathbf{x} y=1; M_2) = (p(y=1 \mathbf{x}; D_1)p(\mathbf{x}; G_2))/(p(y=1 \mathbf{x}; D_2)p(\mathbf{x}; G_1))</math></li> </ul>  |
|              | 17. Tournament Win Rate and Skill Rating [45]            | <ul style="list-style-type: none"> <li>• Implements a tournament in which a player is either a discriminator that attempts to distinguish between real and fake data or a generator that attempts to fool the discriminators into accepting fake data as real.</li> </ul>   |
|              | 18. Normalized Relative Discriminative Score (NRDS) [32] | <ul style="list-style-type: none"> <li>• Compares n GANs based on the idea that if the generated samples are closer to real ones, more epochs would be needed to distinguish them from real samples.</li> </ul>   |
|              | 19. Adversarial Accuracy and Divergence [16]             | <ul style="list-style-type: none"> <li>• Adversarial Accuracy: Computes the classification accuracies achieved by the two classifiers, one trained on real data and another on generated data, on a labeled validation set to approximate <math>P_g(y \mathbf{x})</math> and <math>P_r(y \mathbf{x})</math>. Adversarial Divergence: Computes <math>\text{KL}(P_g(y \mathbf{x}), P_r(y \mathbf{x}))</math></li> </ul>   |
|              | 20. Geometry Score [47]                                  | <ul style="list-style-type: none"> <li>• Compares geometrical properties of the underlying data manifold between real and generated data.</li> </ul>  |
|              | 21. Reconstruction Error [48]                            | <ul style="list-style-type: none"> <li>• Measures the reconstruction error (e.g. <math>L_2</math> norm) between a test image and its closest generated image by optimizing for <math>z</math> (i.e. <math>\min_z \ G(\mathbf{z}) - \mathbf{x}^{(test)}\ ^2</math>)</li> </ul>   |
|              | 22. Image Quality Measures [49, 50, 51]                  | <ul style="list-style-type: none"> <li>• Evaluates the quality of generated images using measures such as SSIM, PSNR, and sharpness difference</li> </ul>   |
|              | 23. Low-level Image Statistics [52, 53]                  | <ul style="list-style-type: none"> <li>• Evaluates how similar low-level statistics of generated images are to those of natural scenes in terms of mean power spectrum, distribution of random filter responses, contrast distribution, etc.</li> </ul>   |
|              | 24. Precision, Recall and $F_1$ score [23]               | <ul style="list-style-type: none"> <li>• These measures are used to quantify the degree of overfitting in GANs, often over toy datasets.</li> </ul>   |
| Qualitative  | 1. Nearest Neighbors                                     | <ul style="list-style-type: none"> <li>• To detect overfitting, generated samples are shown next to their nearest neighbors in the training set</li> </ul>  |
|              | 2. Rapid Scene Categorization [18]                       | <ul style="list-style-type: none"> <li>• In these experiments, participants are asked to distinguish generated samples from real images in a short presentation time (e.g. 100 ms); i.e. real v.s fake</li> </ul>   |
|              | 3. Preference Judgment [54, 55, 56, 57]                  | <ul style="list-style-type: none"> <li>• Participants are asked to rank models in terms of the fidelity of their generated images (e.g. pairs, triples)</li> </ul>  |
|              | 4. Mode Drop and Collapse [58, 59]                       | <ul style="list-style-type: none"> <li>• Over datasets with known modes (e.g. a GMM or a labeled dataset), modes are computed as by measuring the distances of generated data to mode centers</li> </ul>  |
|              | 5. Network Internals [1, 60, 61, 62, 63, 64]             | <ul style="list-style-type: none"> <li>• Regards exploring and illustrating the internal representation and dynamics of models (e.g. space continuity) as well as visualizing learned features</li> </ul>   |

## Pros and cons of GAN evaluation measures

<https://arxiv.org/abs/1802.03446>

# Concluding Remarks

Introduction of Generative Models

Generative Adversarial Network (GAN)

Theory behind GAN

Tips for GAN

Conditional Generation

Learning from unpaired data

Evaluation of Generative Models

**Q&A**